

Tech Together Powered by Dify

Dify. AI Developer Ecosystem Director

Liu Xinrui



Traction & Impact

POPULARITY

90,000+

GitHub Stars

LOCAL DEPLOYMENT

3,725,000+

Docker Pulls

OPEN-SOURCE CONTRIBUTION

700+

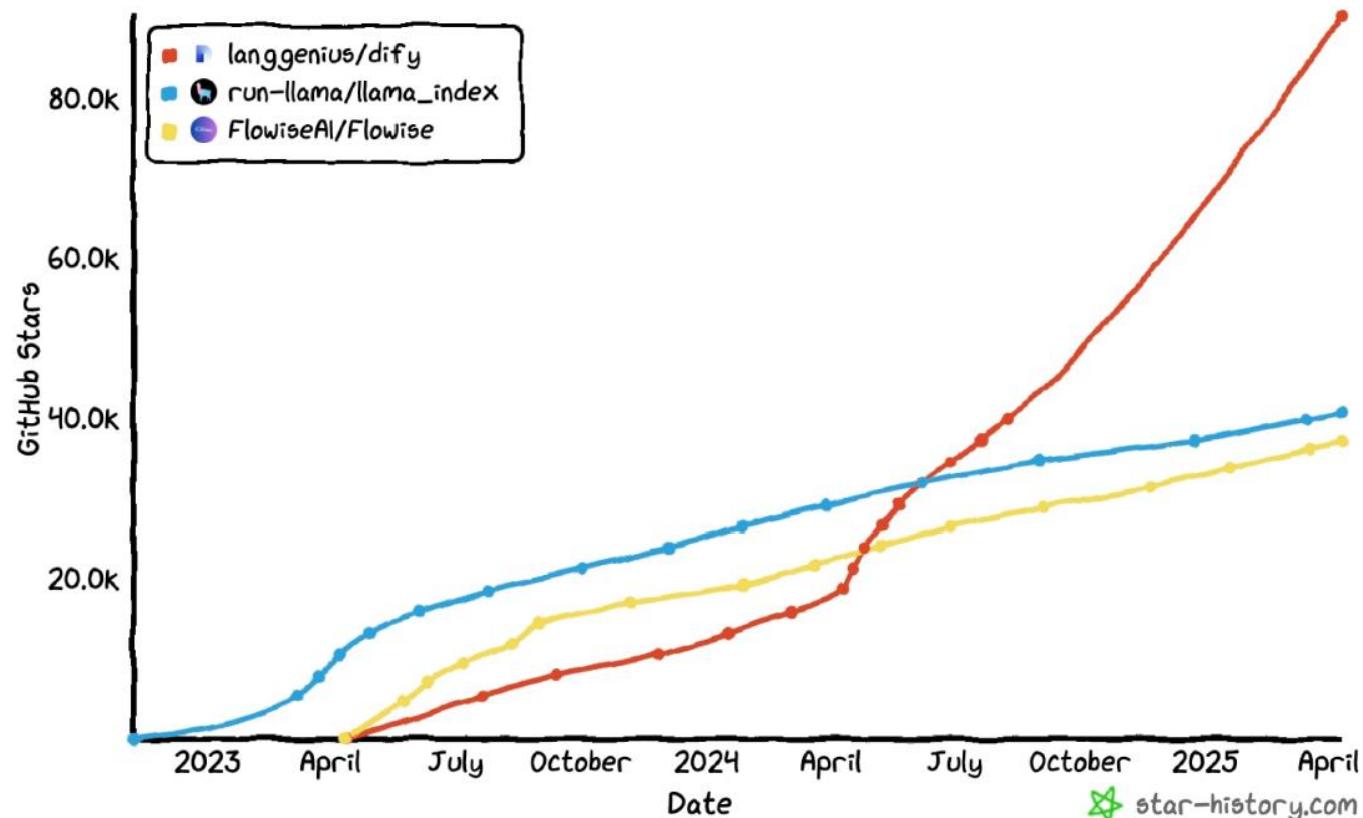
Contributors

BUILD APPS

800,000+

Apps Powered by Dify

Star History



Trusted and Used by Leading Enterprises Across Industries Worldwide

VOLVO

Panasonic

infineon

Yum!

ANKER

Deloitte.

ThermoFisher
SCIENTIFIC

PERSOL

MAERSK

NOVARTIS

GLENCORE

TCS
TATA
CONSULTANCY
SERVICES

SHAKUDO

ADVANTECH

価格.com

Lilly

NTTEAST

ETS

McDonald's

RICOH

CEQUENS

GovTech 東京

Midea

TDSE

JIP
日本電子計算株式会社

MUFG

Daiichi-Sankyo

AIA



Built on Trust and Compliance

Since day one of product design, Dify has made user data security, privacy protection, and compliance a top priority.



SOC2 Type I & Type II ✓



ISO 27001:2022 ✓



GDPR ✓

**All LLM Stacks
You Need in One Place**

Dify is a powerful open-source platform for developing GenAI applications. Experience seamless integration of **AI workflows, RAG pipelines, agent capabilities, model management, built-in observability** and more. Build production-ready application from day one.

Dify Across the Stack

Innovation Engine for
GenAI Applications

POWERED BY DIFY

Industry & Solution

Automobile

Electronics

Semiconductors

Education

Consulting

E-commerce

Finance

Use case & Scenarios

Automation

Customer Support

Marketing

Knowledge Search

Sales

and More...

DIFY

Service Layer

Backend-as-Service

Backendless frontend application development via API integration.

Apps

Chatbot

Completion

Agent

Chatflow

Workflow

Embedded in Website

Batch Run

Orchestration Layer

ORCHESTRATION

Agent Workflow

Custom agents that equipped with various tools to solve complex tasks, working side-by-side.

Prompt IDE

Efficient prompt design, testing, optimization, and management.

RAG Pipeline

Improve the quality and accuracy of LLM responses by retrieving relevant external knowledge.

MONITORING

LLMOps

Comprehensive tracking and deep evaluation for full LLM application lifecycle support.

COLLABORATION

Dify Standard DSL

Collaborate effectively with Dify standard DSL, a YAML-based file format designed for seamless workflow sharing and teamwork.

GENAI INFRA ECOSYSTEM

PLUGINS ECOSYSTEM

Connection Layer

Plugins

Enhance AI application capabilities by integrating models and tools.

Models Provider, MaaS, Self-host



INFRASTRUCTURE & DATA

Infrastructure



Datasource



Dify Agentic Workflow

Visually build and orchestrate your AI applications, from sketches into reality.



Cross-Multimodal Support

Optimal AI models (LLM, TTS, STT, etc.) for every task and business need.



Precise Memory Management

Conversation variables store temporary data, enabling multi-round dialogues.



Connect, Share, and Collaborate

Natively embed workflows within workflows. Share with the community and collaborate with your team.



Fault-Tolerant & High-Performance

Error handling, featuring backup branches and retries ensures workflow continuity.

The screenshot displays the Dify Agentic Workflow interface. The top navigation bar includes 'EXPLORE', 'STUDIO / CHATBOT PLAYGROUND', 'KNOWLEDGE', 'TOOLS', and 'PLUGINS'. The main workspace shows a workflow diagram with nodes: 'DOC EXTRACTOR' (retrieving 'sys.files File'), 'LLM' (invoking 'GPT-4o'), and 'CODE' (containing a 'Fail Branch'). A 'Test Run#5' panel on the right shows the execution log:

STEP	TIME	STATUS
START	0.02 s	✓
DOC EXTRACTOR	0.27 s	✓
LLM	258 tokens · 0.12 s	✓
CODE	1.53 s	✗

The 'CODE' node error message states: "Node exception, will automatically execute the fail branch. The node output will return an error type and error message and pass them to downstream." The 'INPUT' for the 'CODE' node is:

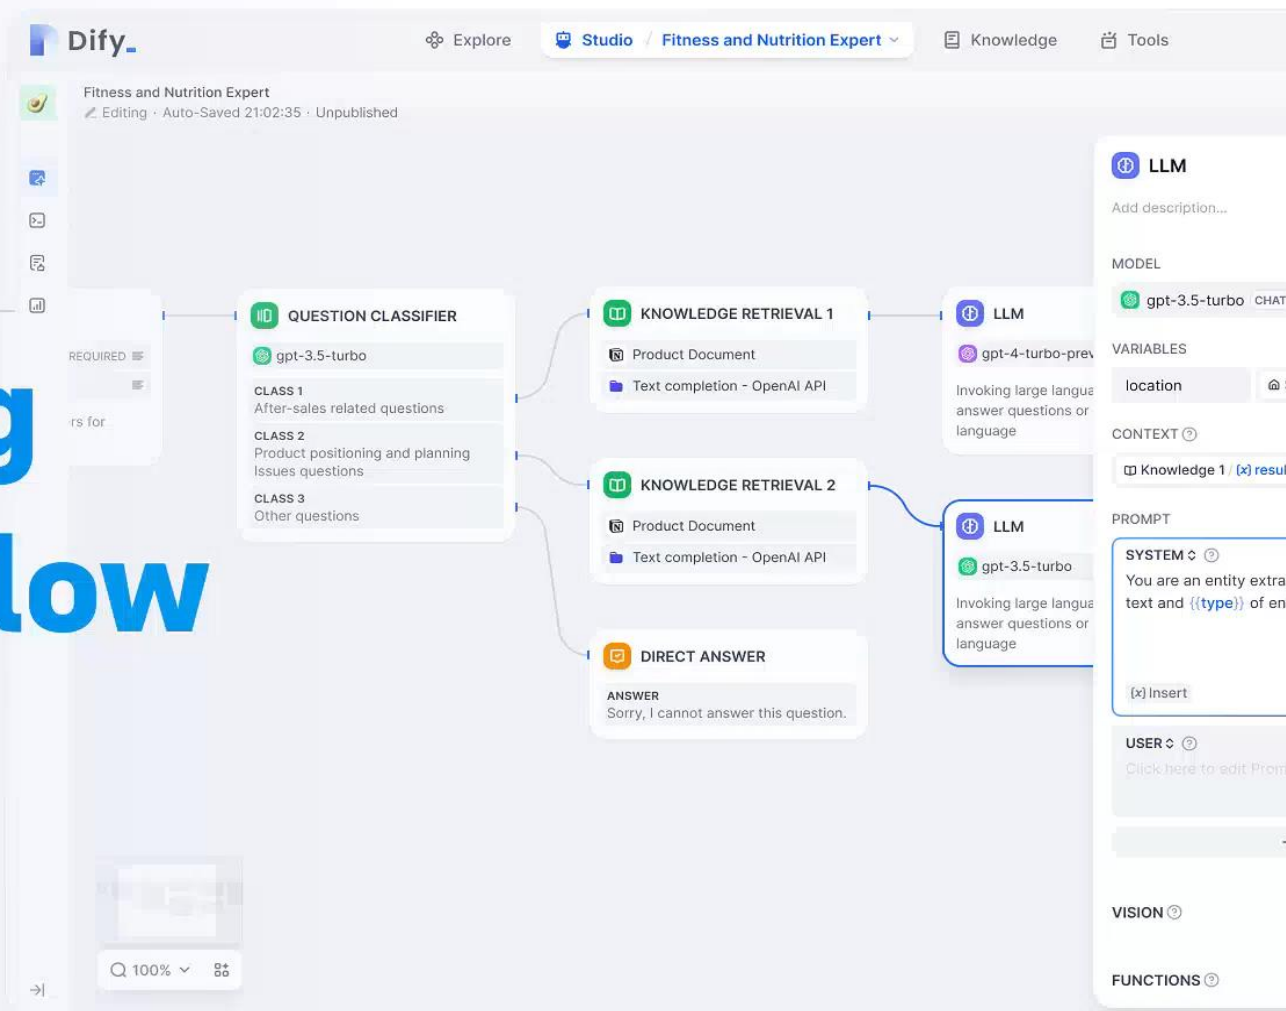
```
{
  "query": "What model are you?"
}
```

The 'OUTPUT' for the 'CODE' node is:

```
{
  "error type": "moduleTypeError",
  "error message": "system error"
}
```

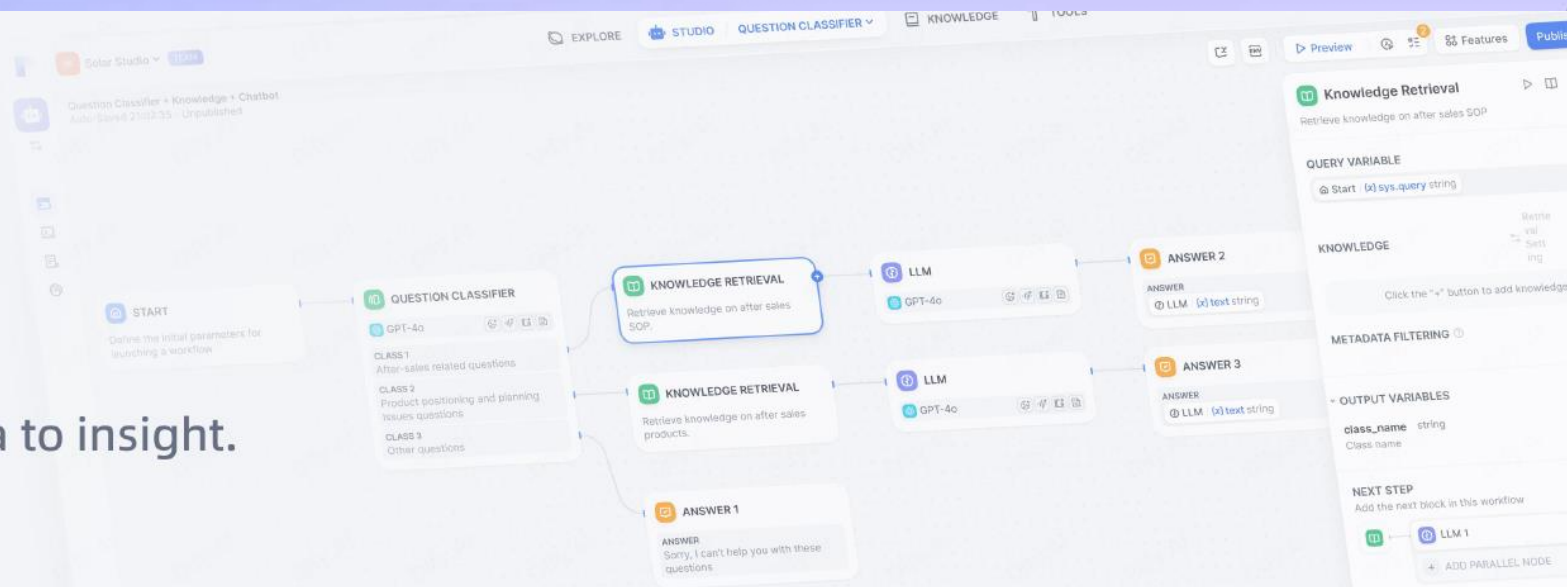



Introducing Dify Workflow



RAG Pipeline

Beyond search:
transform your data to insight.



Integrate Diverse Data Sources

Upload local files in various formats, import pages from Notion, sync content from websites, or connect to external knowledge base like AWS Bedrock.



Get Your Data LLM-ready

Prepare your data for optimal LLM performance with Dify's RAG Pipeline—extract, process, transform and index seamlessly in various structures.



Modify, Organize, Update

Experience seamless knowledge base management with intuitive preview and editing, along with multiple APIs for effortless modification and updates.



Advanced Search Engine

Combine with LLMs for query classification, retrieve from multipath with metadata filtering and hybrid search, optimize results with Reranker, and cite sources for clear reference.

Model Capability in Dify

Unleash innovation with global LLMs.



Support 1000+ LLMs

Integrate thousands of proprietary/open-source LLMs from dozens of inference providers and self-hosted solutions.



LLM API Load Balance

Optimize API requests distribution to increase model invocation capacity for fast response and stable performance.



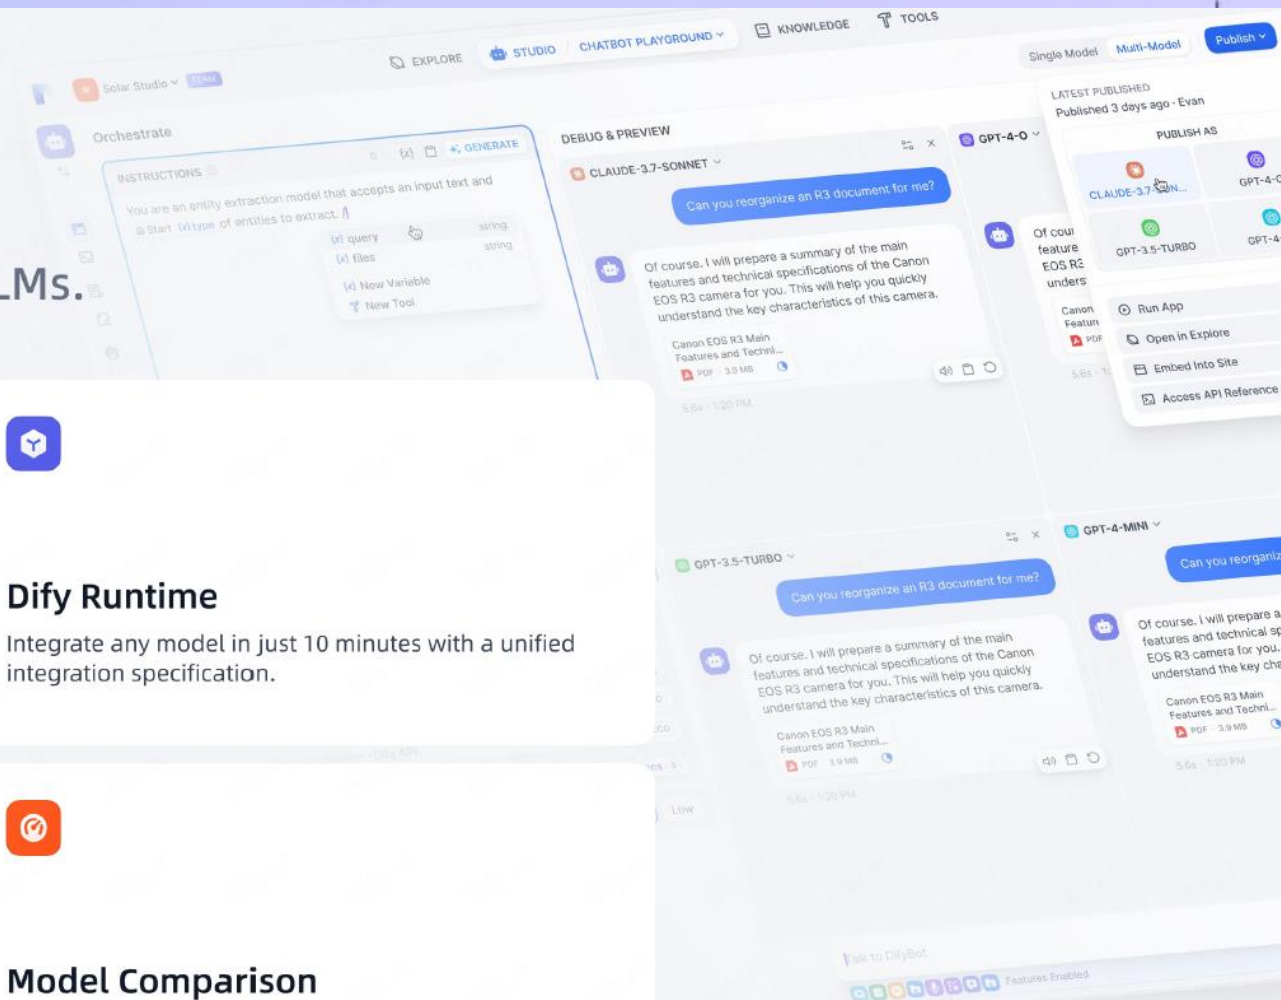
Dify Runtime

Integrate any model in just 10 minutes with a unified integration specification.



Model Comparison

Compare performance from multiple LLMs to the same query, enabling you to select the most suitable model based on your needs.



Marketplace / Tool Integration

Unlock boundless potential of your AI applications by integrating plugins from across the industries.

Empower your AI development

Discover [models](#), [tools](#), [Agent Strategy](#), [extensions](#) and [bundles](#) in Dify Marketplace

All Tags Search all 8677 plugins...

All Models Tools Agent Strategy Extensions Bundles

Featured

Our top picks to get you started

Notion Page Search  notion / notion-page-search TOOL

Search Notion pages and open visited ones faster. No admin access required.

31,711 · # Productivity # Utilities # Search

DuckDuckGo Search  duckduckgo / duckduckgo-search EXTENSION

A privacy-focused search engine.

126,720 · # Productivity # Utilities # Search

Google Search  google / google-search TOOL

Integrate real-time information Google Search directly into your Dify workflows.

6,711 · # Productivity # Utilities # Search

Microsoft  microsoft

Connect to Microsoft s automate tasks, and er

56,798 · # Product

WolframAlpha  wolframalpha / wolframalpha EXTENSION

WolframAlpha is a powerful computational knowledge engine.

502 · # Productivity # Utilities # Search

Wikipedia  wikipedia / wikipedia EXTENSION


Wikipedia is a free online encyclopedia, created and edited by volunteers around the world.

1,235 · # Productivity # Utilities # Search

DALL-E  openai / dall-e TOOL

DALL-E is an AI system that can create realistic images and art from a description in natural language.

2,721 · # Productivity # Utilities # Search

Bing Search  microsoft

Search Notion pages a admin access required

216,712 · # Produc


Popular

Explore the library and discover the incredible work of our community

Zapier  zapier / zapier-zzz TOOL


AI gives you automation superpowers, and Zapier puts them to work.

3,216 · # Productivity # Utilities # Search

YouTube  google / youtube EXTENSION

Retrieve YouTube video info, and integrate video content directly into your workflows.

3,216 · # Productivity # Utilities # Search

Adobe Firefly  adobe / adobe-firefly MODEL

Firefly models and services power generative AI features in Adobe creative apps.

905 · # Productivity # Utilities # Search

Copy to  figma / figma


Automate design tasks

182 · # Productivit

Gmail Checker  google / Gmail-checker EXTENSION

Instantly monitor your inbox with real-time notifications and message previews without leaving your current t...

182 · # Productivity # Utilities # Search

Google Calendar  google / google-calendar EXTENSION

Helps you register your work to Google Calendar

892 · # Productivity # Utilities # Search

Linear  organization / plugin-name TOOL

The Linear extension brings the speed, quality and joy of the app to every corner of your Mac.

892 · # Productivity # Utilities # Search

Moonshot  moonshot

Revolutionizing Conver

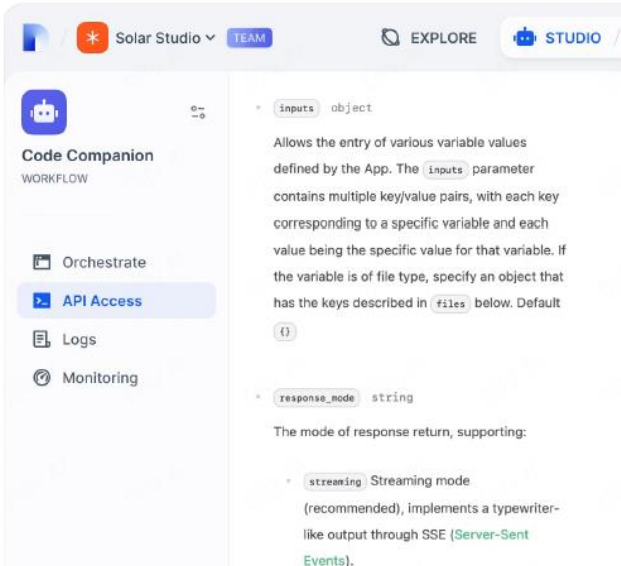
905 · # Productiv

Application Distribution

Flexible options for various needs.

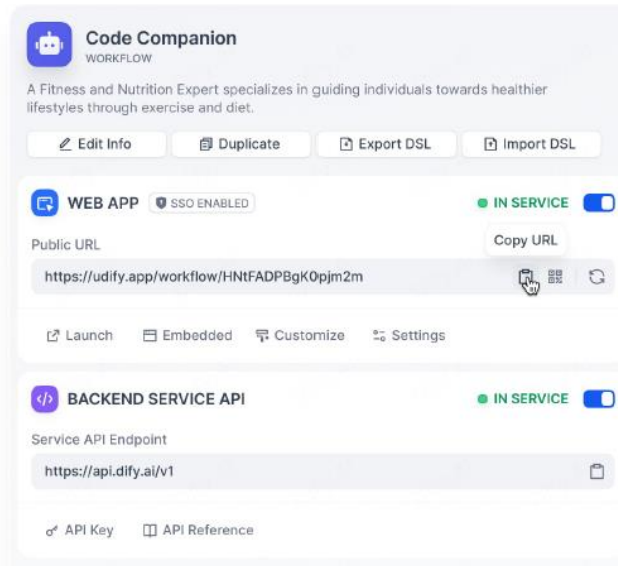
Integrate into Business Systems via API

Seamlessly integrate into your existing business systems for process optimization and operational efficiency.



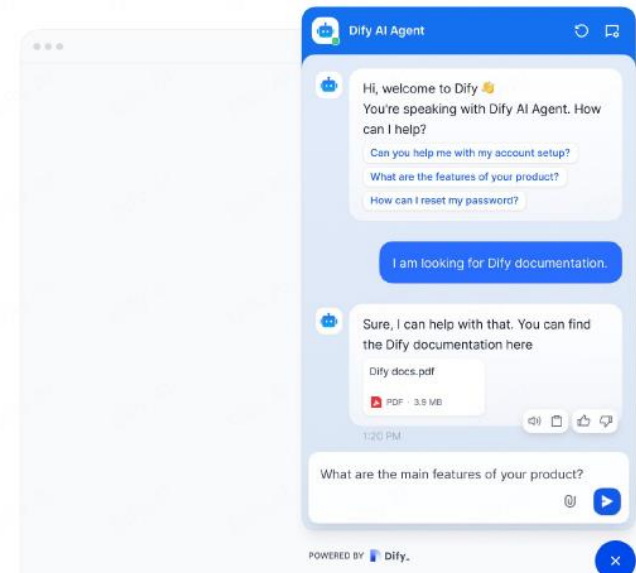
Share via WebApp to Reach a Wider Audience

Launch and make your AI applications easily accessible to anyone, anywhere.



Embed Applications into Web Pages for Better User Experience

Transform your website with AI-powered customer services in minutes.



LLMOps

Bridging the gap:
from experiment to production.

Experiment & Evaluation

Leverage third-party LLMOps tools for a comprehensive traceability of an entire application lifecycle.

LangSmith IN SERVICE

Tracing IN SERVICE

Capture the full context of app execution, including LLM calls, context, prompts, HTTP requests, and more, to a third-party tracing platform.

CONFIGURED

LangSmith IN USE Config

An all-in-one developer platform for every step of the LLM-powered application lifecycle.

MORE PROVIDER

Langfuse Config

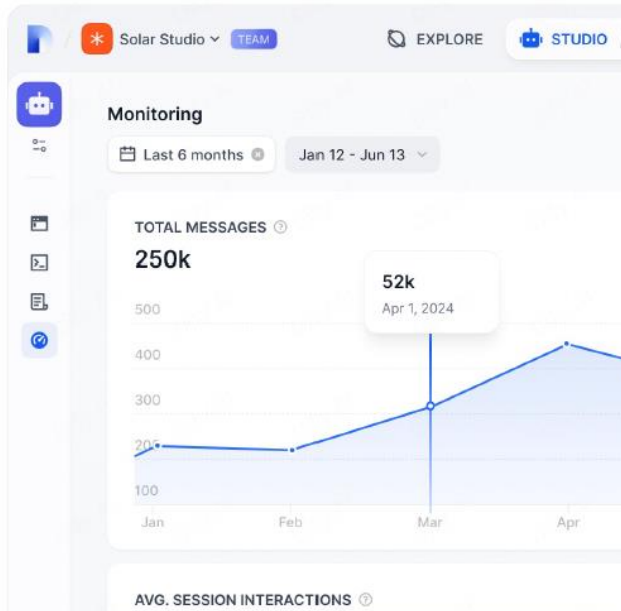
Traces, evals, prompt management and metrics to debug and improve your LLM application.

Opik Config

Opik is an open-source platform for evaluating, testing, and monitoring LLM applications.

Monitoring & Analysis

Gain insights from visualized metrics such as usage costs, latency, user feedback, etc. in the production environment.



Human Feedback & Iteration

Track application performance with logs. Guarantee high-quality responses with manual editing and annotation.

Logs

The logs record the running state of the application.

Last 7 days

ANNOTATED RESPONSE Hit History 3 Hits

User Query

Let's play a decryption game today. You go

[Edit](#)

Storyteller Bot

Lara, the Caesar cipher is a simple substitut technique, where each letter in the alphabe forward or backward a fixed number of posi shifting the letters back 13 positions.

[Your Answer](#)

Type your answer here

[Save](#) [Cancel](#)

[Remove this Cache](#) CREATED AT

Can Dify integrate with my existing CRM? SUCCESS 94924171-e600-

For Enterprise

Enterprise-grade AI infrastructure: powering the future of business.



High Availability

Support stable production for enterprise-grade applications.



Security & Compliance

Enterprise-grade security and compliance to protect critical data assets.



SSO Management (SAML, OIDC, OAuth2)

Centrally manage organizational member access to the Dify Enterprise platform and all web applications using enterprise-grade Single Sign-On (SSO).



Multi-Tenant Management

Data isolation between different groups within a team through multiple Workspaces.



Audit Logs

Comprehensive, auditable, and immutable operation log.



Fast Deployment, Private Control

Deploy to the cloud (AWS, Azure) or on-premises with one-click.

For Startup

Accelerate your business growth with Dify.



AI podcast by dify

Auto-Saved 21:37:16 · Published 2 days ago

Preview

Features

Publish

Hi there, I will introduce you to how to use this chatflow. This chatflow involves file uploading, file parsing, file referencing, and file generation. First, you need to configure the podcast's tone, guest name, podcast materials, and language in the start form. Since files are uploaded through a custom file type in the start form, there's no need to open the features in the upper right corner.

The doc extractor is a new node that supports the parsing of document types. File type variables cannot be used directly by the LLM; They must first be converted into strings or string arrays (when multiple files are uploaded) through the document parser before they can be utilized.

The inspiration for the prompt comes from the open-source project opennotebook: <https://github.com/gabrielchua/open-notebook/blob/main/prompts.py#L8>. First, analyze the insights of the material and brainstorm some ideas for generating subsequent scripts. Then, create a podcast script and finally generate a podcast summary based on that script, using template nodes to piece them together.

For the podcast generator, the specific implementation is carried out using OpenAI's TT model. Each line of the generated script will be converted to speech one by one to allow for a sw in tone, and ultimately stitched together into a complete audio file. Therefore, the script must n the requirement: it should contain strings of alternating lines for two hosts, with each line separated by a newline character.

START

(x) file

REQUIRED

(x) tone

REQUIRED

(x) host_name

REQUIRED

(x) guest_name

REQUIRED

(x) language

REQUIRED

DOC EXTRACTOR

INPUT VARIABLE

Start (x) file

ANALYZE THE INPUT

claude-3-5-son... CHAT

CRAFT THE DIALOGUE

claude-3-5-son... CHAT

DRAFT

ANSWER

Craft th... (x) text

CONCLUSION

claude-3-5-son... CHAT

CONCLUSION

ANSWER

conclusi... (x) text

PREVIEW

file

Local upload

Paste file link

tone

host_name

Please input

guest_name

Please input

language

Click the button to start generating the podcast after uploading the file 📁 And choose the corresponding podcast length, tone, as well as the language of the podcast and the name of the host.

click me!

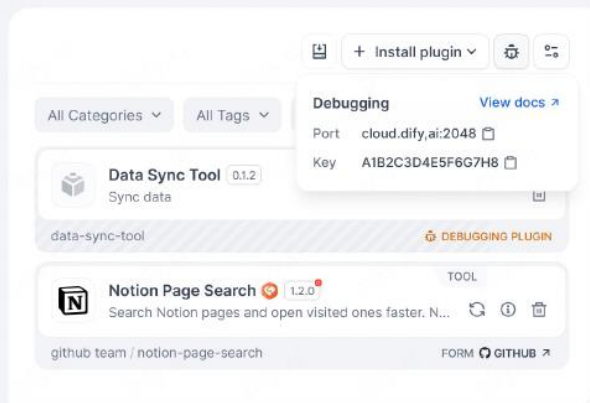
Talk to Bot

Features Enabled

Manage

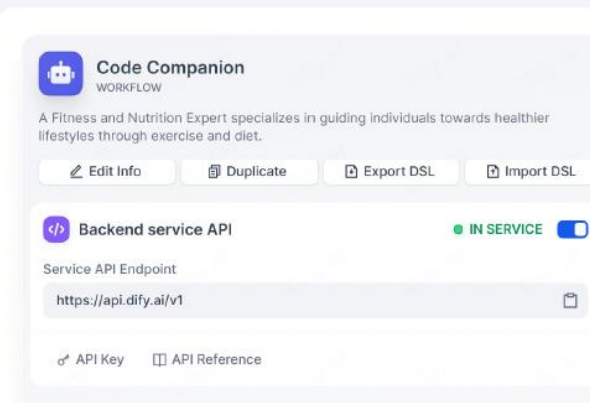
For Developers

Accelerating AI development cycle.



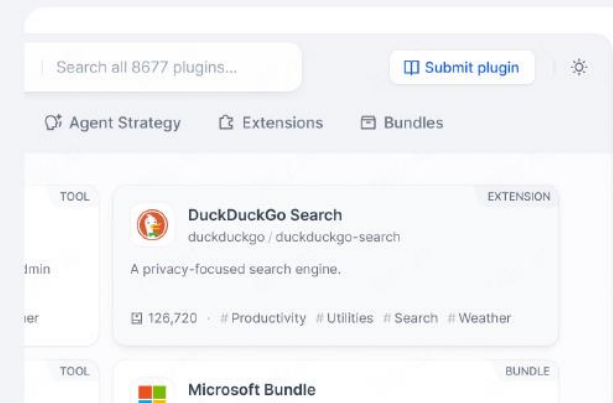
Validation & Debugging

Build AI applications quickly without hard coding and optimize with built-in debugging tools.



Backend as a Service

Developers can quickly and securely access model capabilities and switch model providers with Dify API, without building backend from the scratch.



Plugin extension

Create and submit new plugins to marketplace to enhance development efficiency without touching Dify's source code.



DIFY IN ACTION

ThermoFisher
SCIENTIFIC

Thermo Fisher Scientific is a global leader and supplier of analytical instruments, clinical development solutions, specialty diagnostics, laboratory, pharmaceutical and biotechnology services.



Challenges

The technical support team at Thermo Fisher receives hundreds of support calls for its diverse catalog of scientific instruments each day.

- Extracting information from support calls, organizing tickets, and completing ticket forms is time-consuming.
- Manual ticket processing leads to high repetition and inefficient information transfer.



Solutions

Using Dify, a single IT solutions engineer was able to quickly prototype and deploy to production a ticket auto-filling workflow.



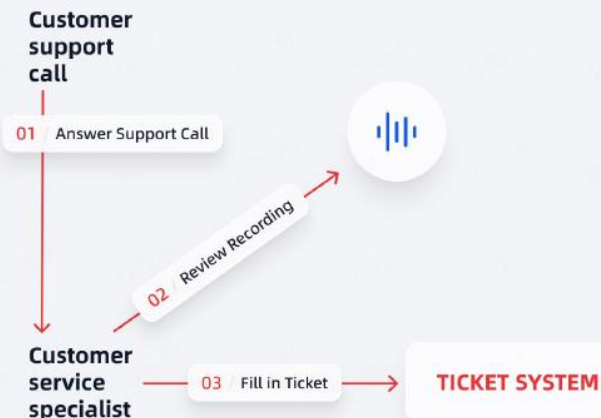
ThermoFisher
SCIENTIFIC

DIFY IN ACTION
One of our customer cases

BEFORE

Manually creating support tickets for diverse catalogs of scientific instruments is time-consuming for customer service specialist. They spend significant time extracting necessary information from lengthy customer calls and often revisit recordings to complete ticket forms.

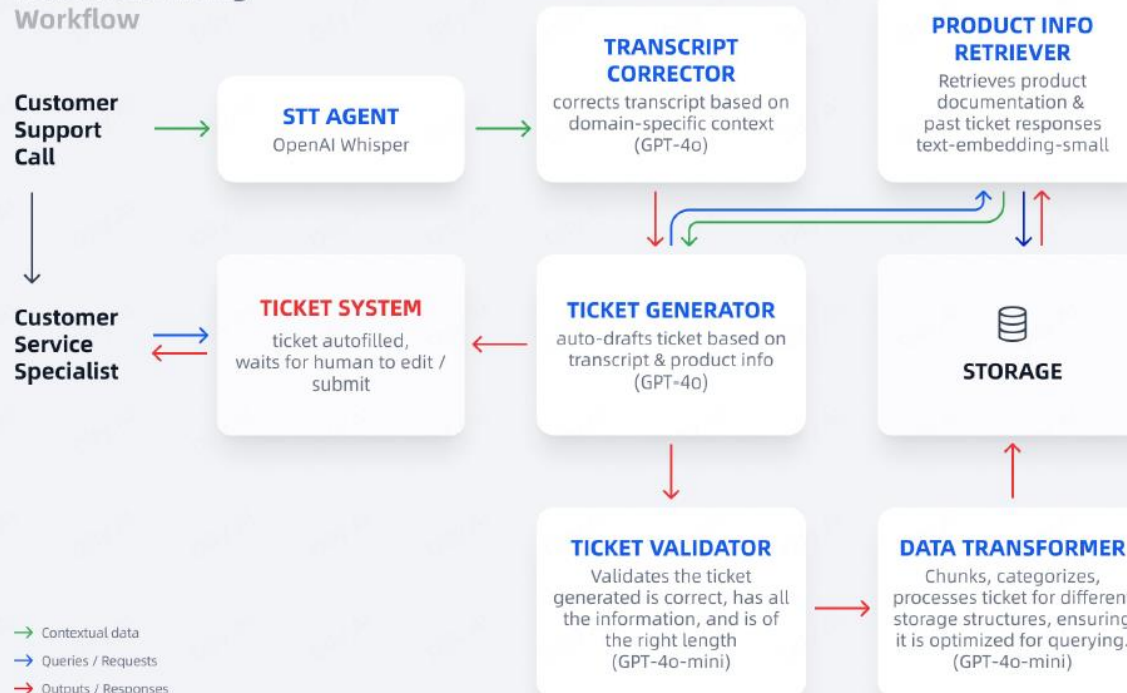
Manually Creating Support Tickets



AFTER

The ticket auto-filling workflow creates structured support tickets from voice calls. Large language models provide intelligent correction, boosting efficiency and accuracy. Human oversight guarantees quality, **saving approximately 200,000 man-hours of work per month.**

Ticket Auto-filling Workflow



DIFY IN ACTION

**MAERSK**

Maersk is an integrated logistics company that offers supply chain solutions for managing shipments and cargo.



Challenges

- Large and complex organizational structure leads to scattered information, making retrieval time-consuming
- Manual review are inefficient, and historical review records are not well-documented



Solutions

- Enterprise-level Q&A bot built with Dify for effective information access and query across 20+ departments, 19,000+ employees
- A lease contract pre-audit robot analyzes previous contracts and stores approved terms in a knowledge base, providing support for future lease agreement reviews



MAERSK

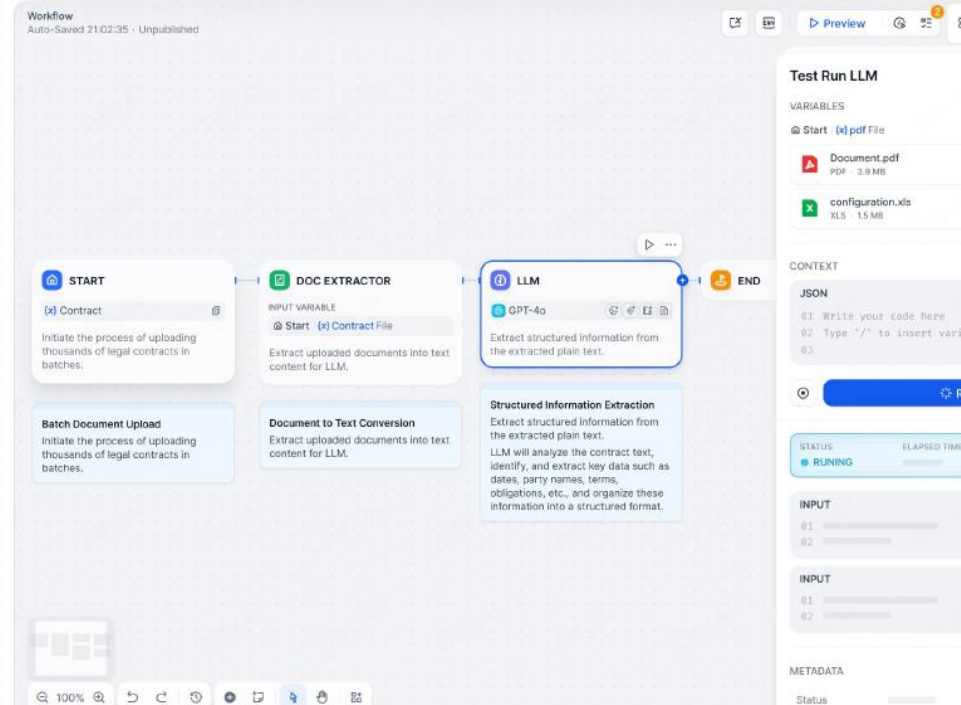
BEFORE

Maersk team needs to process and manage thousands of contracts every day, which brings them management pressure and challenges.



AFTER

- ✓ **Boost Contract Processing Efficiency**
Automated workflows drastically cut processing time and accelerate business response.
- ✓ **Unlock Structured Contract Data**
From massive files to structured data format, establishing foundation for future analysis, retrieval, and strategic decision-making.

DIFY IN ACTION
One of our customer cases

DO IT FOR YOU

THANK YOU

