Make your LLM Serverless

Guillaume Blaquiere
@gblaquiere

Google Developer Expert - Cloud
Group Data Architect - Carrefour

Agenda

- Mystic computing
- Video Game
- The Andes Trail
- What’s the weather

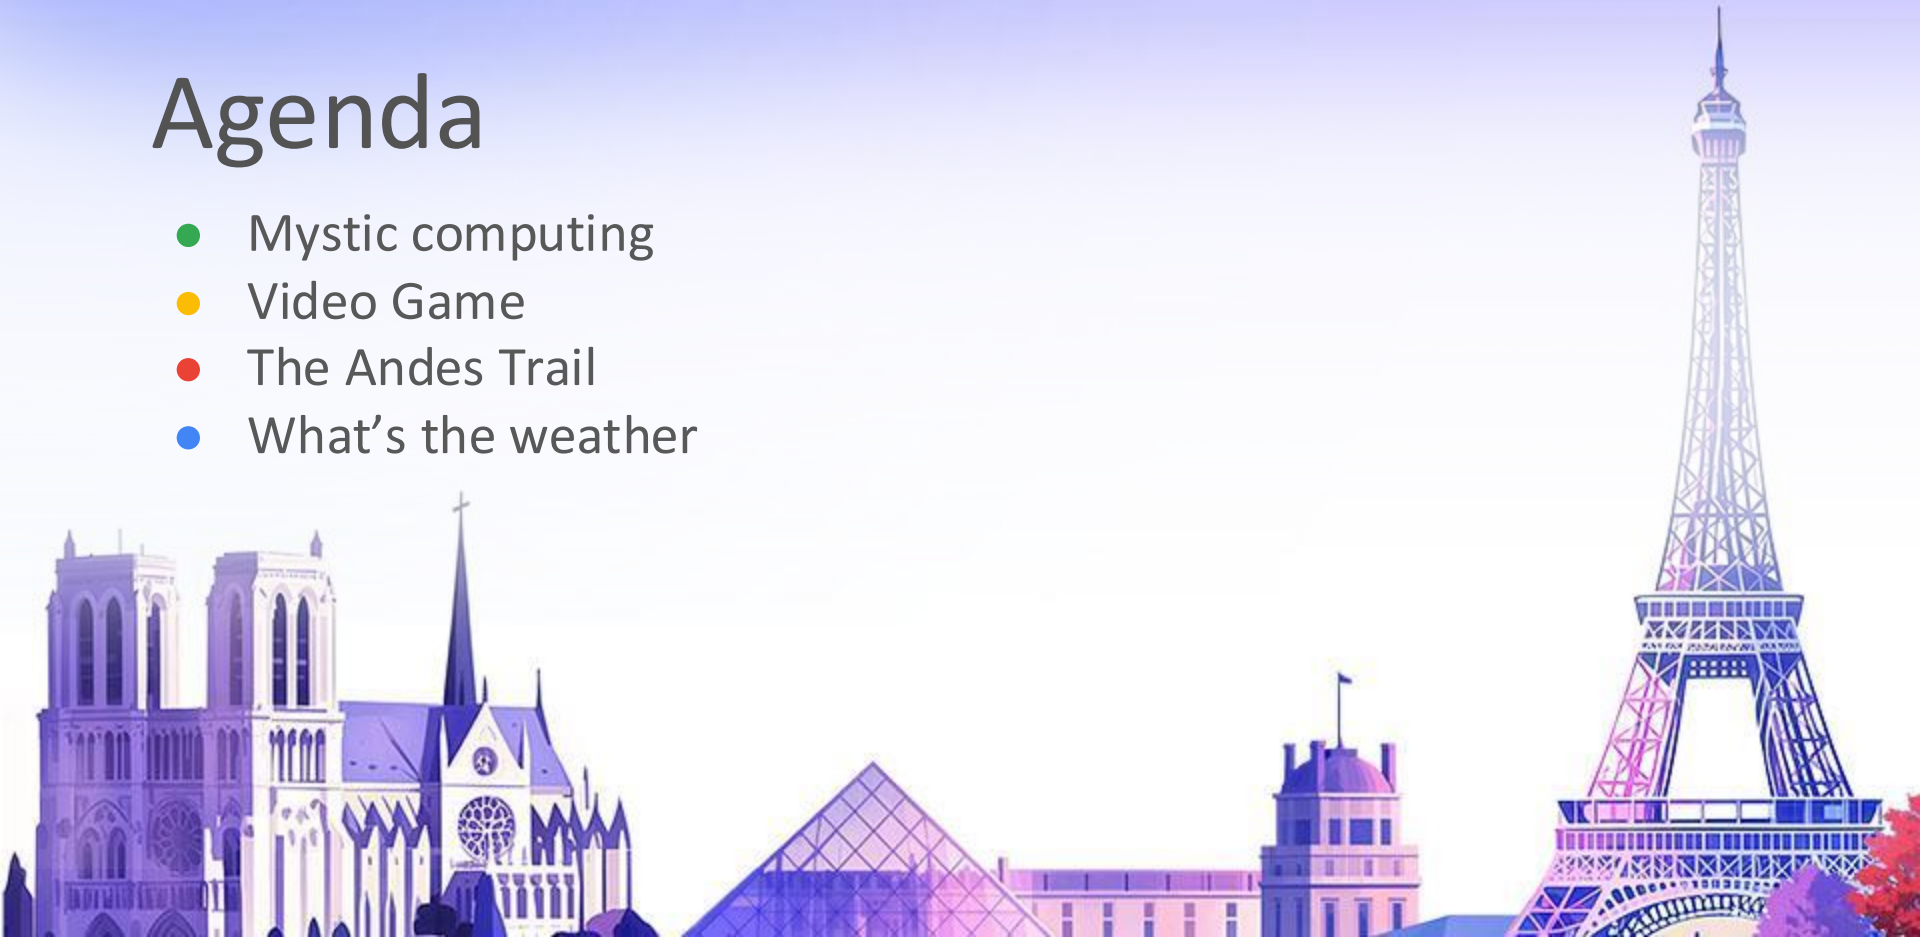

GOSIM

GOSIM AI Paris 2025

Mystic Computing

Serverless computing

A shared responsibility model

Run your application without **worrying about** the servers

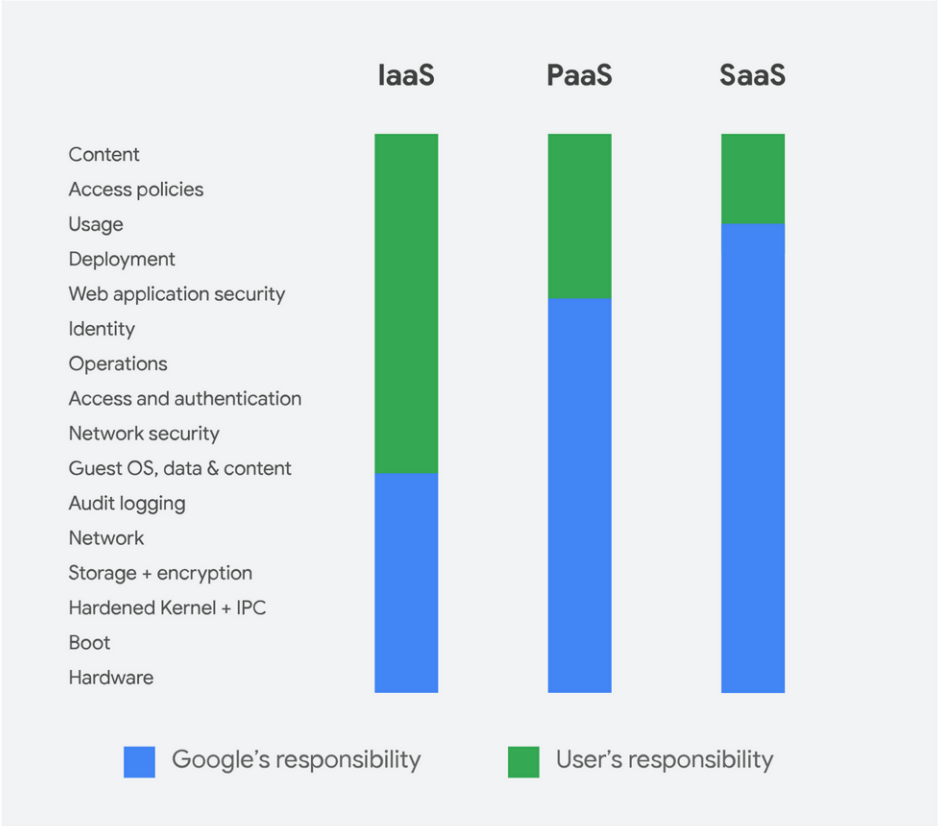

Cloud run service

Serverless container platform

Zero config deployments
gcloud run deploy

Auto-scaling
to support
peak traffic spikes

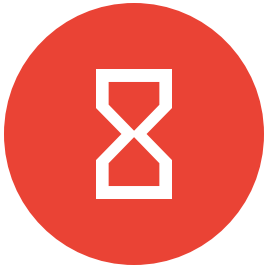

Pay only
while your code runs

Video Game foundations

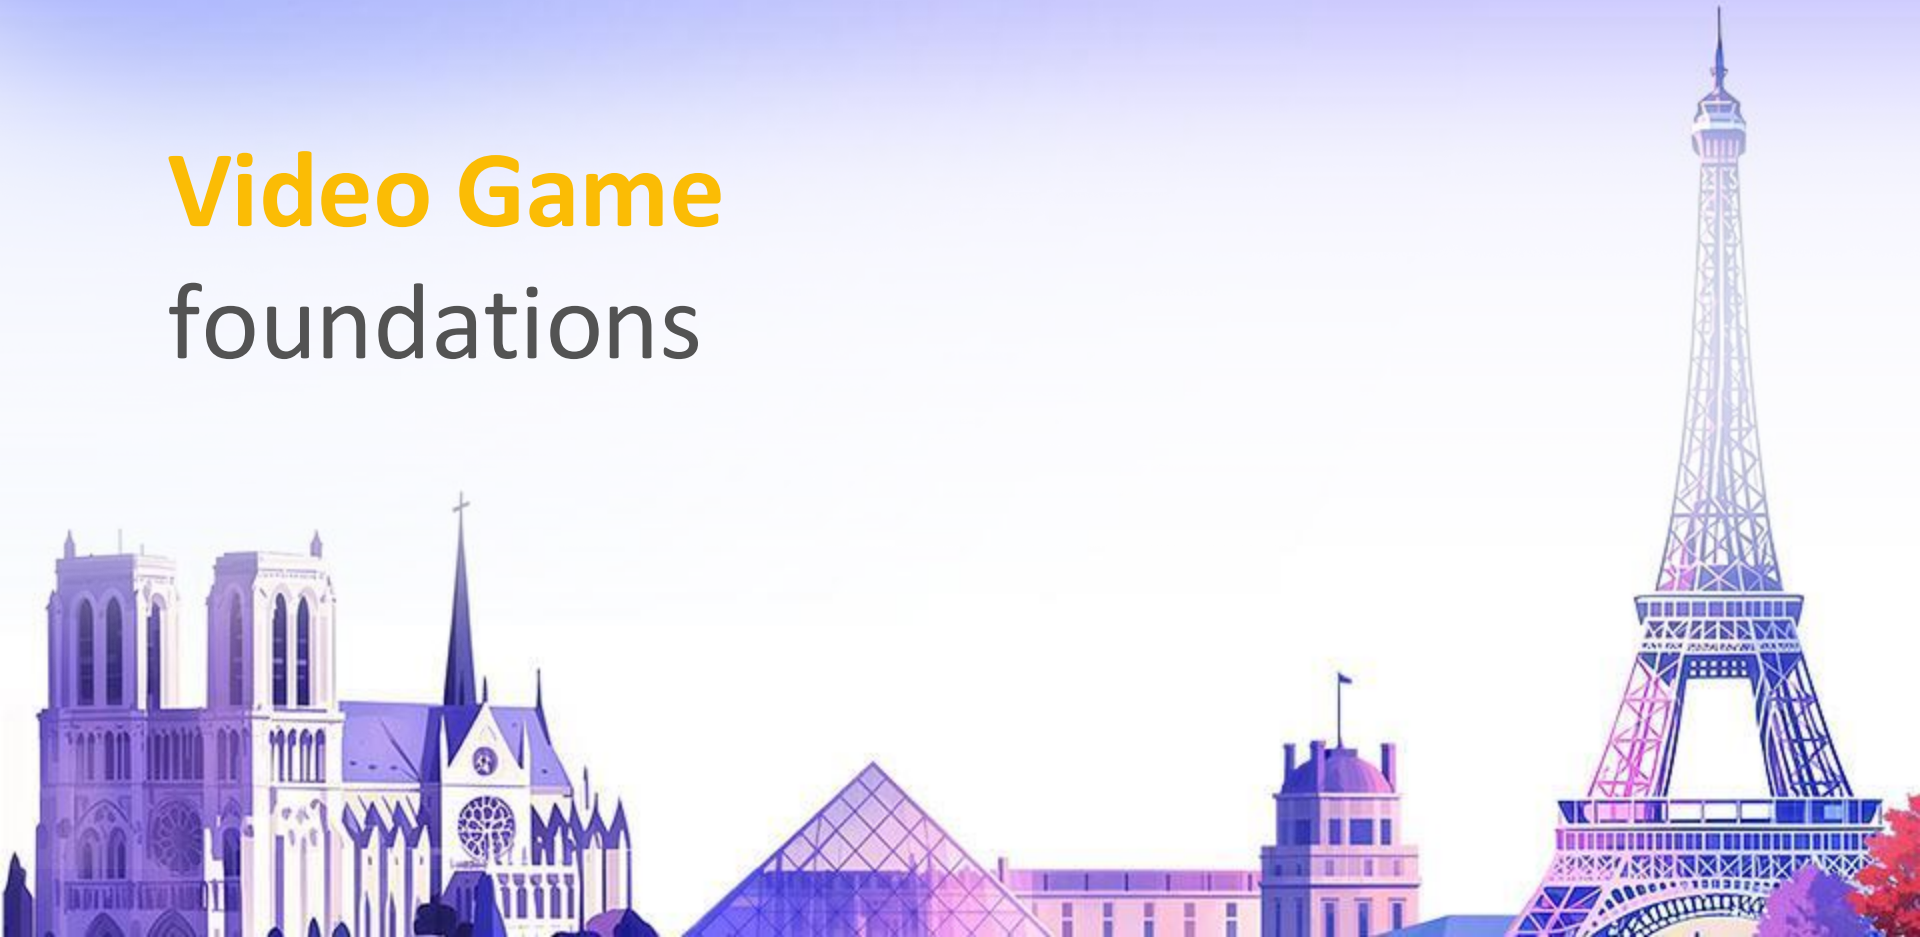

Mathematics in the core

Vertex and 3D computation

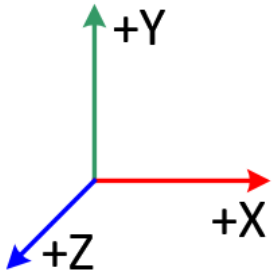

$$\begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix}$$

Vector

translation

$$\begin{pmatrix} 1 & 0 & 0 & t_x \\ 0 & 1 & 0 & t_y \\ 0 & 0 & 1 & t_z \\ 0 & 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} v_x \\ v_y \\ v_z \\ 1 \end{pmatrix} = \begin{pmatrix} v_x + t_x \\ v_y + t_y \\ v_z + t_z \\ 1 \end{pmatrix}$$

Scaling

$$\begin{pmatrix} s_x & 0 & 0 & 0 \\ 0 & s_y & 0 & 0 \\ 0 & 0 & s_z & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix} \begin{pmatrix} v_x \\ v_y \\ v_z \\ 1 \end{pmatrix} = \begin{pmatrix} s_x v_x \\ s_y v_y \\ s_z v_z \\ 1 \end{pmatrix}$$

Rotation

$$\mathbf{R}_X(\theta) = \begin{pmatrix} 1 & 0 & 0 & 0 \\ 0 & \cos(\theta) & \sin(\theta) & 0 \\ 0 & -\sin(\theta) & \cos(\theta) & 0 \\ 0 & 0 & 0 & 1 \end{pmatrix}$$

Mathematics in the core

Light computation

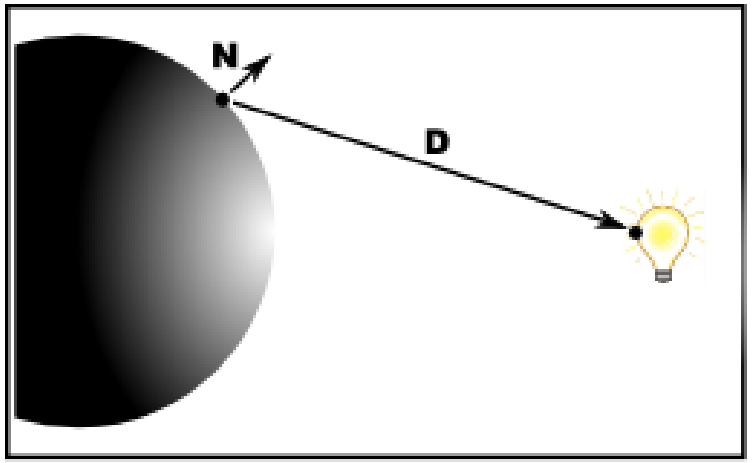

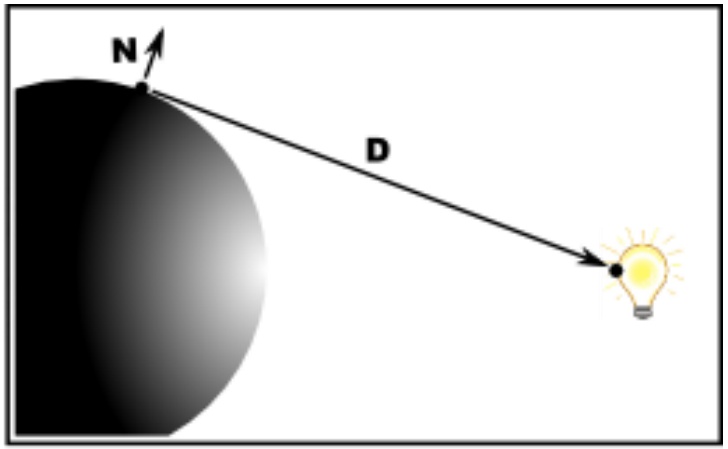

Scalar product to determine the light & reflexion surface alignment

GOSIM

GOSIM AI Paris 2025

Math recap and secret superpower

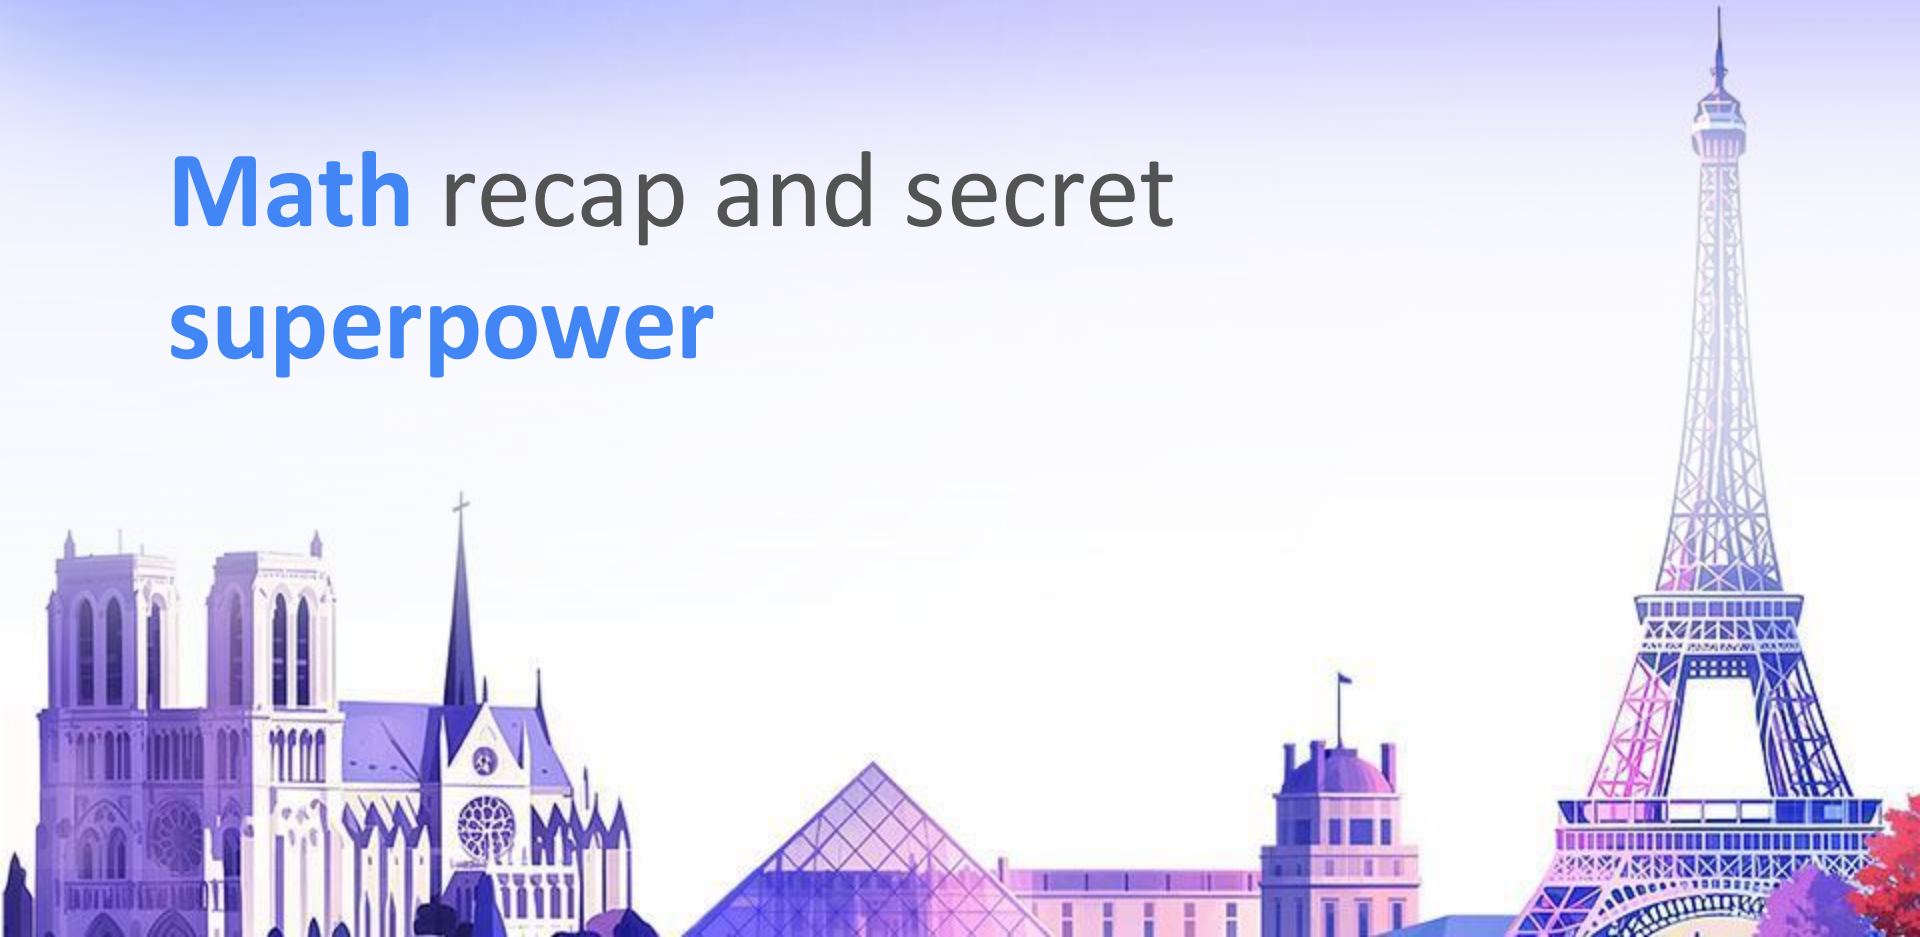

School reminder

Addition and multiplication in the core

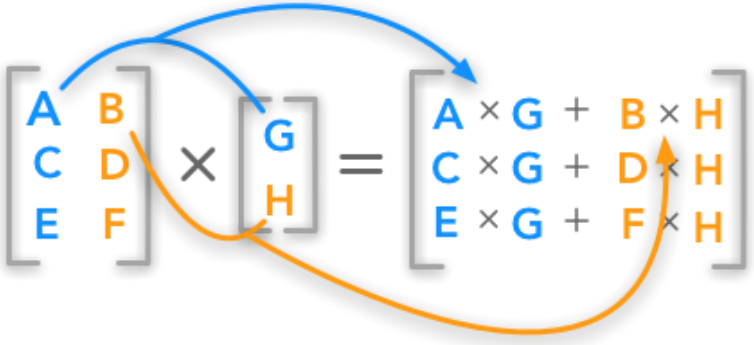

Matrix - vector multiplication

Scalar product

$$\begin{bmatrix} A_1 & A_2 & A_3 \end{bmatrix} \begin{bmatrix} B_1 \\ B_2 \\ B_3 \end{bmatrix} = A_1B_1 + A_2B_2 + A_3B_3 = \vec{A}.\vec{B}$$

GPUs revolution

Video game super power

GOSIM

NVIDIA®

AMD

LLM core computation

LLMs, matrices and vectors

Feel the similarities

"This is a prompt"

Prompt

Array of bytes

This
is
a
pro
mpt

Vector of tokens

0.5	1.2	0.35	8.9	0.54	1
0.4	0.01	1.6	5.1	1.3	0.57
1.8	0.4	5.6	10	4.4	0.24
5.1	0.25	0.7	0.92	0.3	0.02
0.25	0.68	0.08	1.25	5.91	0.99

Vector of token array values

Matrix

13545
5645
3515
12
1354
18
1561531
1812
15644

Embeddings

Neuron activation

Scalar product and activation function

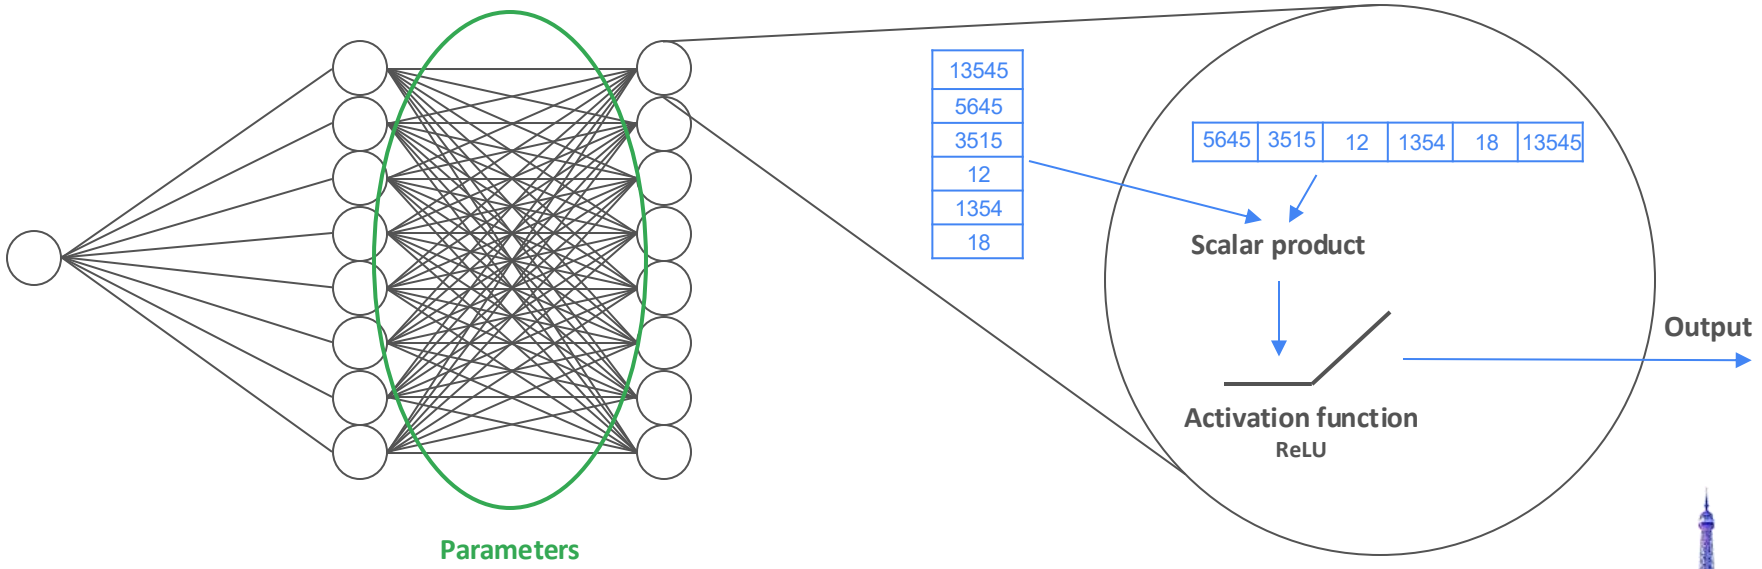

Not so different

GPUs in the core

	Video Games	LLMs
Matrix - vector multiplication →	Vertices transformation	Tokenization and embeddings
Scalar product →	Light effect	Neuron alignement & activation

Cloud run & GPUs

GOSIM

Since September 2024

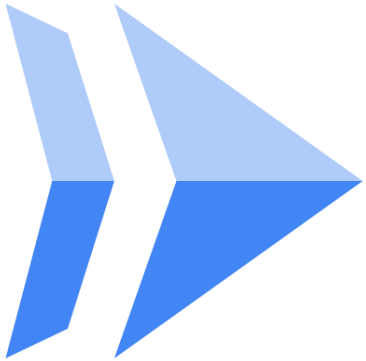

NVIDIA®

The Andes trail


GEMMA
LLAMA
MISTRAL
OH!

Ollama swiss knife

Serving LLM, easily

GOSIM

Multi LLM support

Gemma

LLaMA by Meta

MISTRAL AI_

Adaptive runtime

NVIDIA

Open Source

Secure

Weather **mix**

Running LLM, serverless

Mix all the ingredients

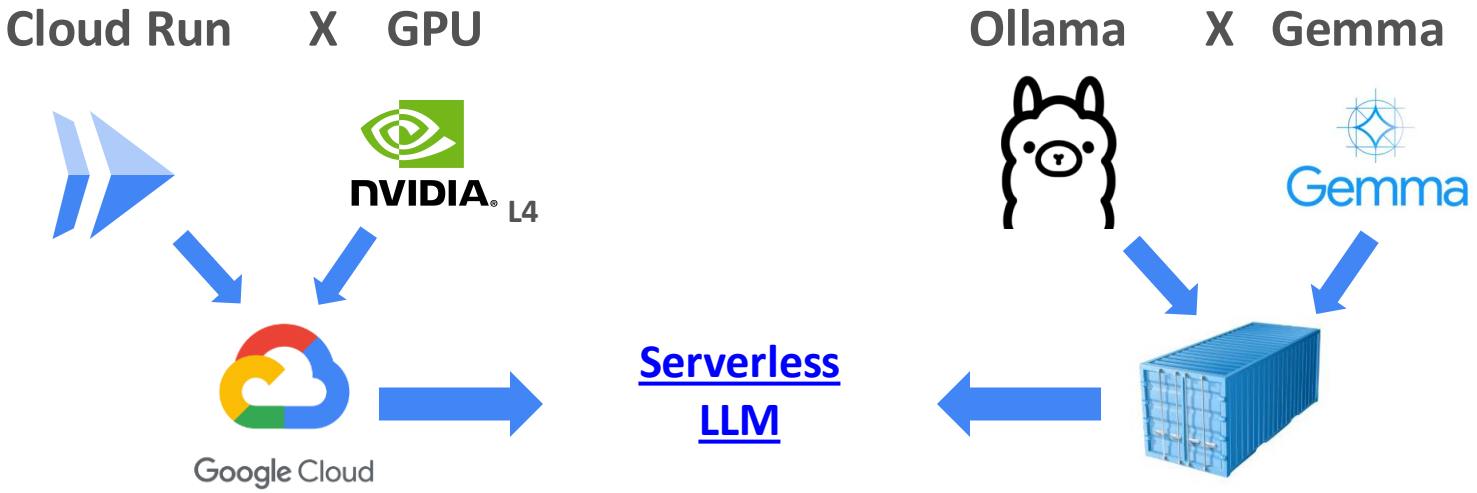

No solution's perfect

Pros and Cons

No overprovisioning

scale to 0

Scale with the traffic

pay as you use

Easy to use

Auto driver installation

Cold start

First request latency

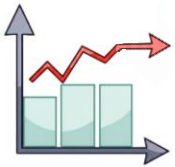

Limitation

Regions & max instances

GPUs available

only NVidia L4

Give Cloud Run GPU a **try**,

and move to **serverless LLM**

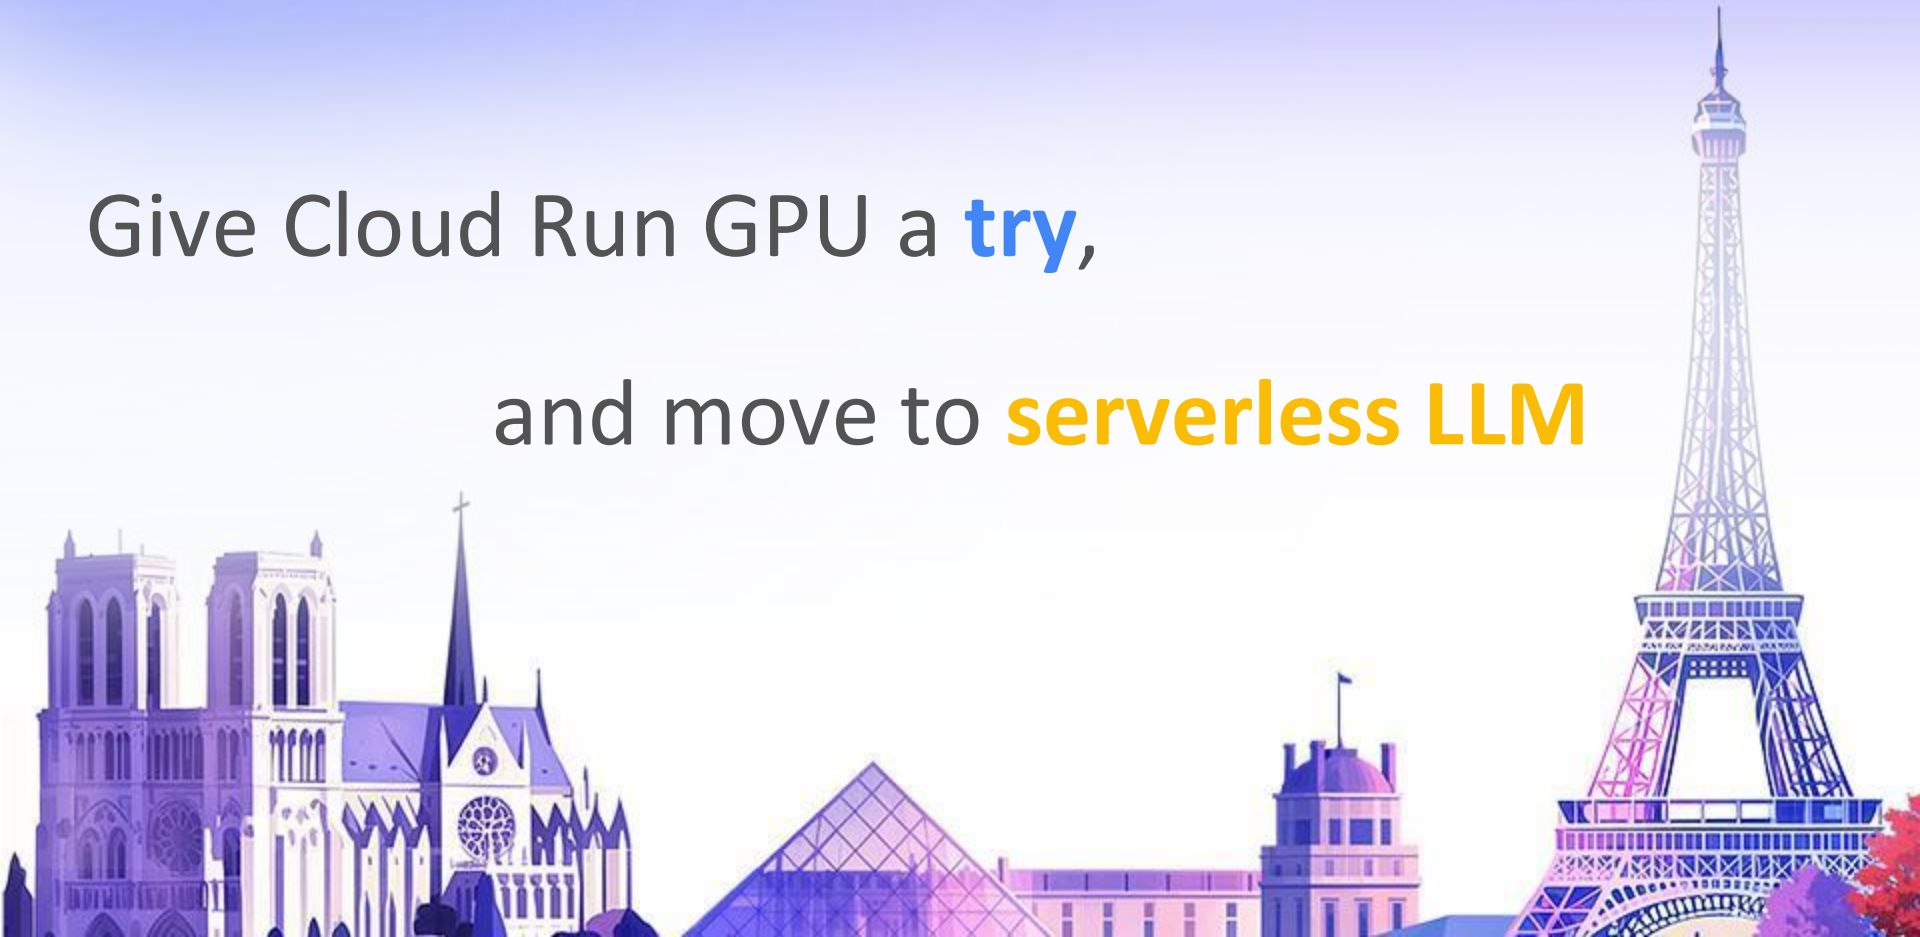

THANK YOU

Article https://medium.com/google-cloud/cloud-run-gpu-make-your-llms-serverless-5188caacc667

Find me on :

Twitter @gblaquiere

Medium @guillaume-blaquiere

GitHub guillaumeblaquiere

LinkedIn guillaume blaquiere