

RAGFlow—Leading the Open-Source Revolution in Enterprise-Grade RAG

Yingfeng@InfiniFlow.ai



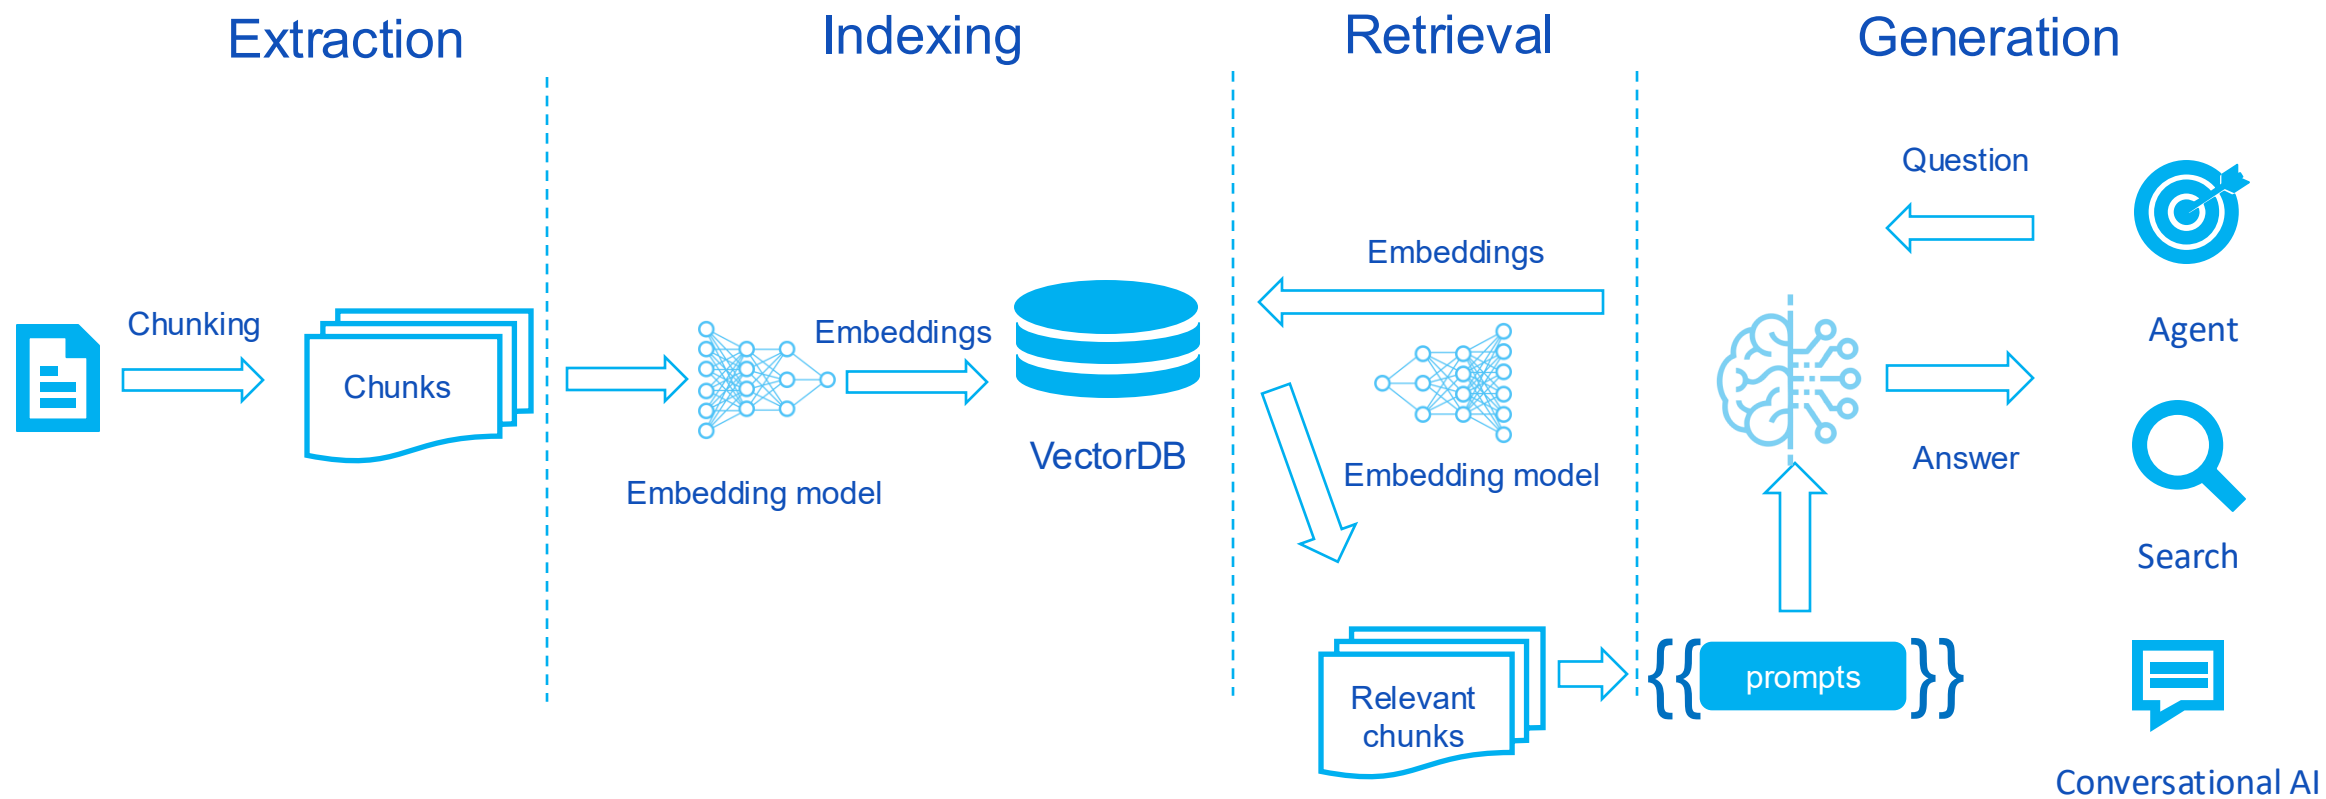
- ① Introduction
- ② Chunking
- ③ Retrieval
- ④ Agentic and future
- ⑤ Summary



Introduction



Paradigm of RAG



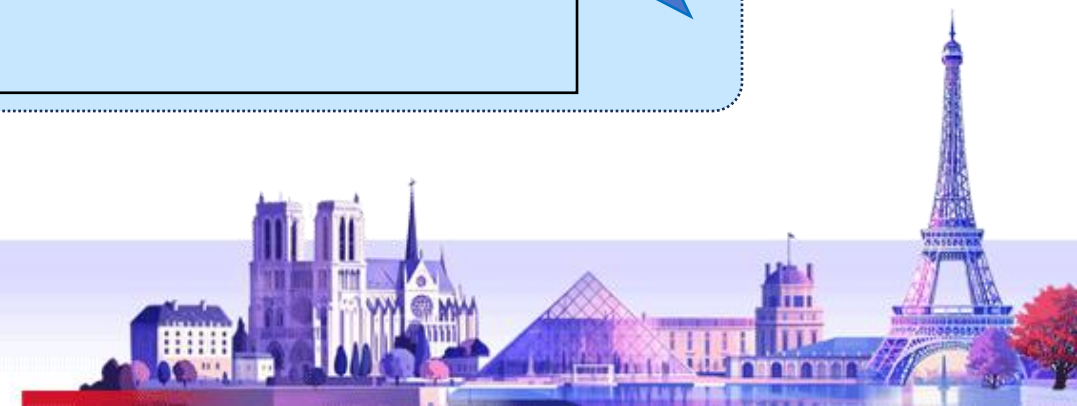
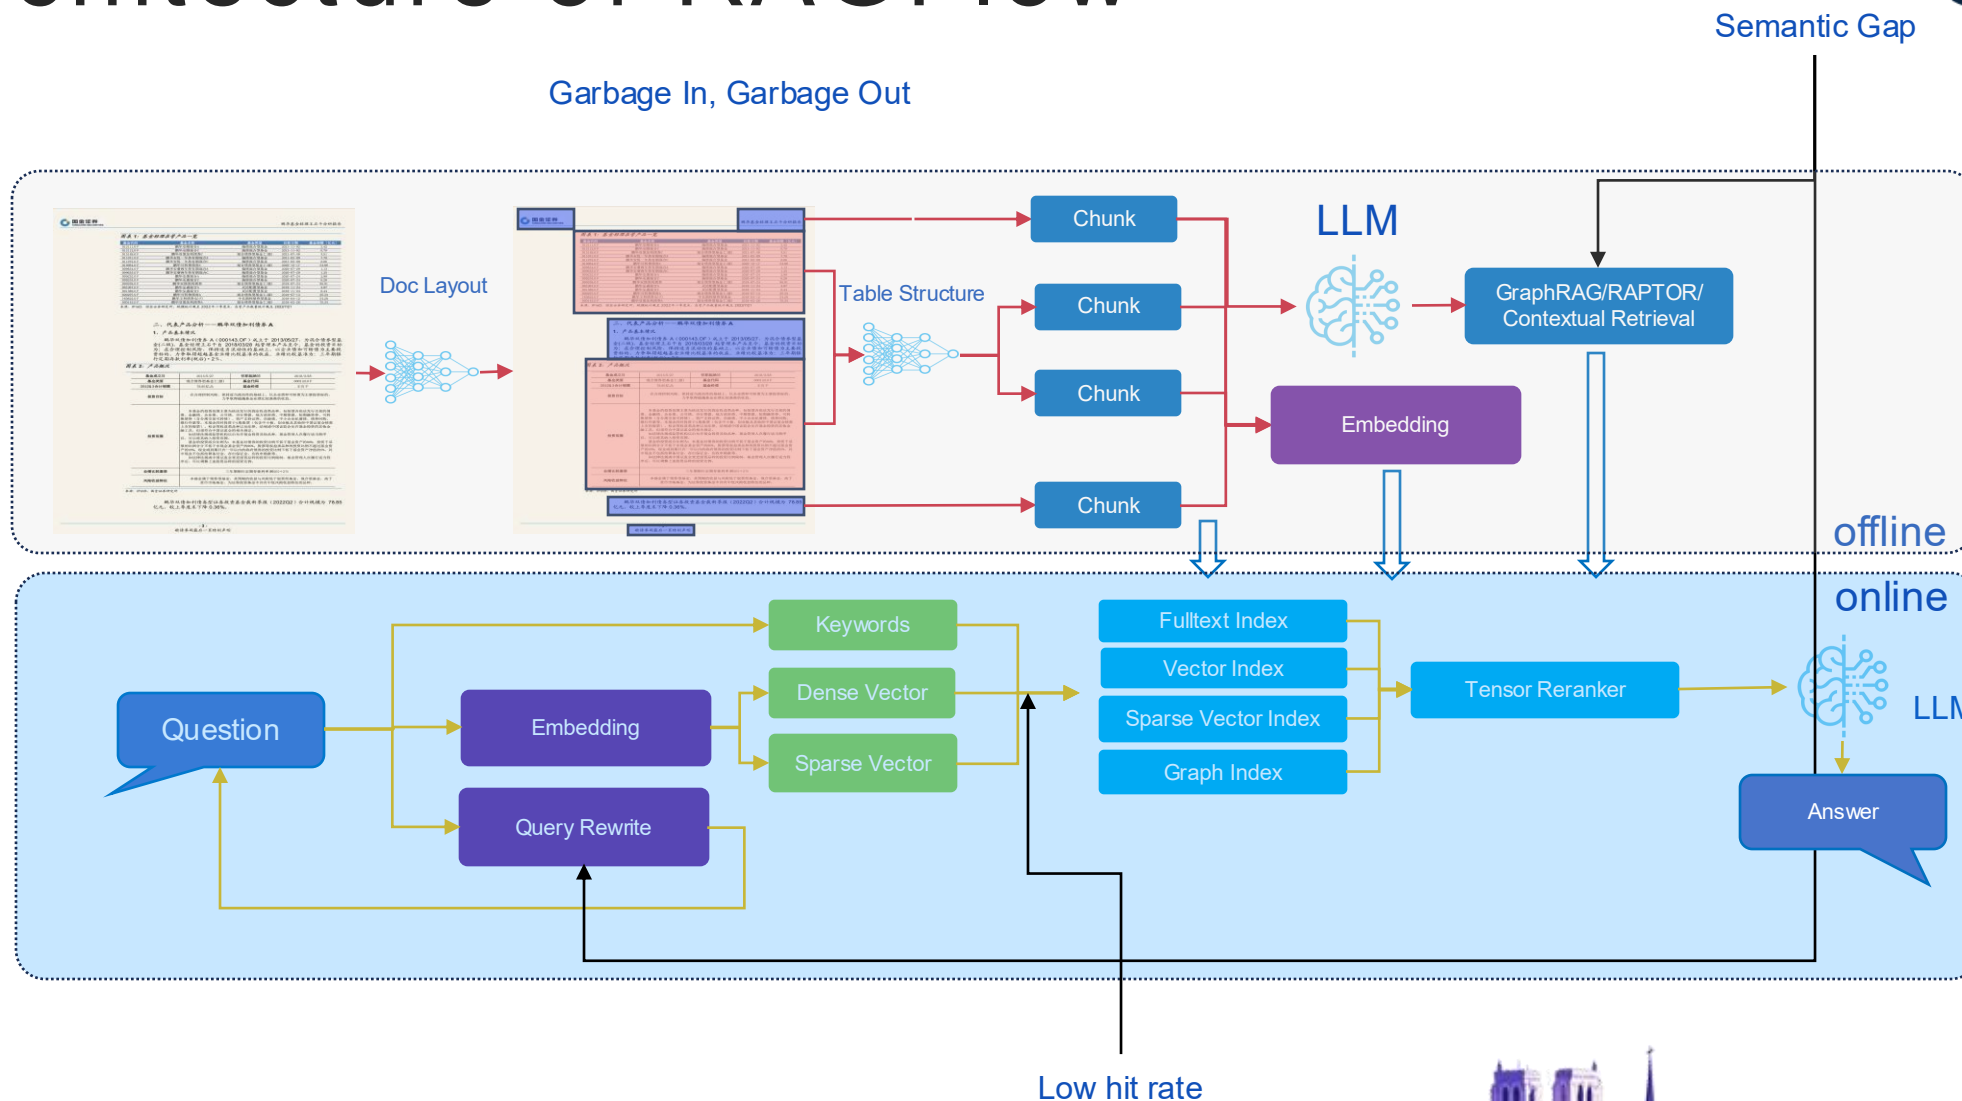
Challenges of RAG

- ✗ Garbage in, garbage out
- ✗ Low hit rate
- ✗ Semantic gap between question and answer

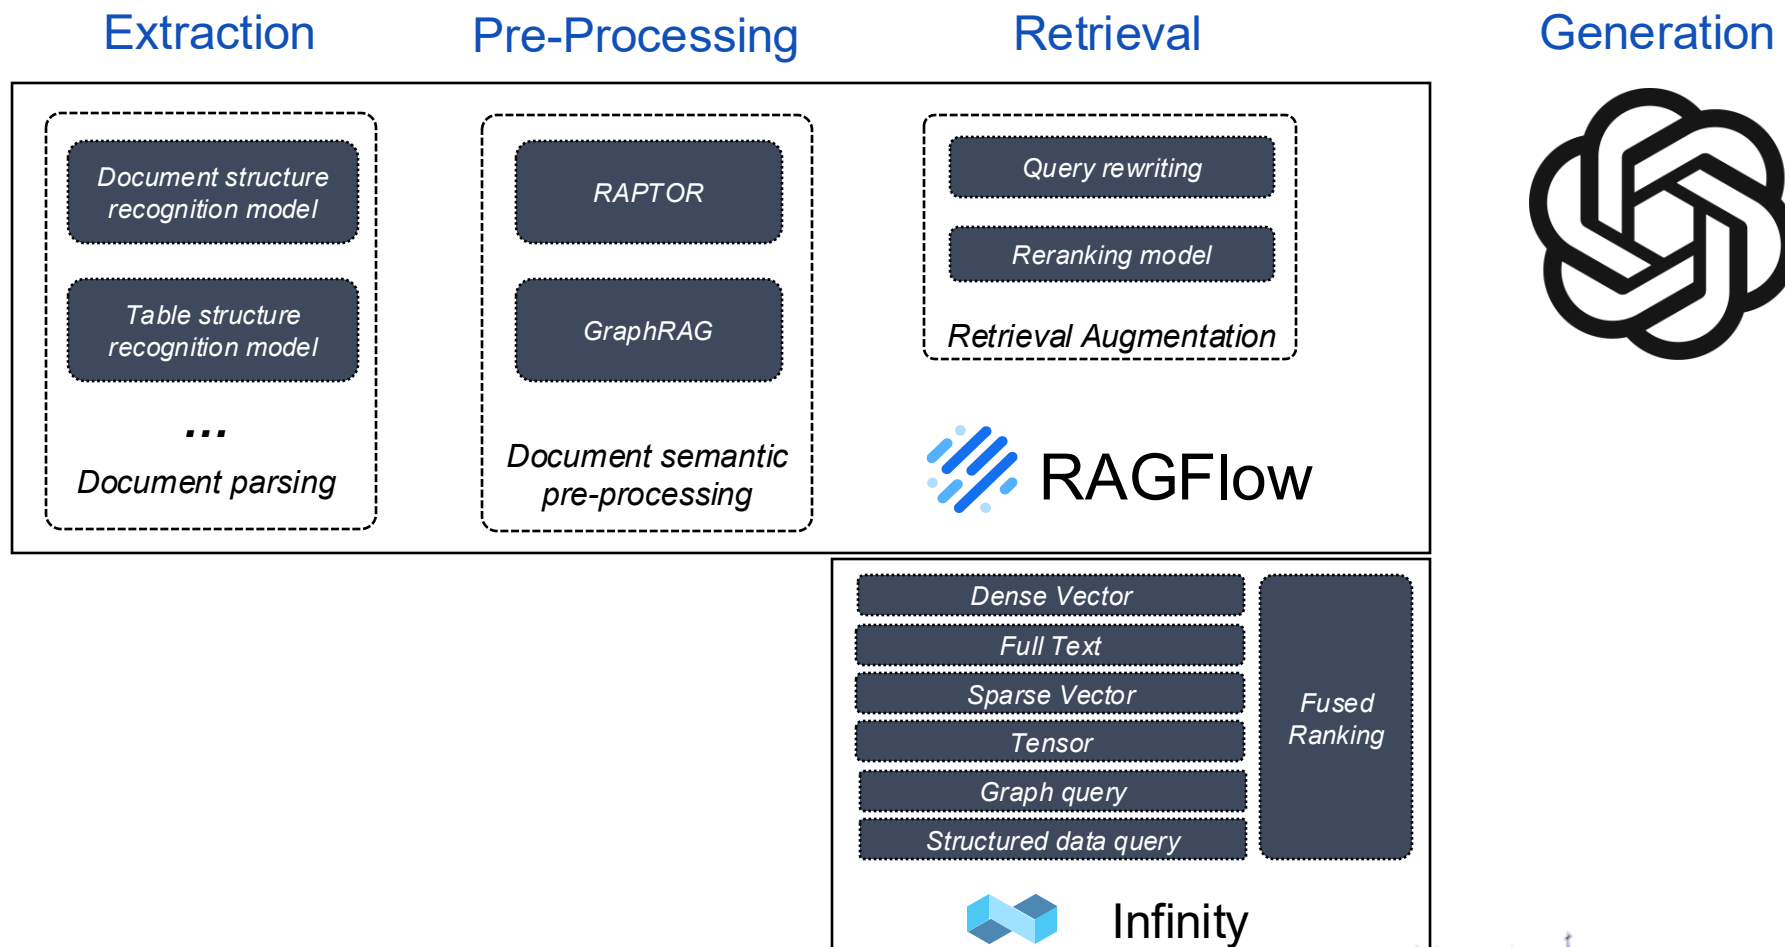


Architecture of RAGFlow

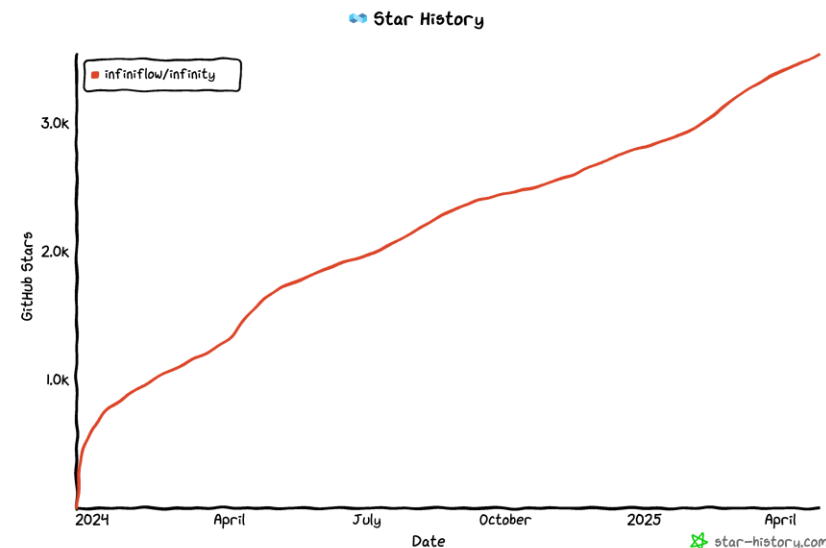
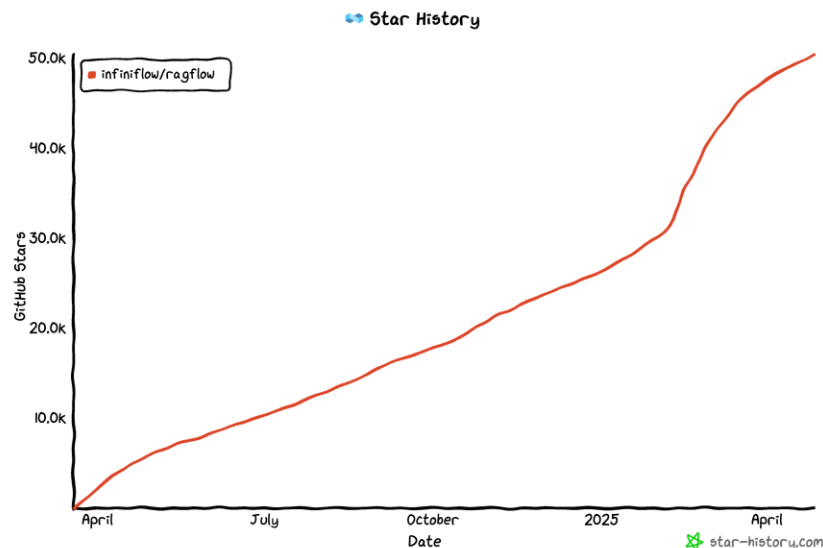
GOSIM



RAG ≠ LLMOps



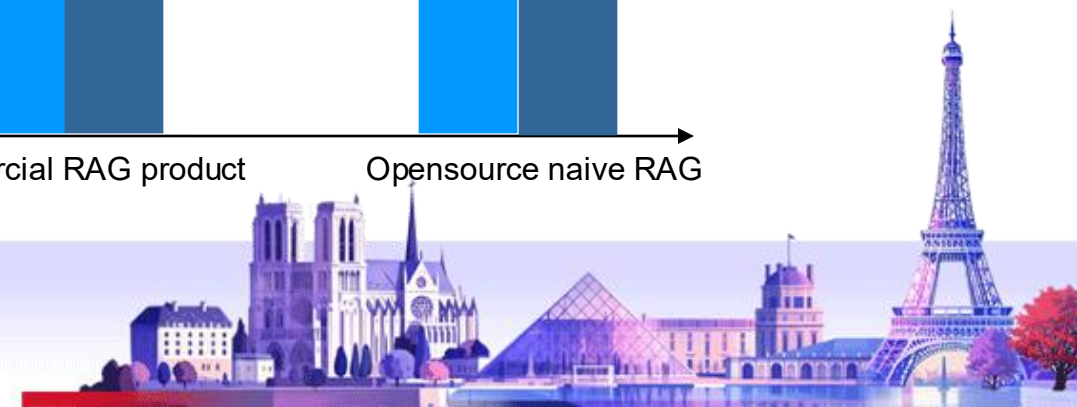
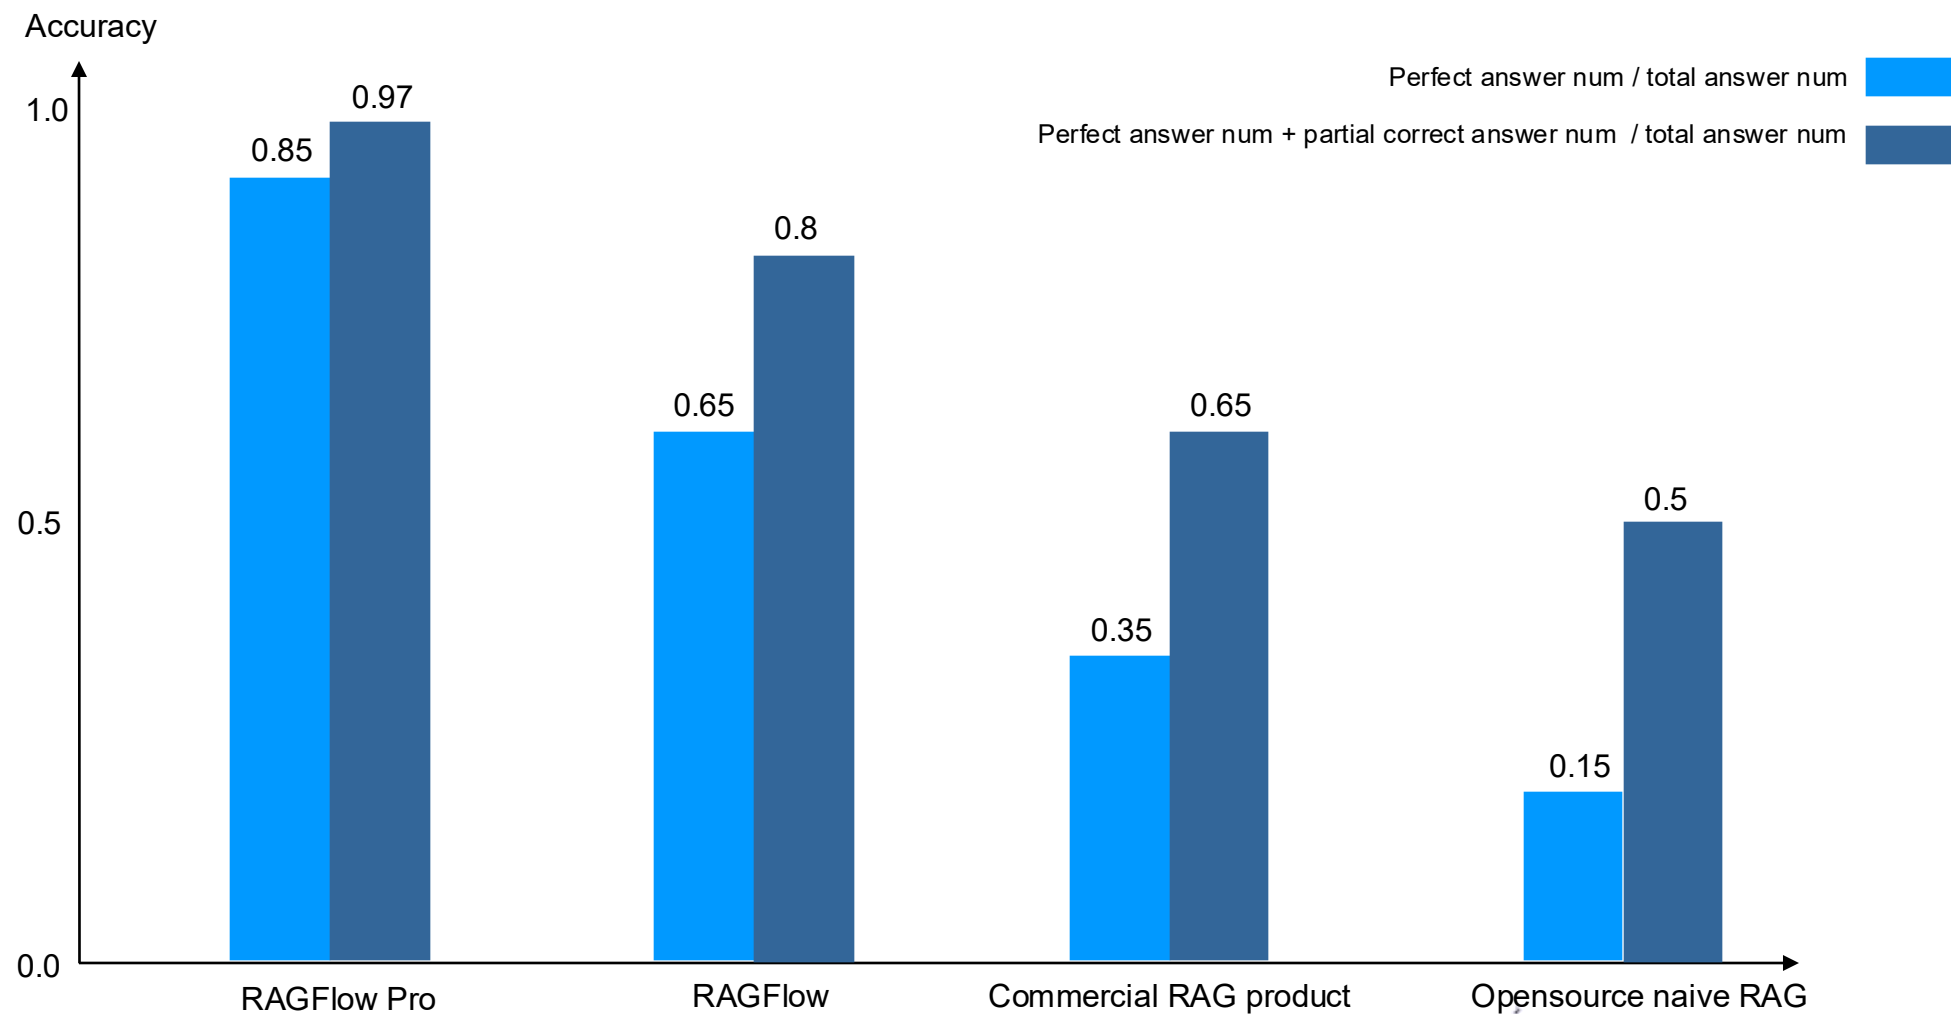
InfiniFlow = Infinity + RAGFlow



Top Star Geo-Locations



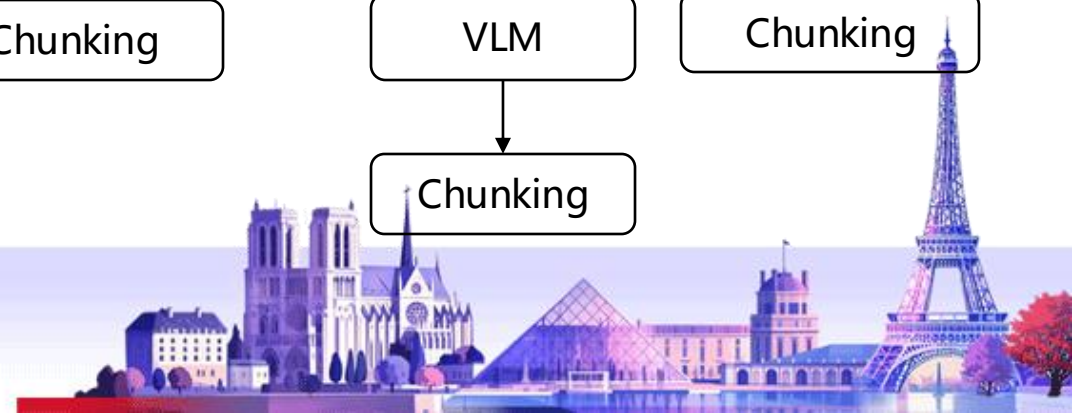
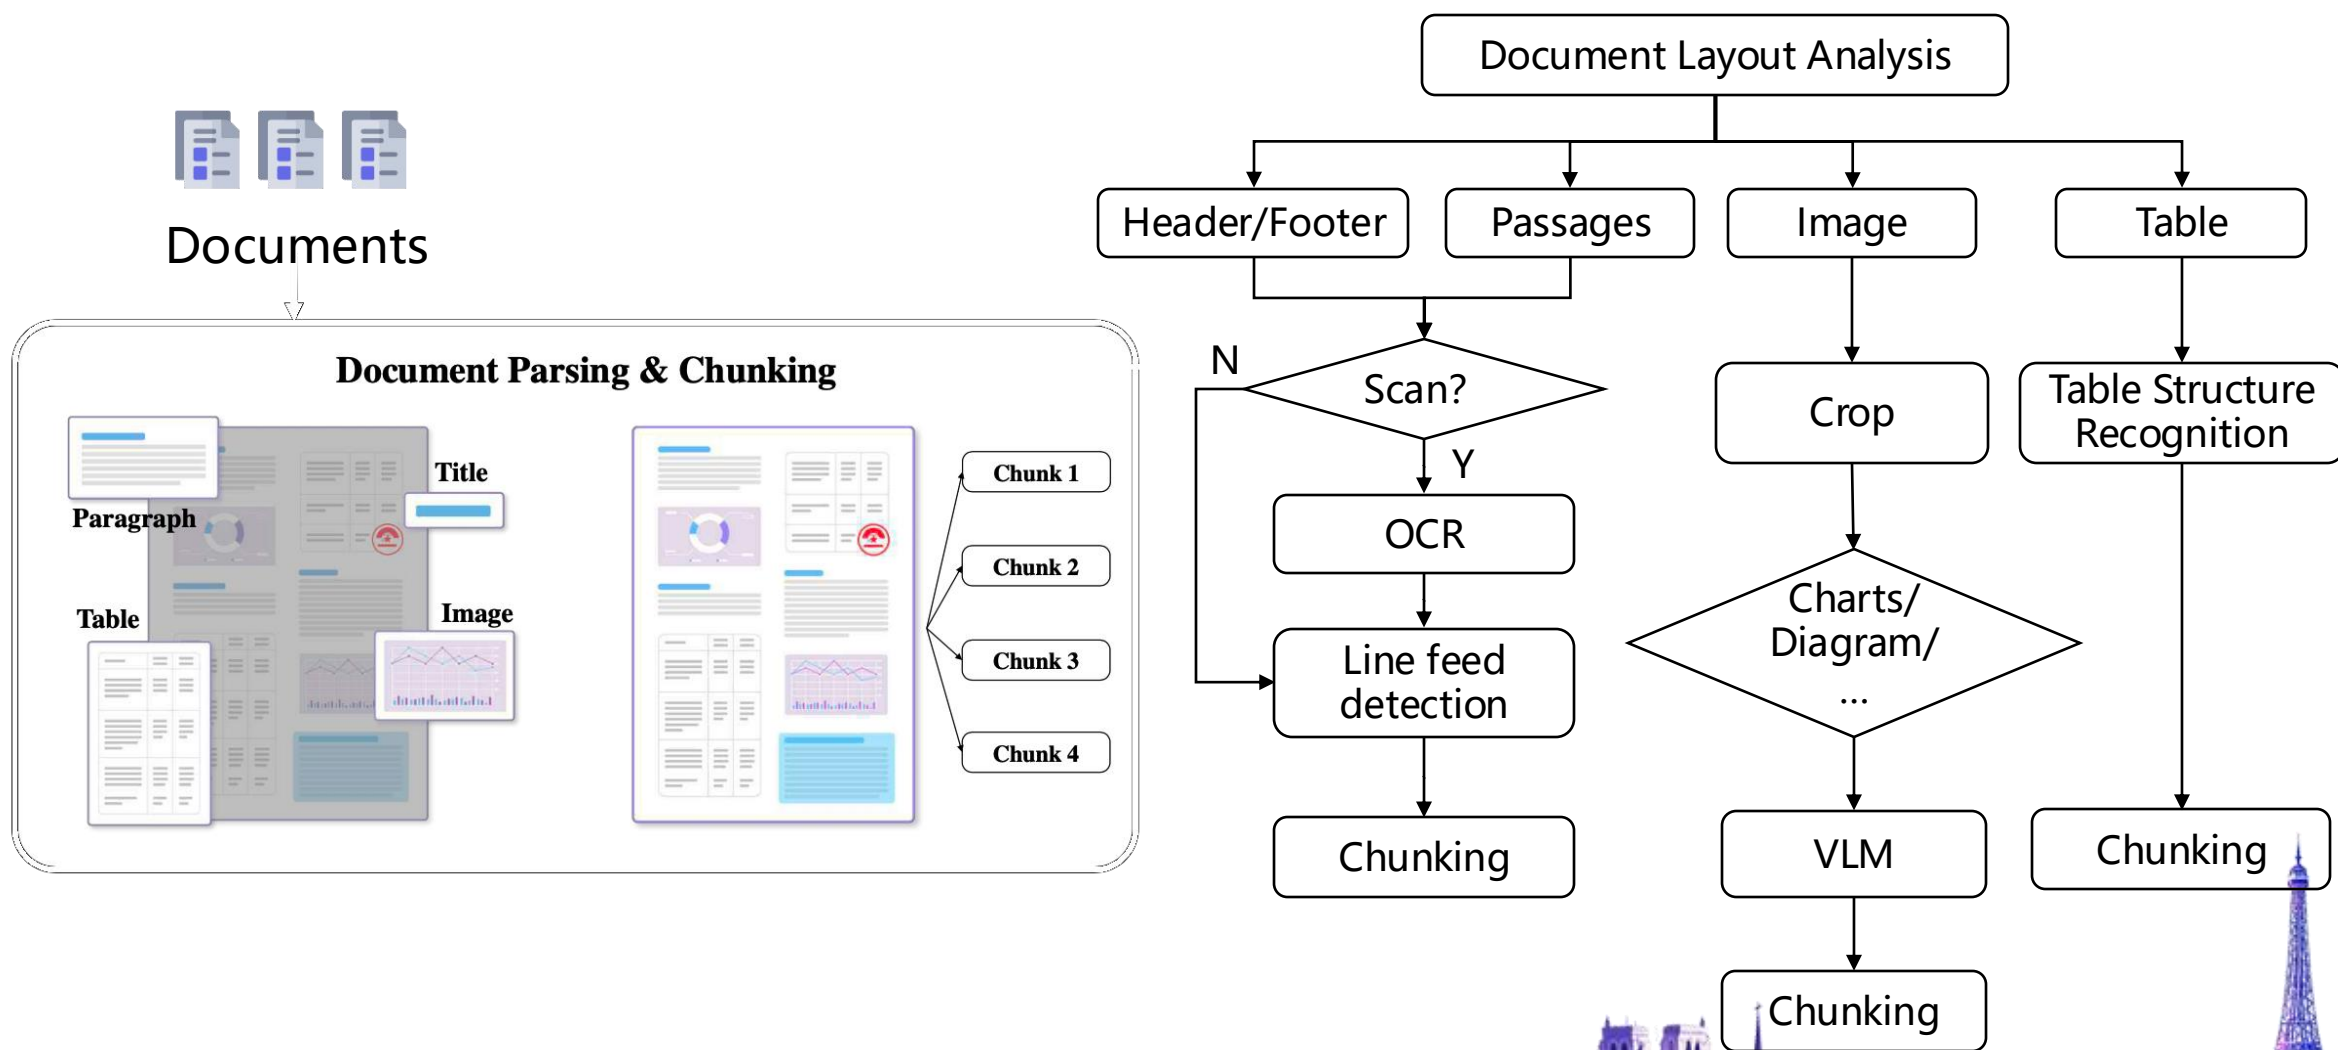
RAGFlow vs Others



Chunking



Deep Document Understanding



Visual Chunking Results

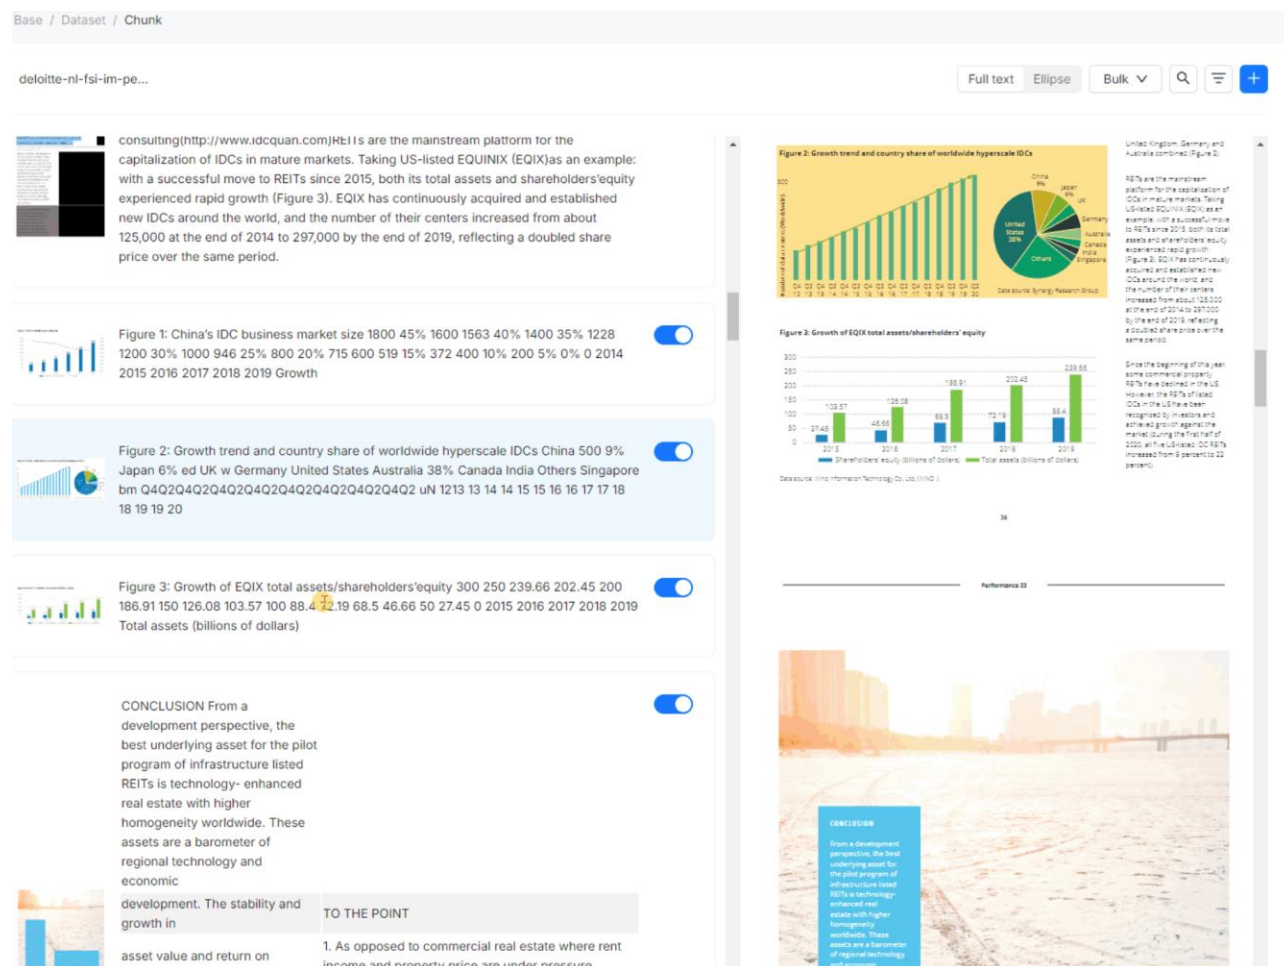


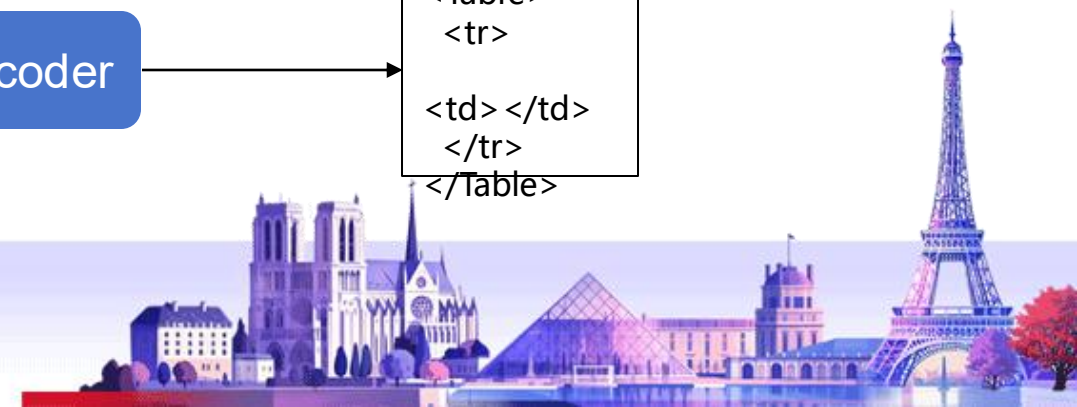
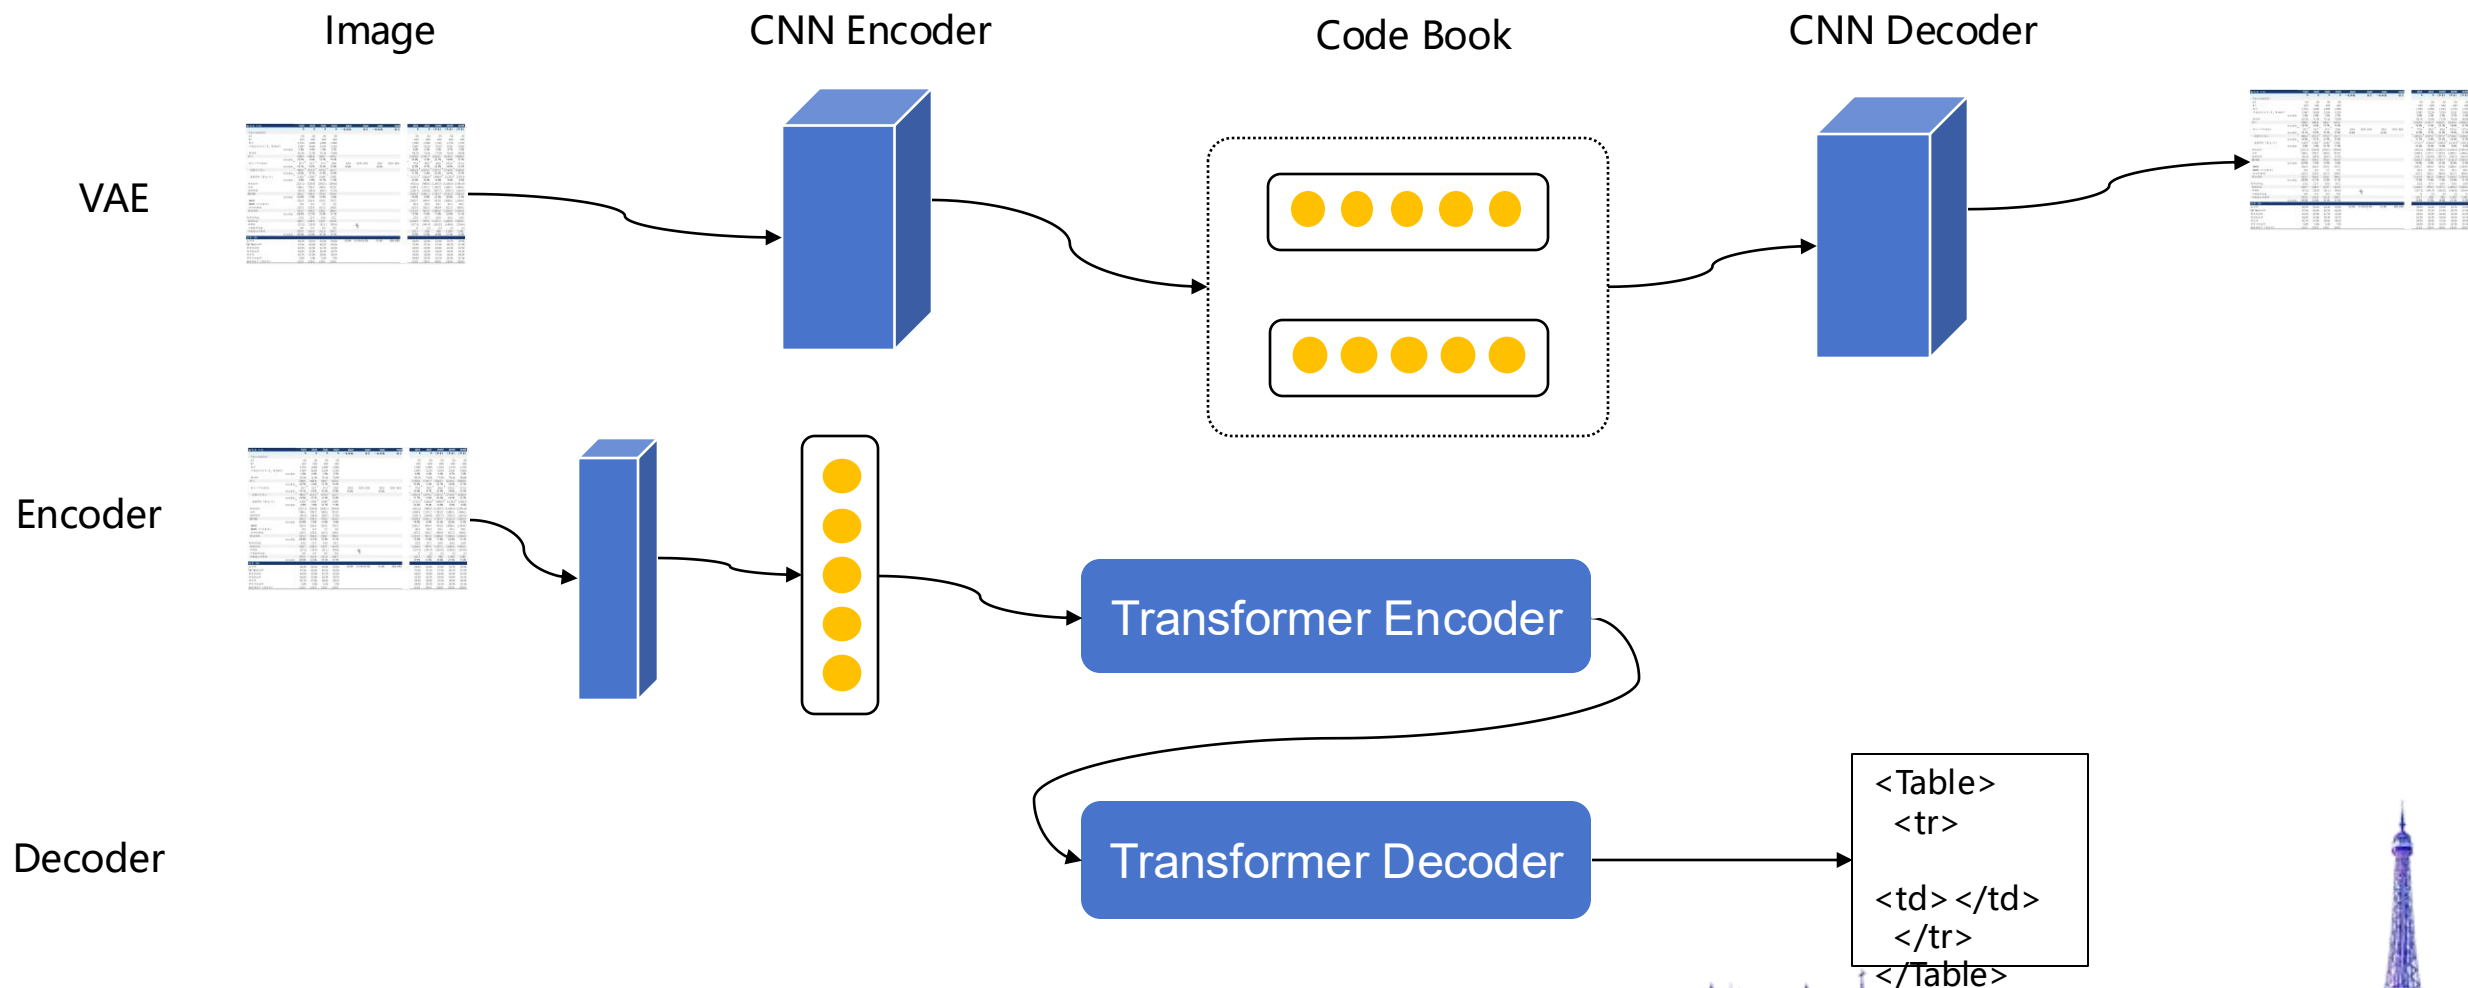
Table Structure Recognition

Model	CAD							MoCA-Mask-TE							ARG
	$S_m \uparrow$	$F_\beta^\omega \uparrow$	MAE \downarrow	$F_\beta \uparrow$	$E_m \uparrow$	mDice \uparrow	mIoU \uparrow	$S_m \uparrow$	$F_\beta^\omega \uparrow$	MAE \downarrow	$F_\beta \uparrow$	$E_m \uparrow$	mDice \uparrow	mIoU \uparrow	
Ours ²³	0.757	0.593	0.020	0.631	0.865	0.599	0.510	0.734	0.476	0.010	0.497	0.736	0.497	0.422	0.00%
w/ difference image	0.724	0.528	0.023	0.569	0.761	0.526	0.441	0.693	0.401	0.017	0.429	0.704	0.425	0.360	↓15.46%
w/ optical flow	0.729	0.534	0.022	0.573	0.765	0.532	0.450	0.697	0.409	0.016	0.436	0.708	0.433	0.370	↓13.41%
w/o intra-frame self-attention	0.741	0.552	0.023	0.588	0.862	0.574	0.490	0.683	0.394	0.013	0.412	0.684	0.413	0.352	↓10.79%
w/o cross-frame cues diffusion	0.736	0.546	0.024	0.584	0.860	0.569	0.481	0.679	0.384	0.013	0.404	0.681	0.405	0.340	↓12.17%
w/o temporal shifting	0.742	0.571	0.021	0.610	0.828	0.576	0.488	0.690	0.408	0.013	0.426	0.692	0.427	0.366	↓8.88%

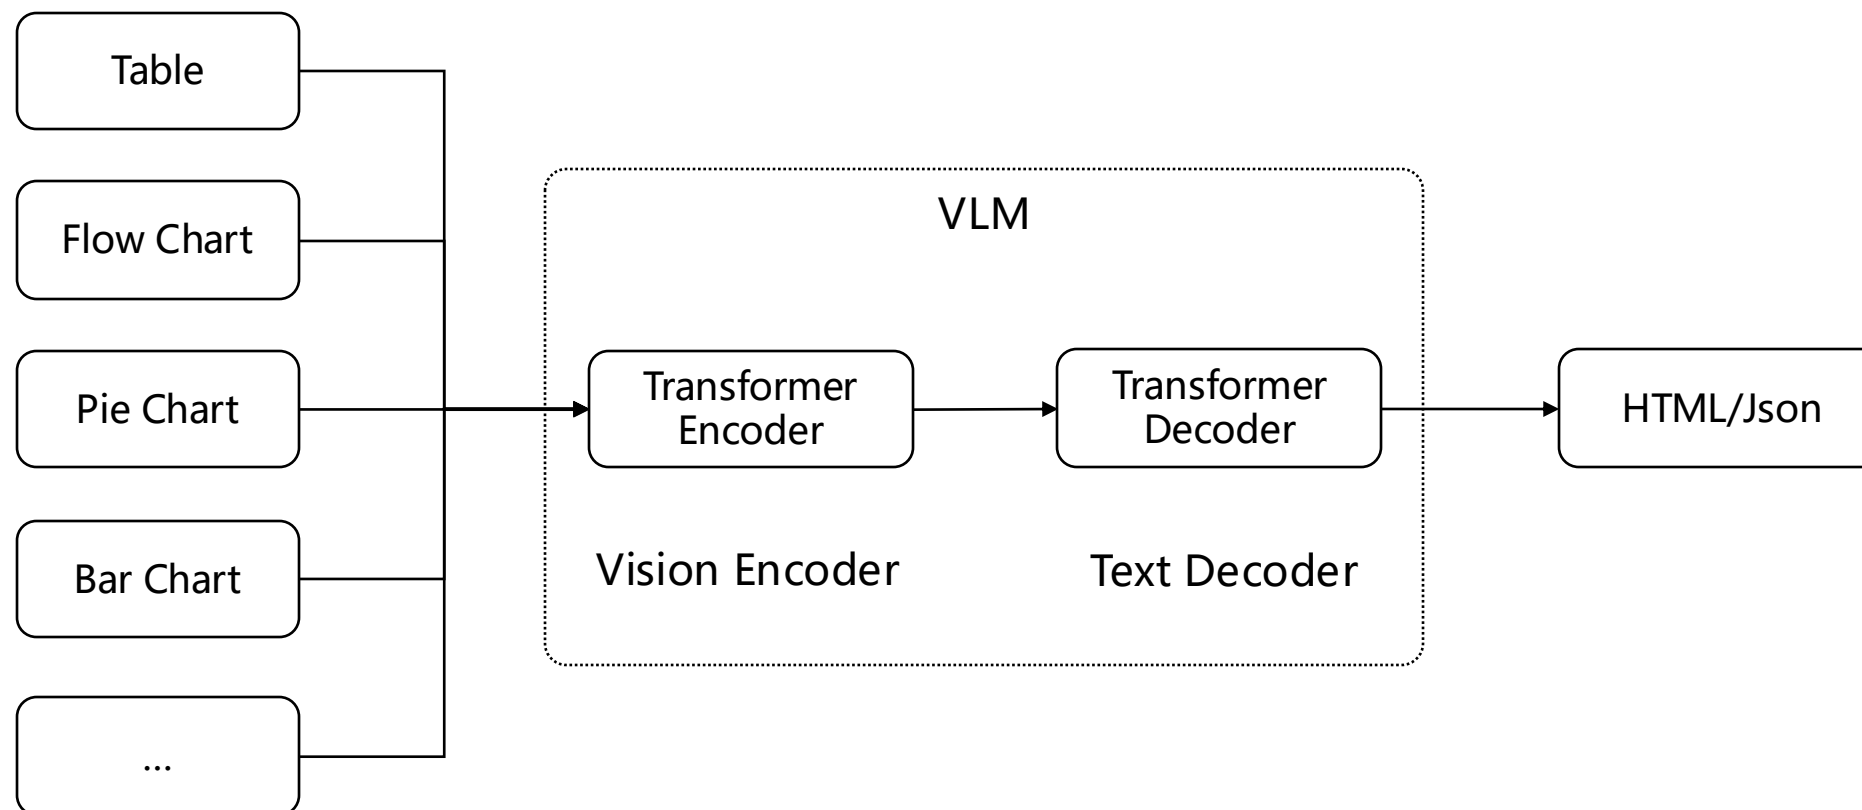
- Multi table header
- Border/Bordless
- Cell merge
- Cross page



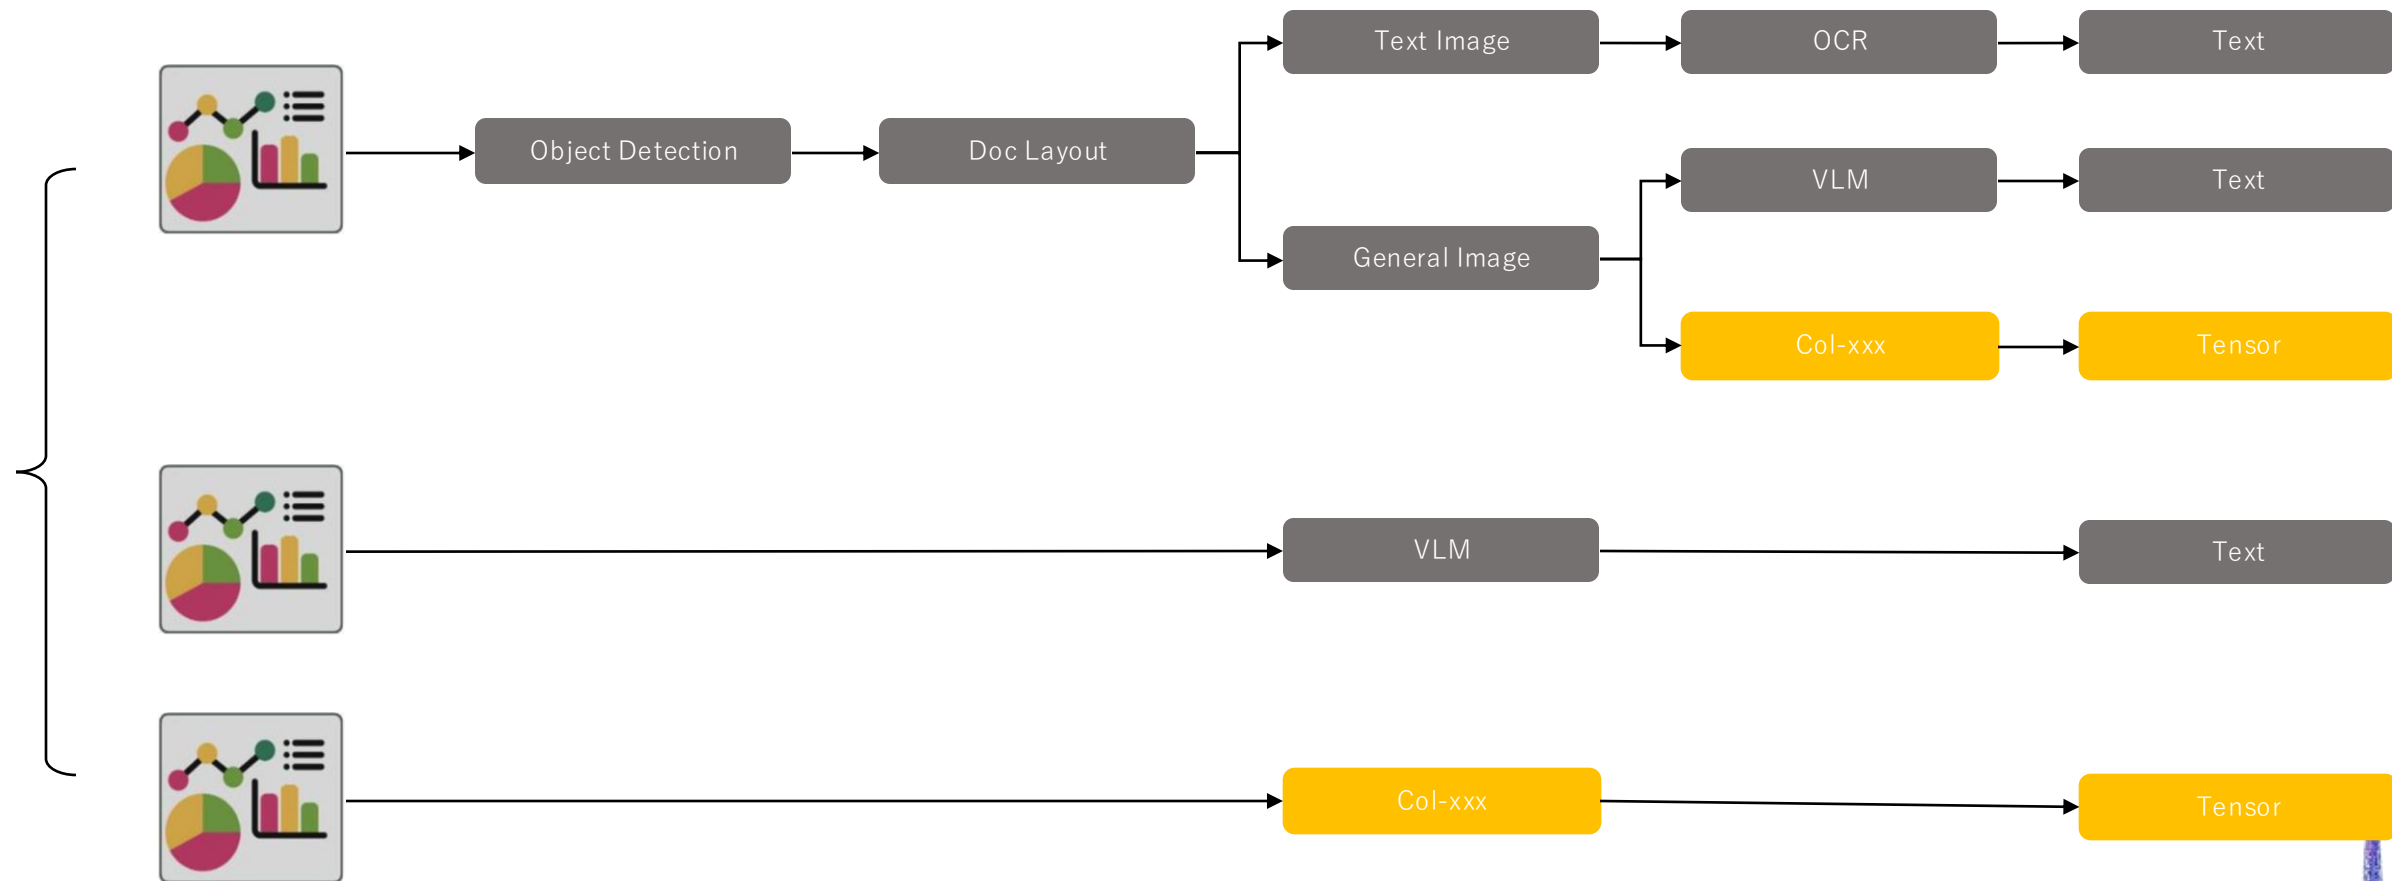
DeepDoc—CNN or GenAI?



DeepDoc—CNN or GenAI?



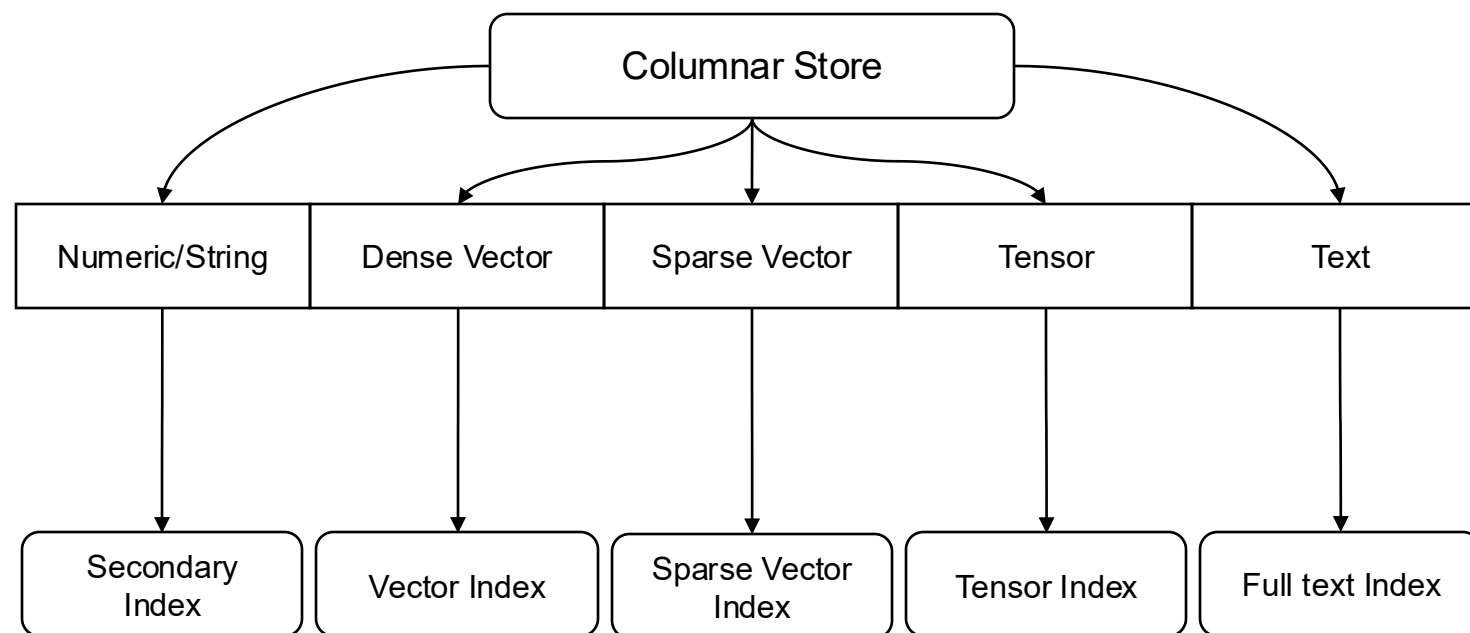
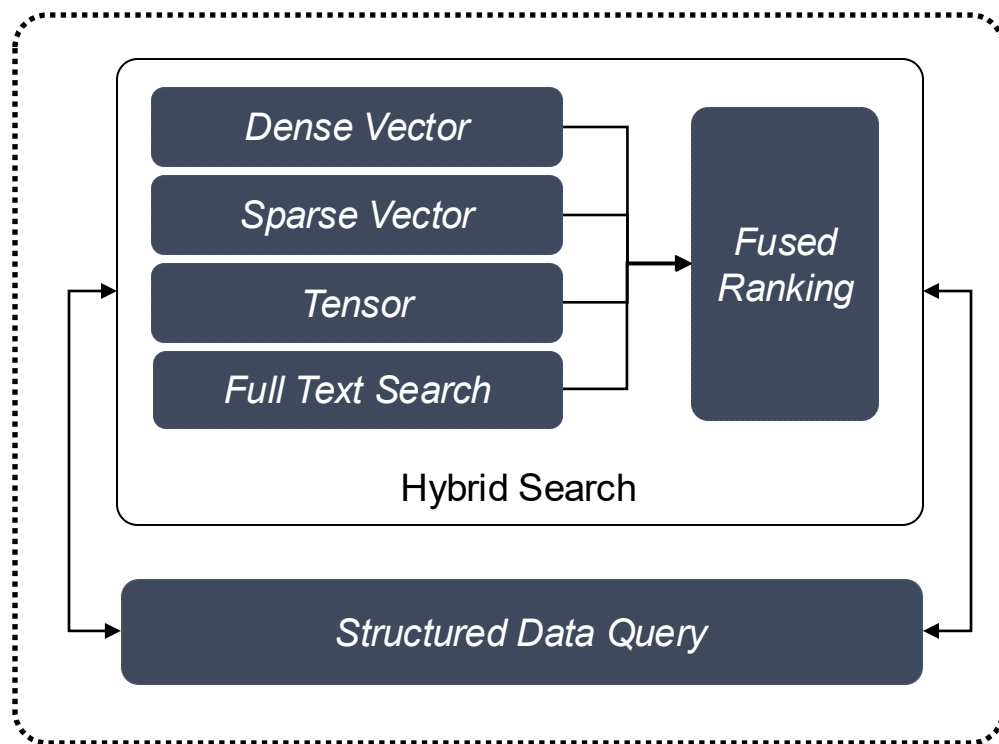
Multi-Modal RAG



Retrieval



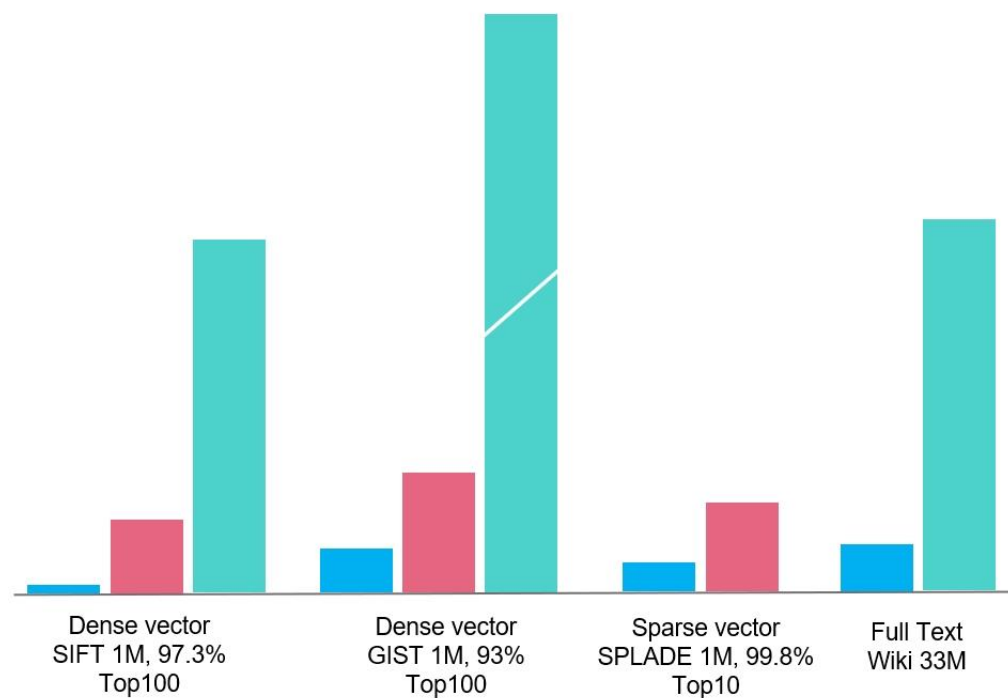
Infinity—Indexing Database



Benchmark

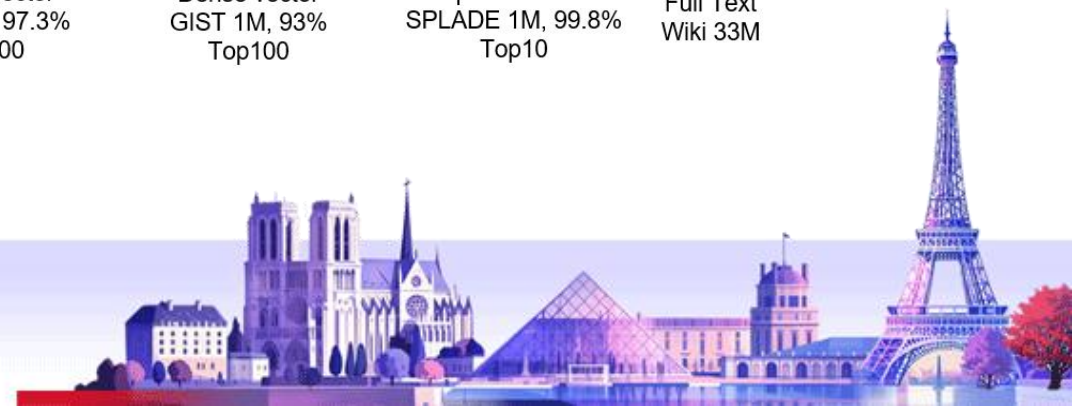
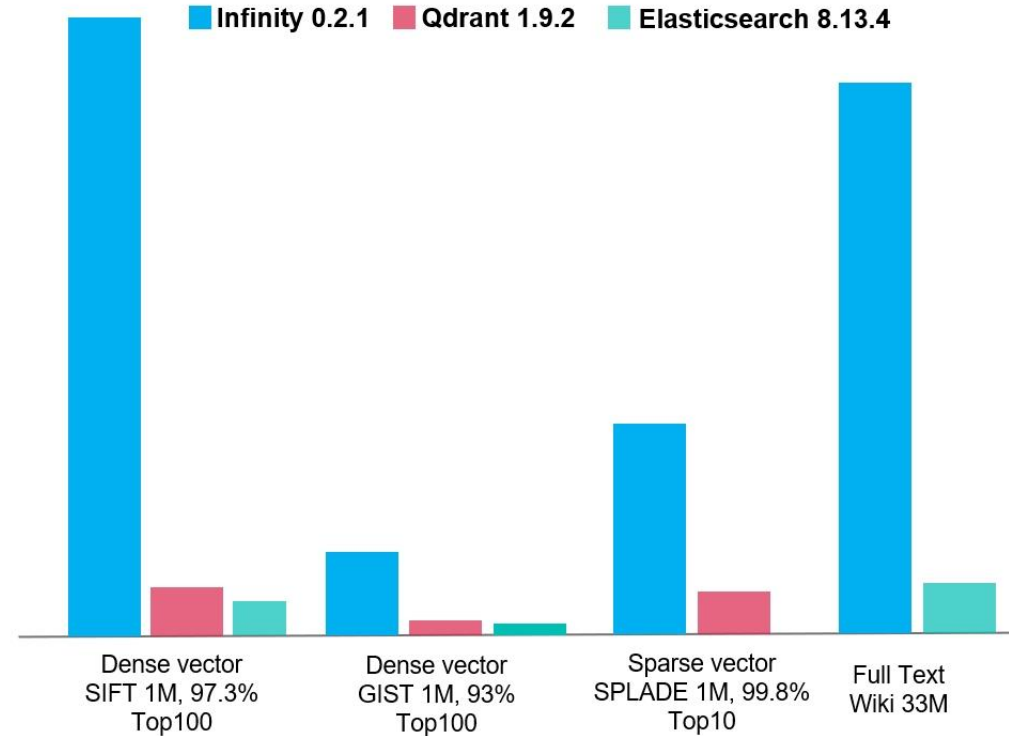
Query latency benchmark: Lower is better

■ Infinity 0.2.1 ■ Qdrant 1.9.2 ■ Elasticsearch 8.13.4

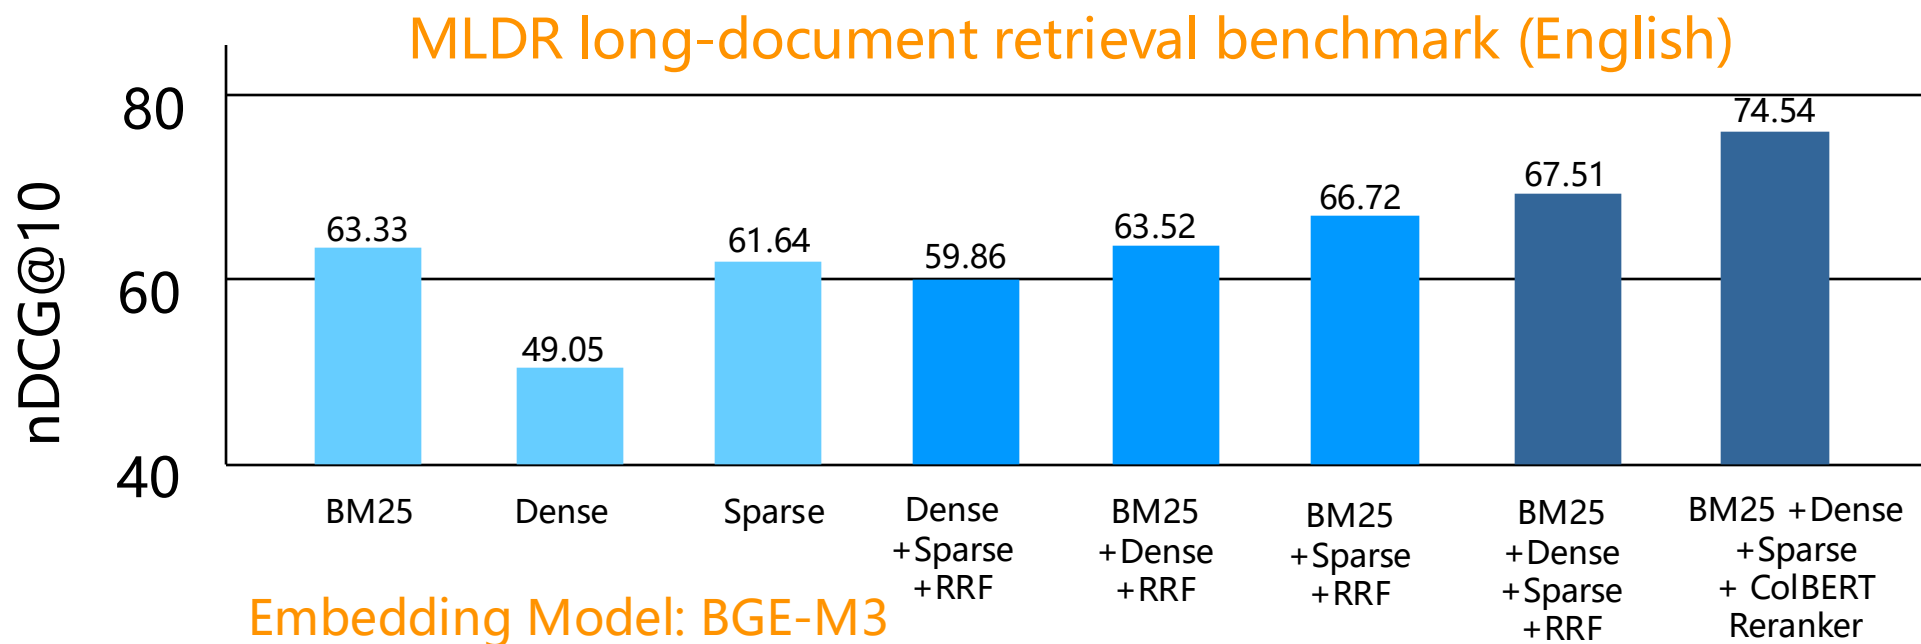


Query QPS benchmark: Higher is better

■ Infinity 0.2.1 ■ Qdrant 1.9.2 ■ Elasticsearch 8.13.4



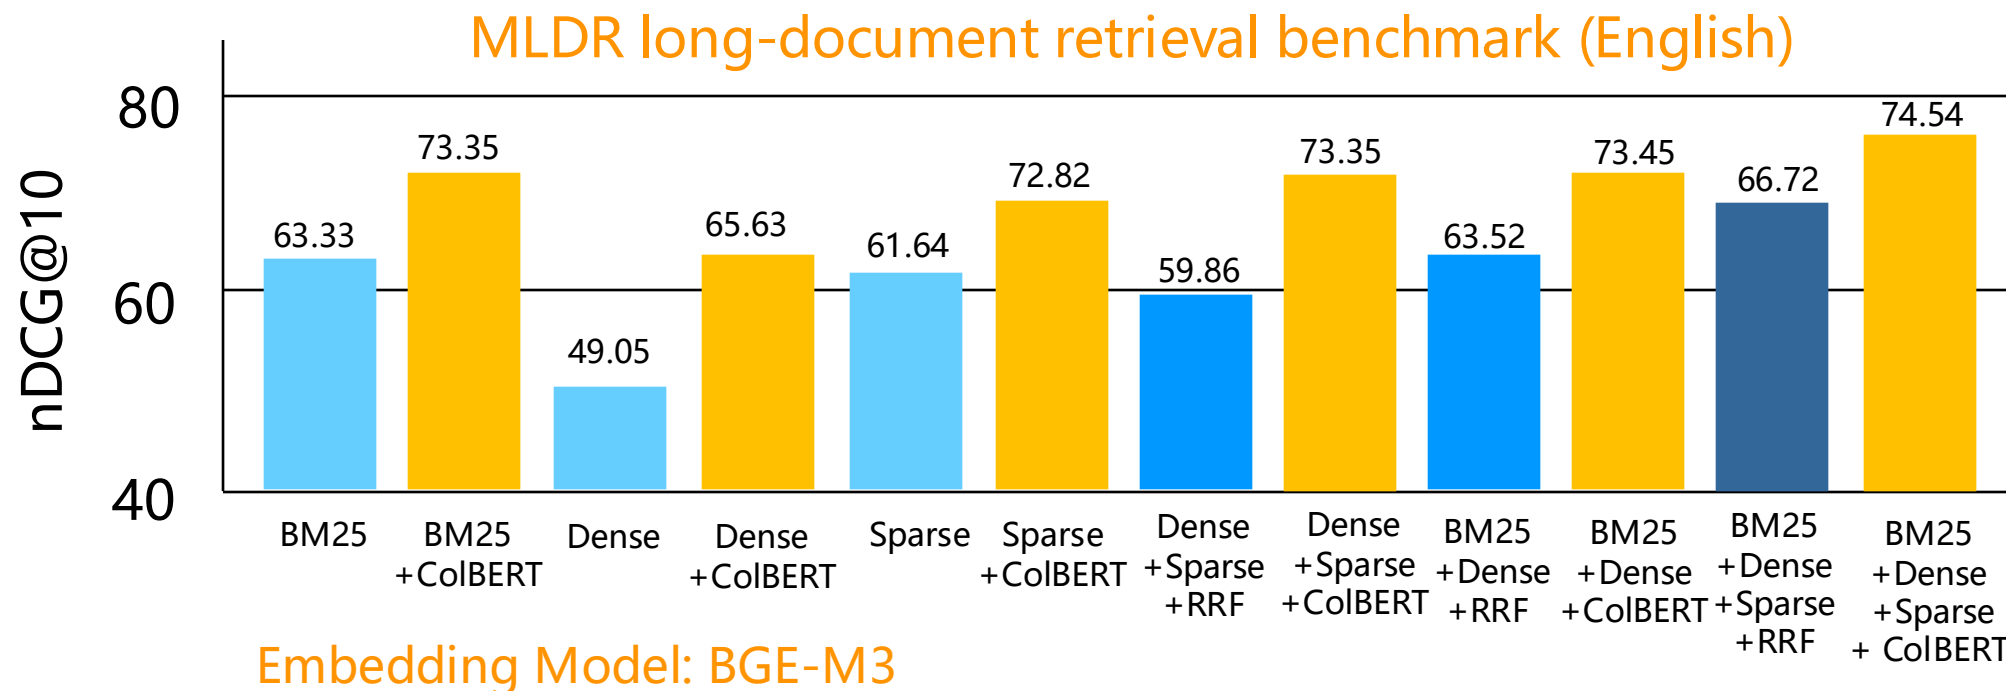
Hybrid Search



IBM Blended RAG
<https://arxiv.org/abs/2404.07220>

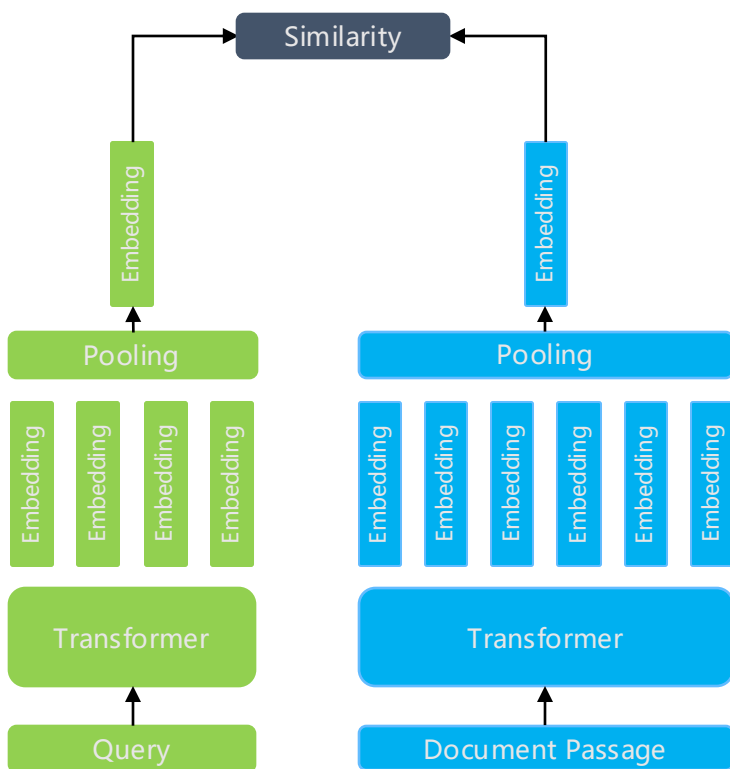


Benefits of ColBERT

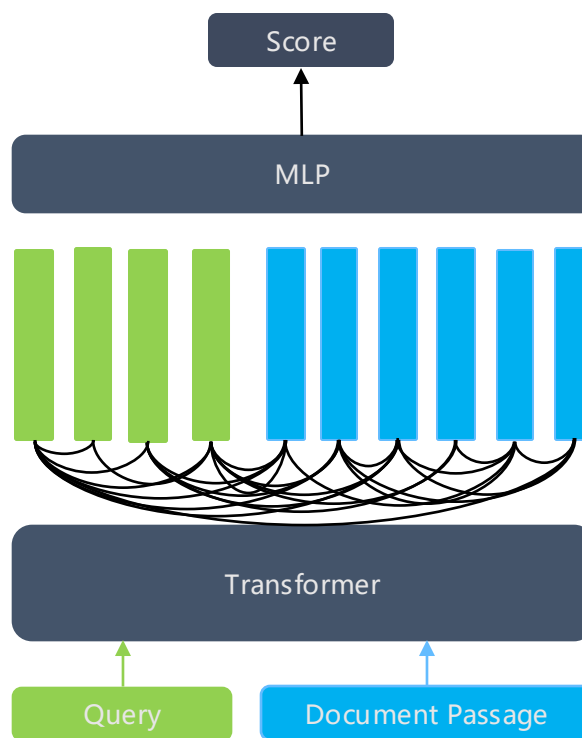


Ranking Model

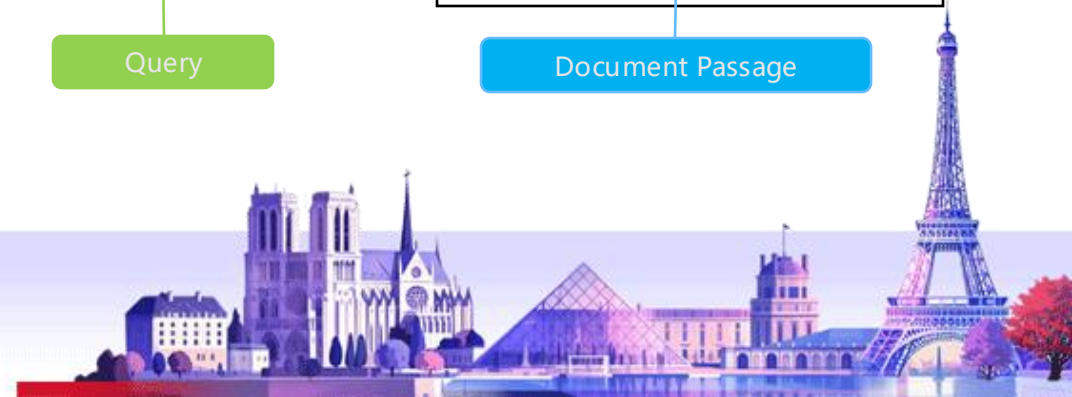
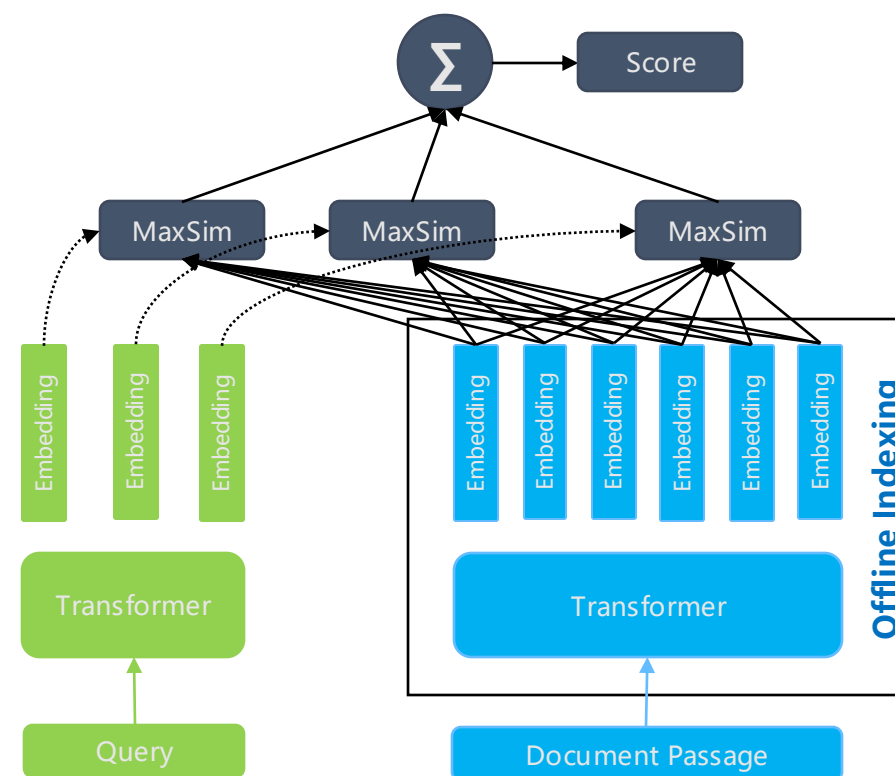
Dual Encoder



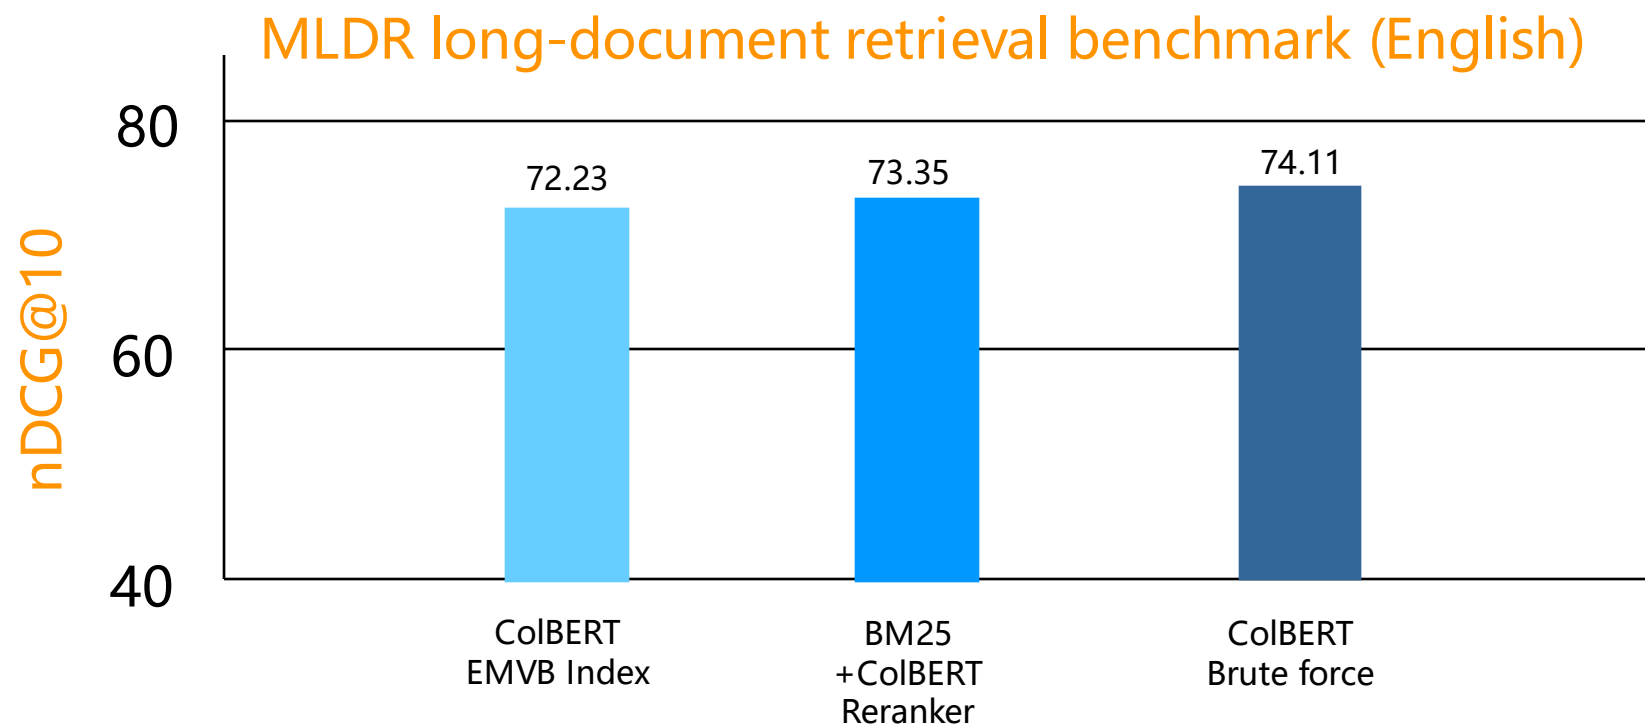
Cross Encoder



Late Interaction Encoder ColBERT



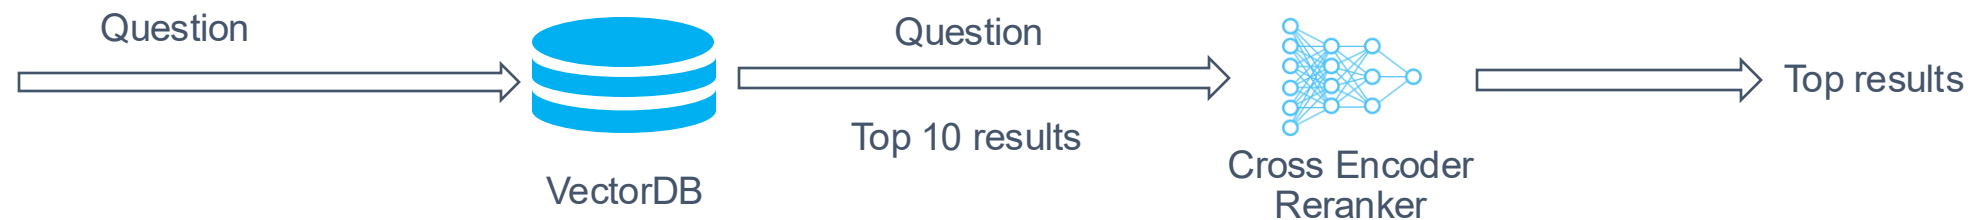
ColBERT Ranker vs Reranker



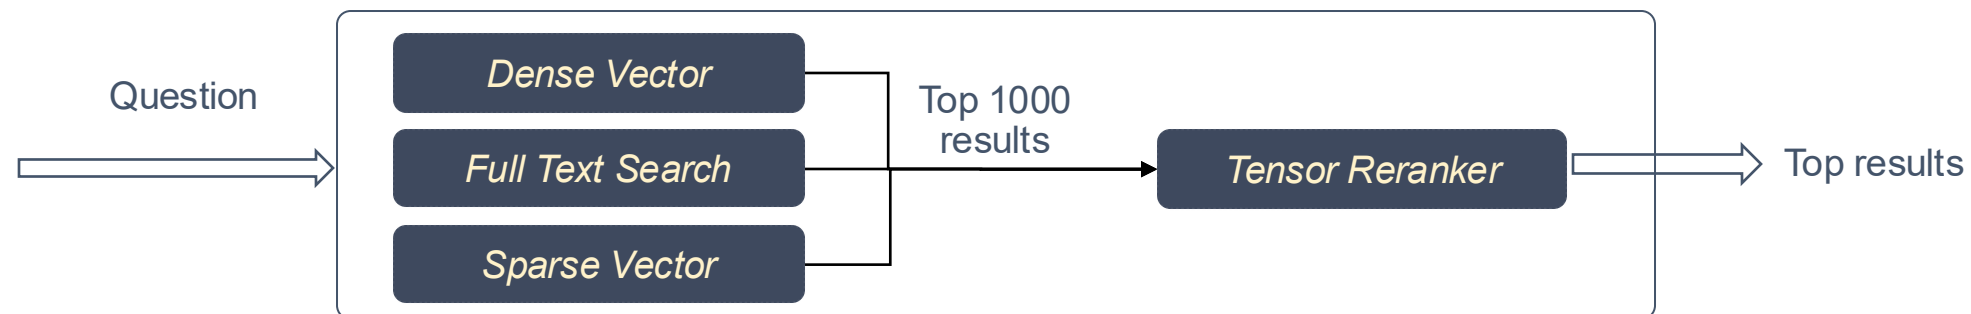
Embedding Model: BGE-M3



Benefits of Col-xxx Rerankers

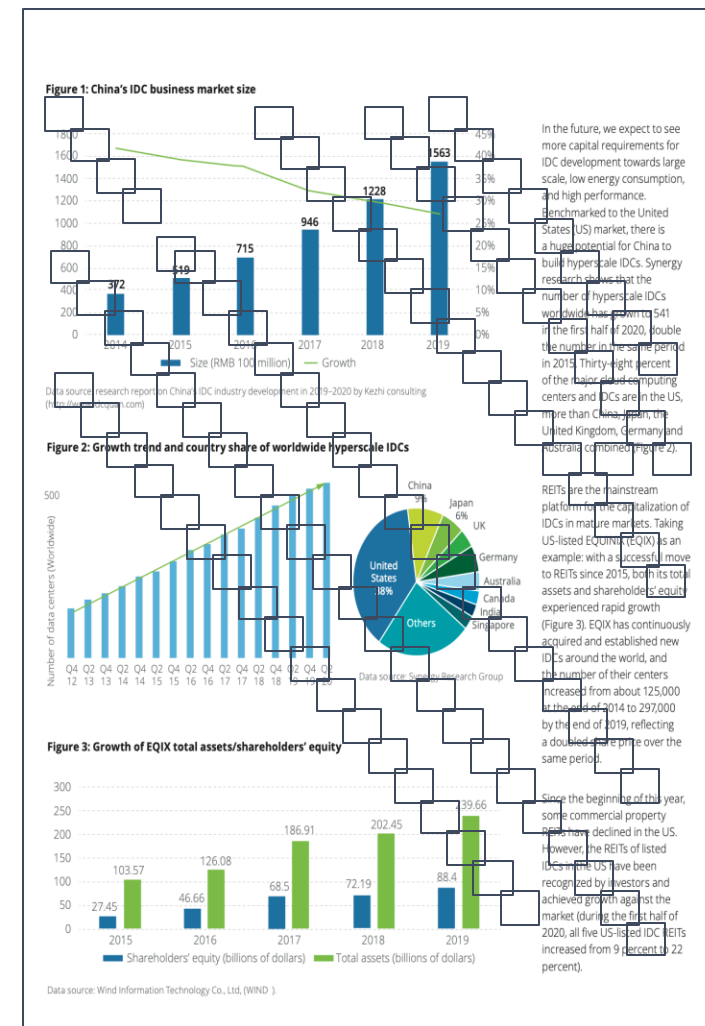


VS



Col-xxx in Multi-Modal-RAG

Divides a PDF page into $32 \times 32 = 1024$ patches
Each patch is represented as a 128-dim vector

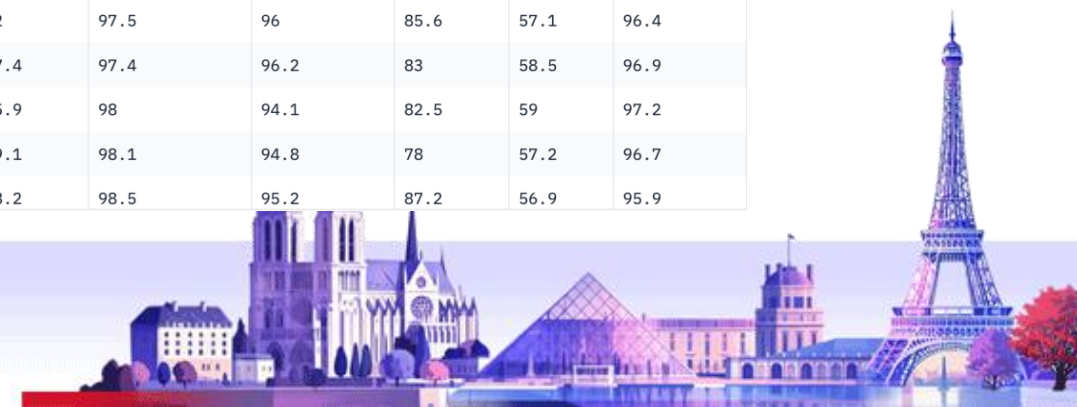


ViDoRe Benchmark

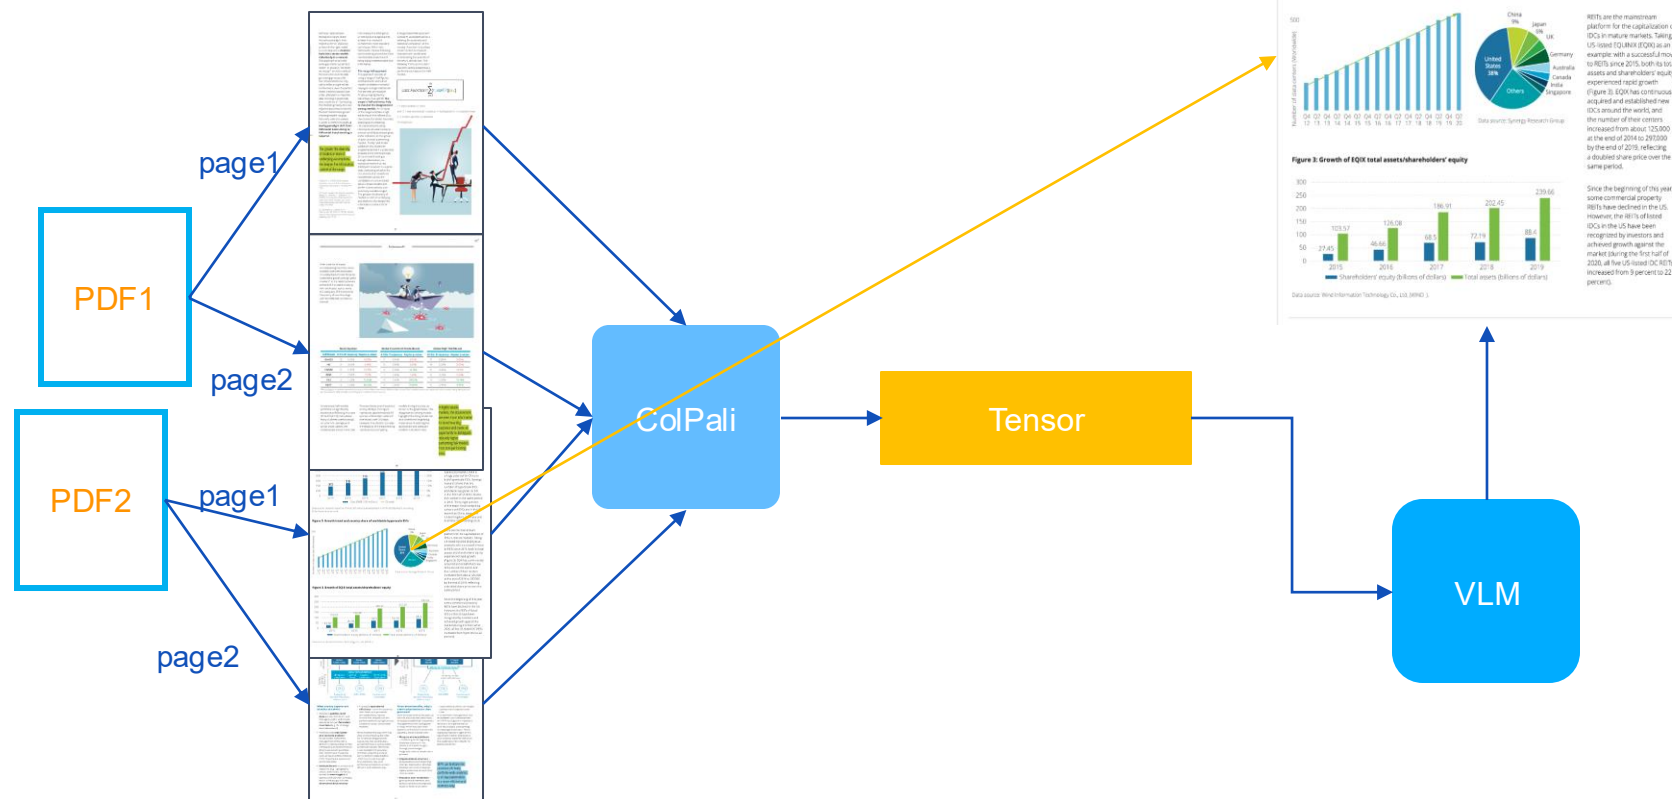


Qwen2 VL 7B => ColQwen2
PaliGemma => ColPali

Rank	Model	Average	TAT-DQA	Shift Project	Artificial Intelligence	Government Reports	ArxivQA	DocVQA	Healthcare Industry
1	Metric-AI_ColQwen2.5-7b-multitask	90.4	80.9	87.7	99.3	95.2	91.7	65.1	97.8
2	Metric-AI_ColQwen2.5-3b-multitask	90.3	80.7	88.4	98.9	96.4	92.2	64.4	98.4
3	yydxlv_colqwen2.5-7b-v0.1	90.2	81.9	89.1	98.9	96.2	91.1	63.1	98.5
4	tsystems_colqwen2-7b-v1.0	90.1	78.6	89.3	99.3	96.3	90.7	64.5	99.3
5	Metric-AI_colqwen2.5-3b-multitask	90	80.7	86.5	99.4	97.7	93	65.1	98.8
6	yydxlv_colqwen2-7b-v1.0	90	81.7	87.8	100	96.7	89.3	63.1	98.7
7	Alibaba-NLP_gme-Qwen2-VL-7B-1	89.9	76.1	97.3	99.6	99.5	87.6	56.6	99.6
8	Metric-AI_ColQwenStella-2b-multitask	89.8	78.7	88.1	99.6	97.1	89.5	62.6	99.6
9	tsystems_colqwen2-2b-v1.0	89.6	79.5	89.9	99.6	95.2	90	60.9	98.9
10	vidore_colqwen2.5-v0.2	89.4	81.1	87.3	99.6	96.4	89.2	63.2	97.9
11	vidore_colqwen2-v1.0	89.3	81.4	90.7	99.4	96.3	88.1	60.6	98.1
12	vidore_colqwen2.5-v0.1	88.8	80.8	85.5	99.3	95.3	88.5	61.9	98.8
13	Alibaba-NLP_gme-Qwen2-VL-2B-1	87.8	71.1	94.3	99	97.9	83.9	54.6	98.9
14	vidore_colqwen2-v0.1	87.3	75.9	86	98.7	92.8	86.1	61.4	98
15	vidore_colmolvlm-v0.1	86.1	79.5	79.5	98.1	96.9	79.4	60	99.6
16	MrLight_dse-qwen2-2b-mxl-v1	85.8	69.4	82	97.5	96	85.6	57.1	96.4
17	vidore_colpali-v1.3	84.8	70.4	77.4	97.4	96.2	83	58.5	96.9
18	vidore_colpali2-3b-pt-448	84.5	68.6	75.9	98	94.1	82.5	59	97.2
19	vidore_colpali-v1.2	83.9	68	79.1	98.1	94.8	78	57.2	96.7
20	yydxlv_colphi3.5	83.7	73.1	68.2	98.5	95.2	87.2	56.9	95.9



Col-xxx based Multi-Modal RAG



What's China's IDC business market size in 2018?

Answer: 1228



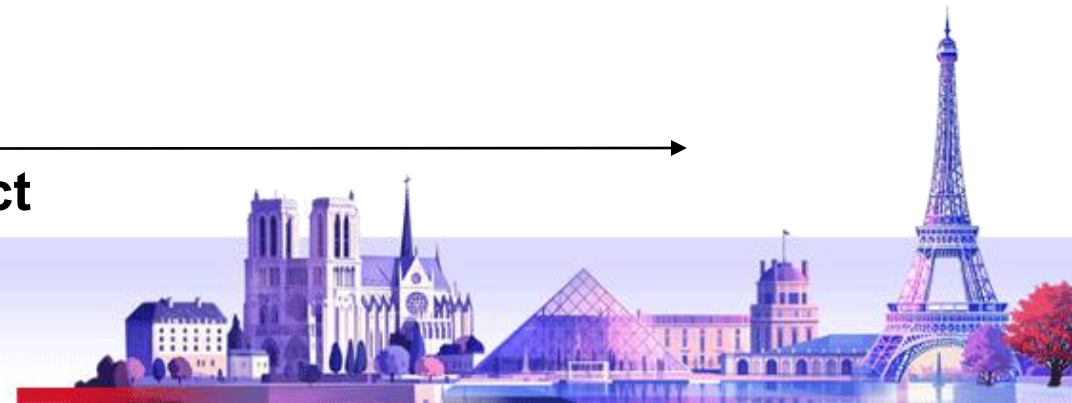
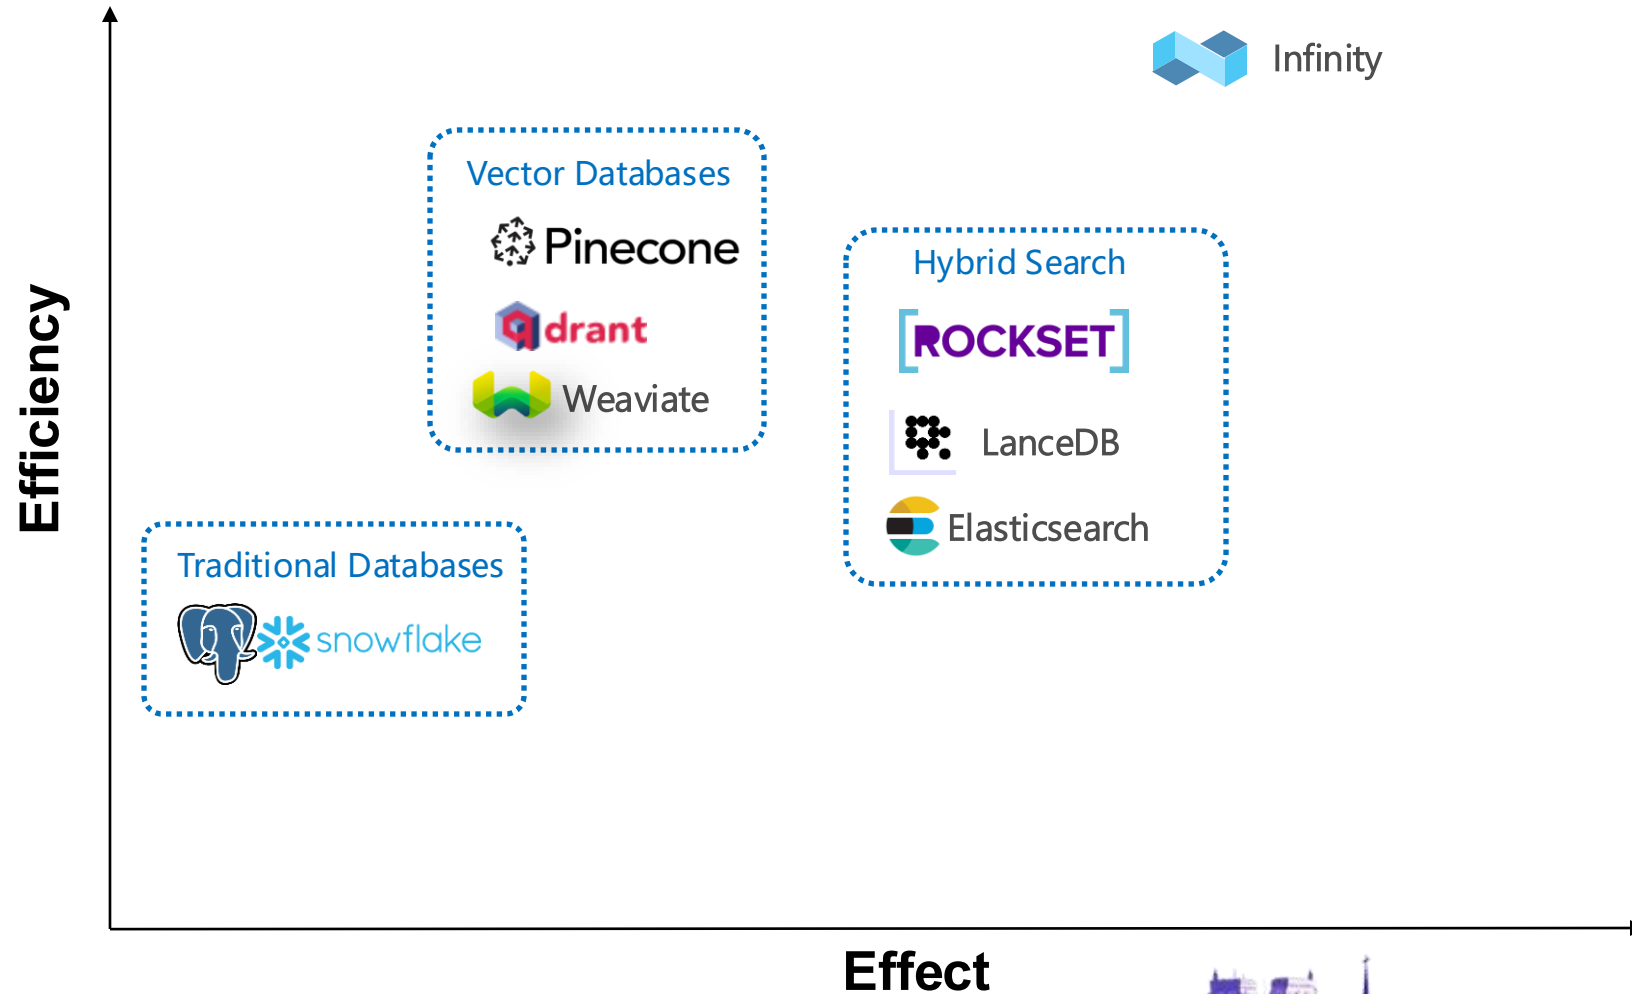
Scaling of Col-xxx Support

- Data base Side
- Tensor Reranker
 - Binary Quantization
 - Multi-Vector Index
 - Full text index

- Model Side
- Dimensional Reduction (MRL)
 - Tokens Reduction (ConstBERT/SVD)



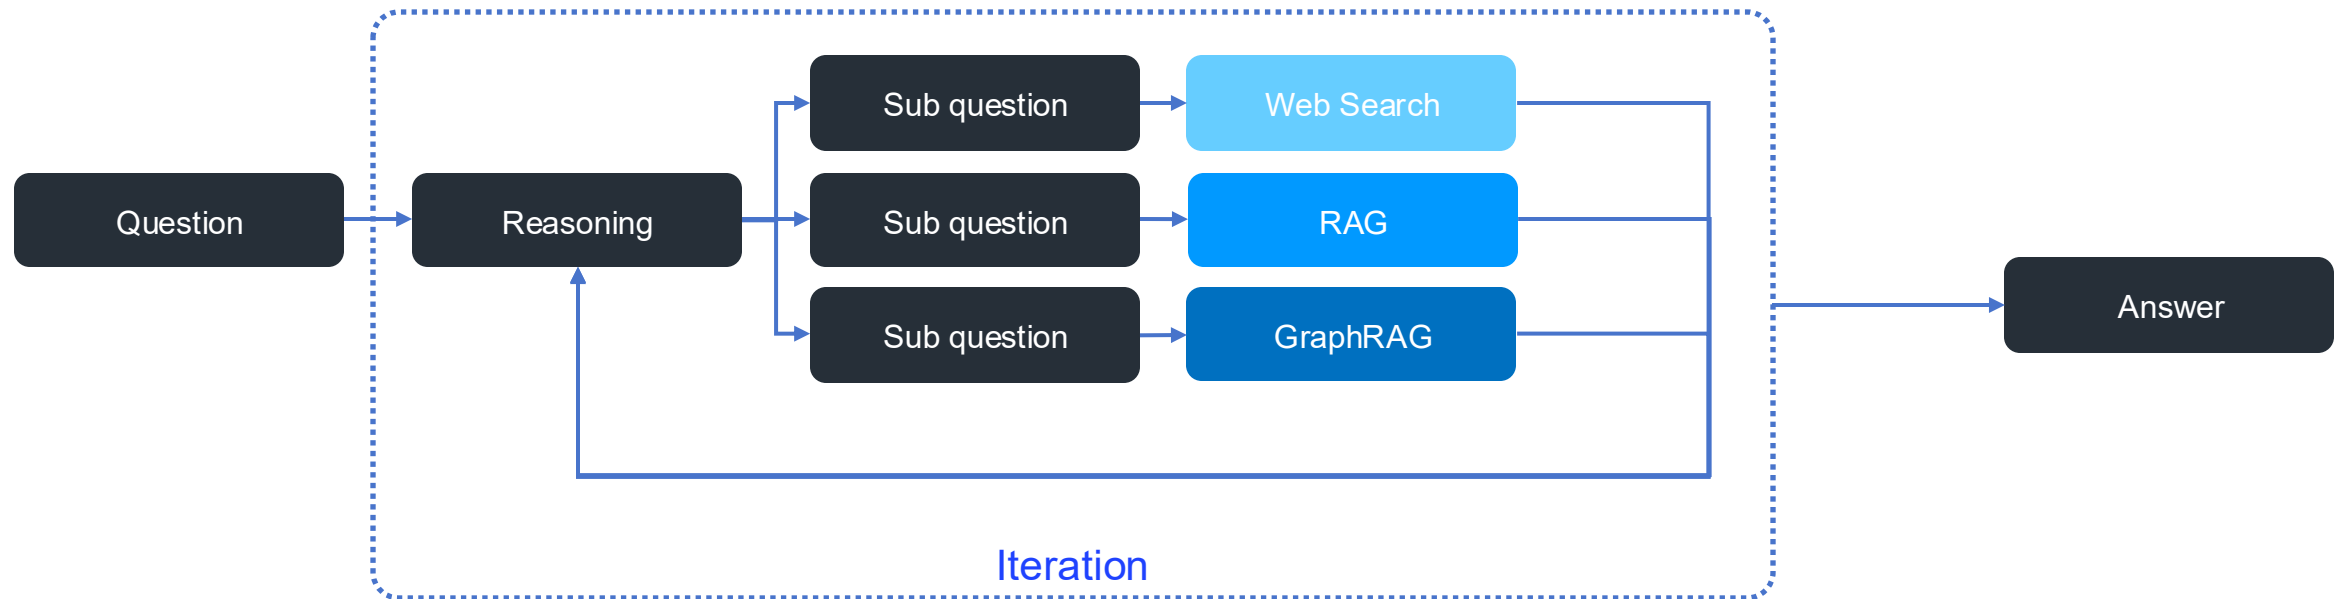
Comparison



Agentic and Future



Agentic Reasoning



Agentic Reasoning



Musk and Trump, who has more sons?



To answer the question, I need to determine how many sons Elon Musk and Donald Trump have. Let's break this down step by step.

Find the number of sons Elon Musk has.

Find the number of sons Donald Trump has.

Compare the two numbers to determine who has more sons.

Let's start with Elon Musk.

How many sons does Elon Musk have? →

Final Information

Elon Musk has a total of 12 children, including sons and daughters. Here is the breakdown of his sons:

With Justine Wilson (first wife):
Musk has six children with Justine Wilson, including twins and triplets. Among these, there are sons, but the exact number of sons is not specified in the provided documents.

With Grimes (ex-girlfriend):
Musk has three children with Grimes:
X Æ A-12 (or X Æ A-Xii), born in May 2020 (son).
Techno Mechanicus (nicknamed **Tau**), born in 2022 (son).
Exa Dark Sideræl, born in December 2021 (daughter).
Total sons with Grimes: 2.

With Shimon Zilis (Neuralink executive):
Musk has three children with Shimon Zilis, including twins born in November 2021. The exact number of sons is not specified in the provided documents.

With Ashley St. Clair (author):
Musk reportedly has one child with Ashley St. Clair, but the gender is not specified in the provided documents.

Based on the available information, Elon Musk has at least **2 confirmed sons** (X Æ A-12 and Techno Mechanicus) with Grimes. The exact number of sons with Justine Wilson, Shimon Zilis, and Ashley St. Clair is not clearly stated in the documents. Therefore, the total number of sons Elon Musk has is at least 2, but it could be higher depending on the gender distribution of his other children. Now, I will search for information about Donald Trump's sons to compare with Elon Musk's sons.

How many sons does Donald Trump have? →

Final Information

Donald Trump has **three sons**: Donald Trump Jr., Eric Trump, and Barron Trump. This information is confirmed across multiple sources, including:

Business Insider: "Former President Donald Trump has two sons, Donald Jr. and Eric, and a daughter, Ivanka, with his first wife, Ivana."
Today Show: "Donald Trump is a father of five children: Don Jr., Eric, Ivanka, Tiffany, and Barron."
VPM News: "Donald Trump Jr., Trump's oldest son, is 47 and his mother is Ivana Trump. Three of Trump's children — Tiffany, Ivanka, and Eric — are currently..."
ABC News: "Donald Trump is a dad of 5, grandfather of 10: What to know about his family."
Wikipedia: The family tree confirms Donald Trump Jr., Eric Trump, and Barron Trump as his sons.

Thus, Donald Trump has **three sons**: Donald Trump Jr., Eric Trump, and Barron Trump. Based on the search results:

Elon Musk has at least **2 confirmed sons** (X Æ A-12 and Techno Mechanicus), with the possibility of more sons among his other children.

Donald Trump has **3 confirmed sons** (Donald Trump Jr., Eric Trump, and Barron Trump).

Therefore, Donald Trump has more sons than Elon Musk.

Based on the information from the knowledge base:

Elon Musk has at least **2 confirmed sons** (X Æ A-12 and Techno Mechanicus), with the possibility of more sons among his other children.

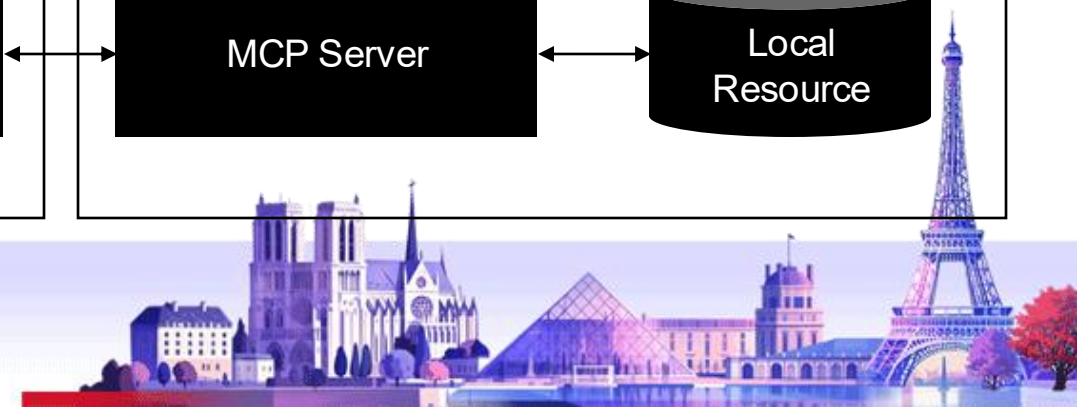
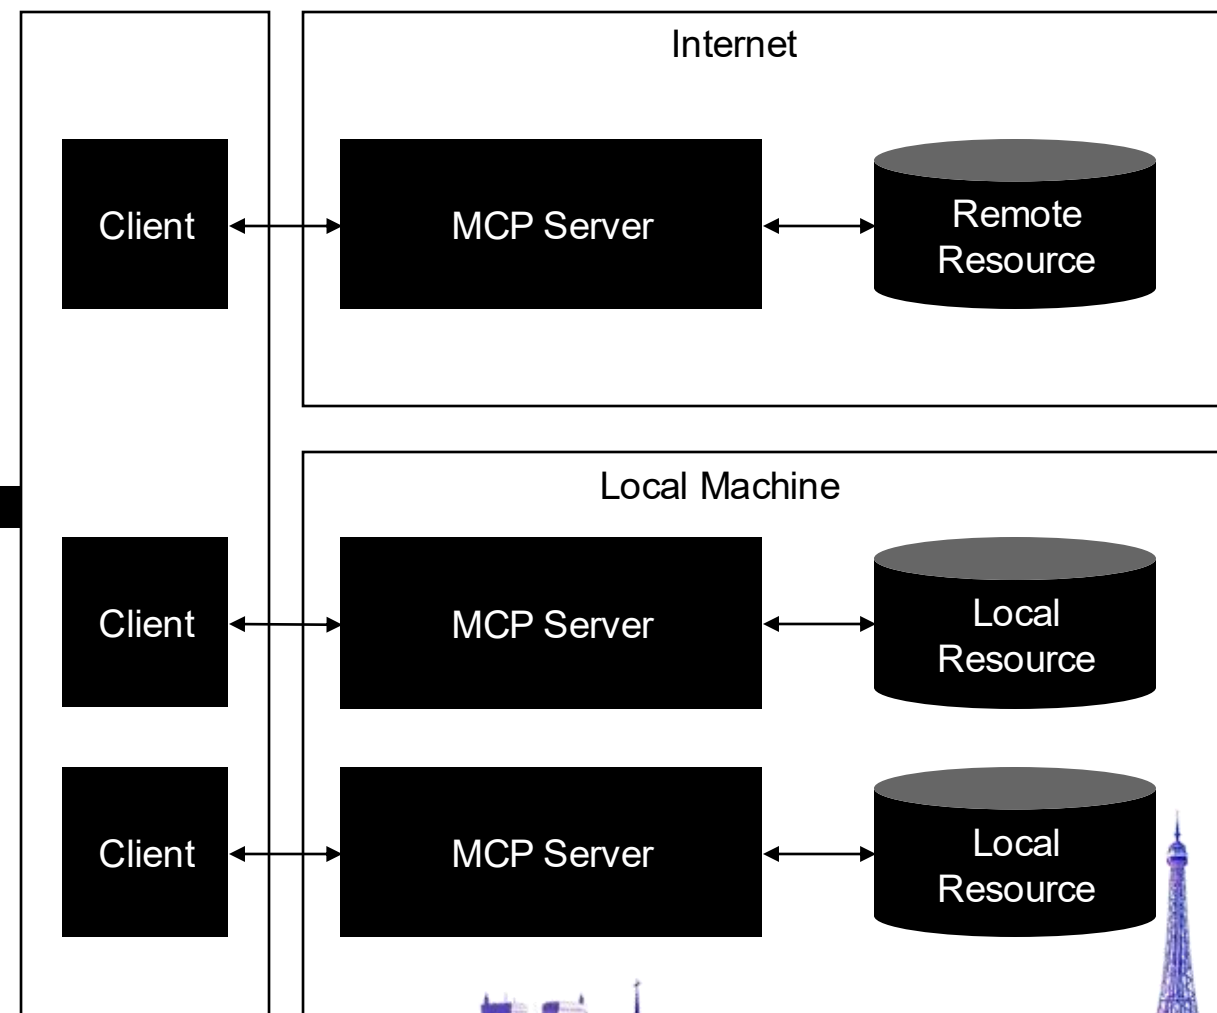
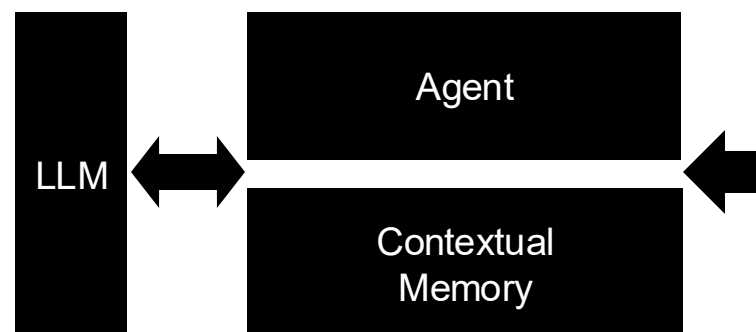
Donald Trump has **3 confirmed sons** (Donald Trump Jr., Eric Trump, and Barron Trump). ⓘ

Therefore, **Donald Trump has more sons than Elon Musk.**

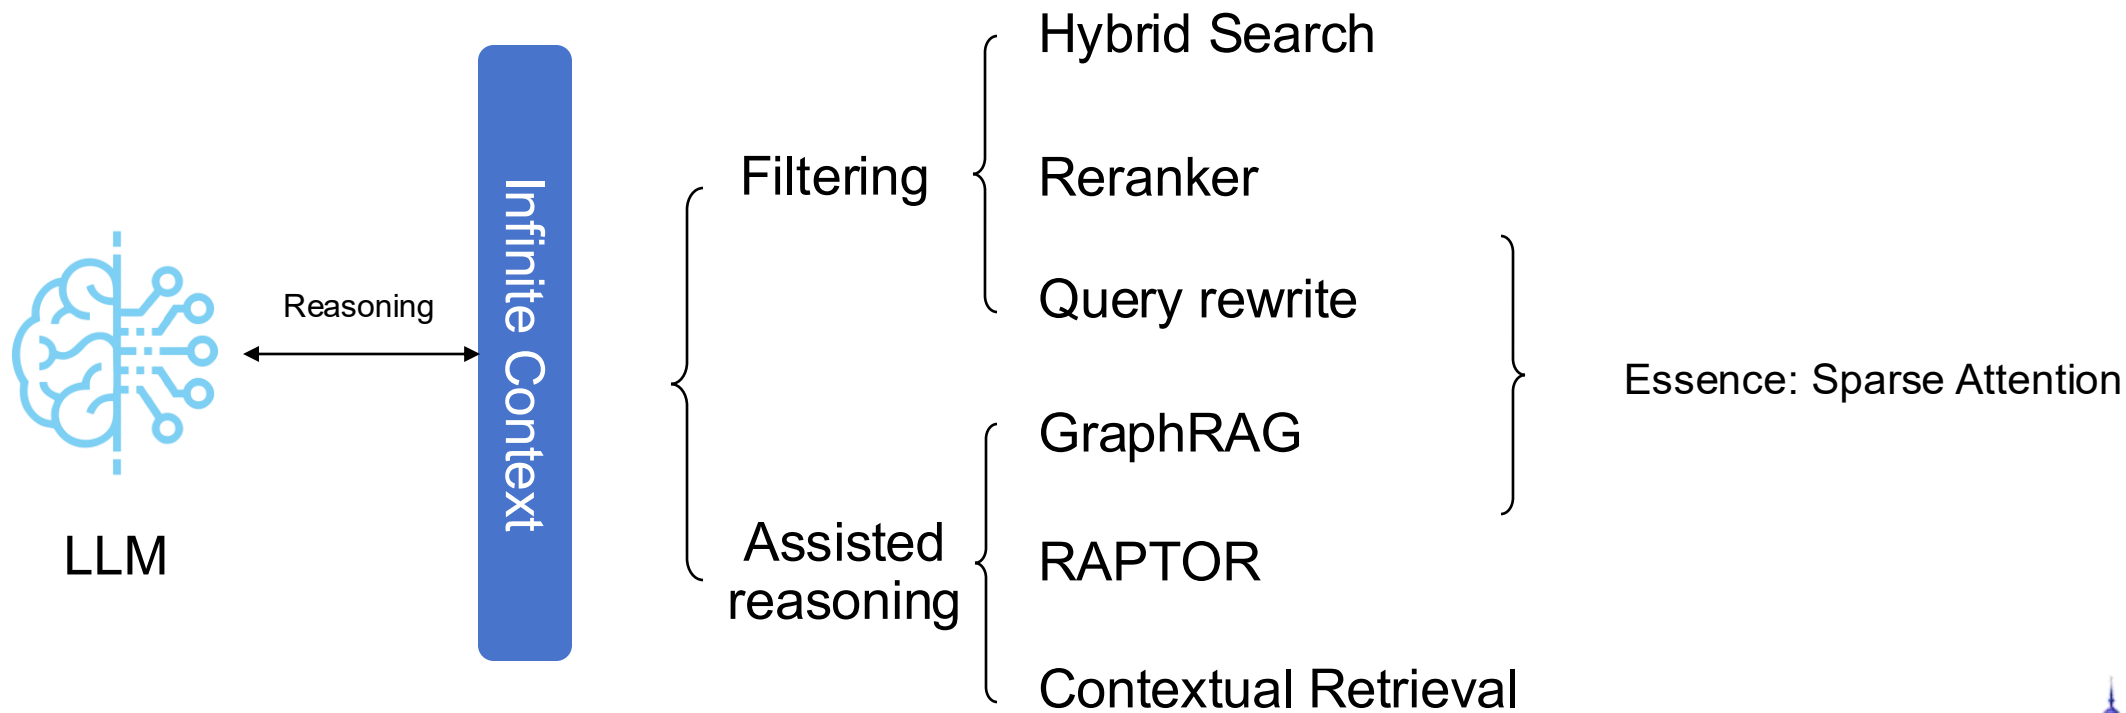


RAG ↔ Agent

- RAG As a Distinct Layer Grows More Vital
- RAG & Agent:
Frequent Interaction, Collaboration and Competition Coexist

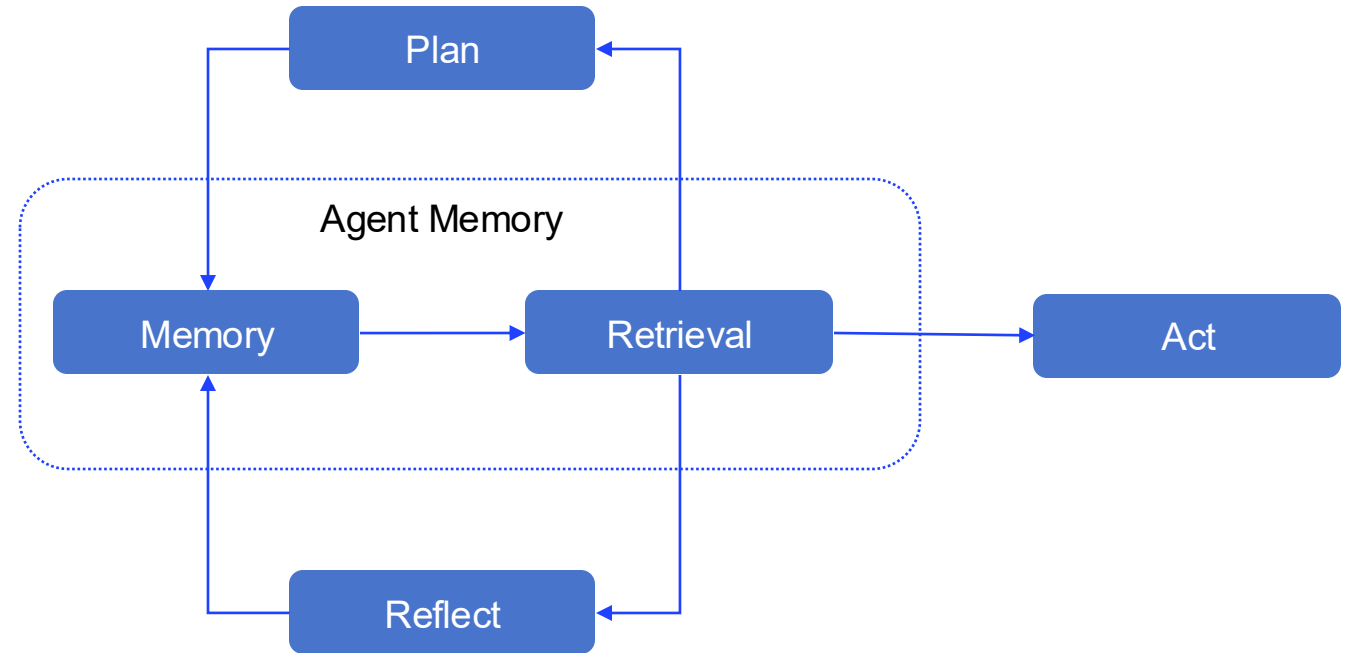


The Essence of Memory

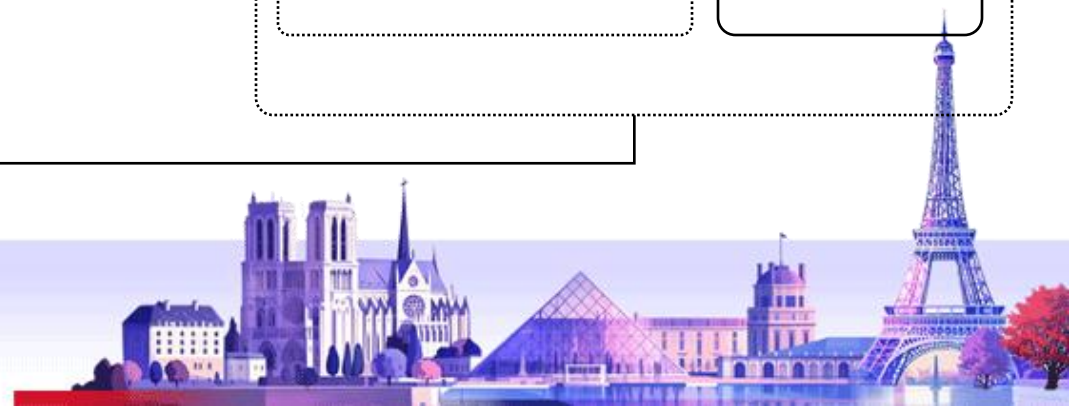
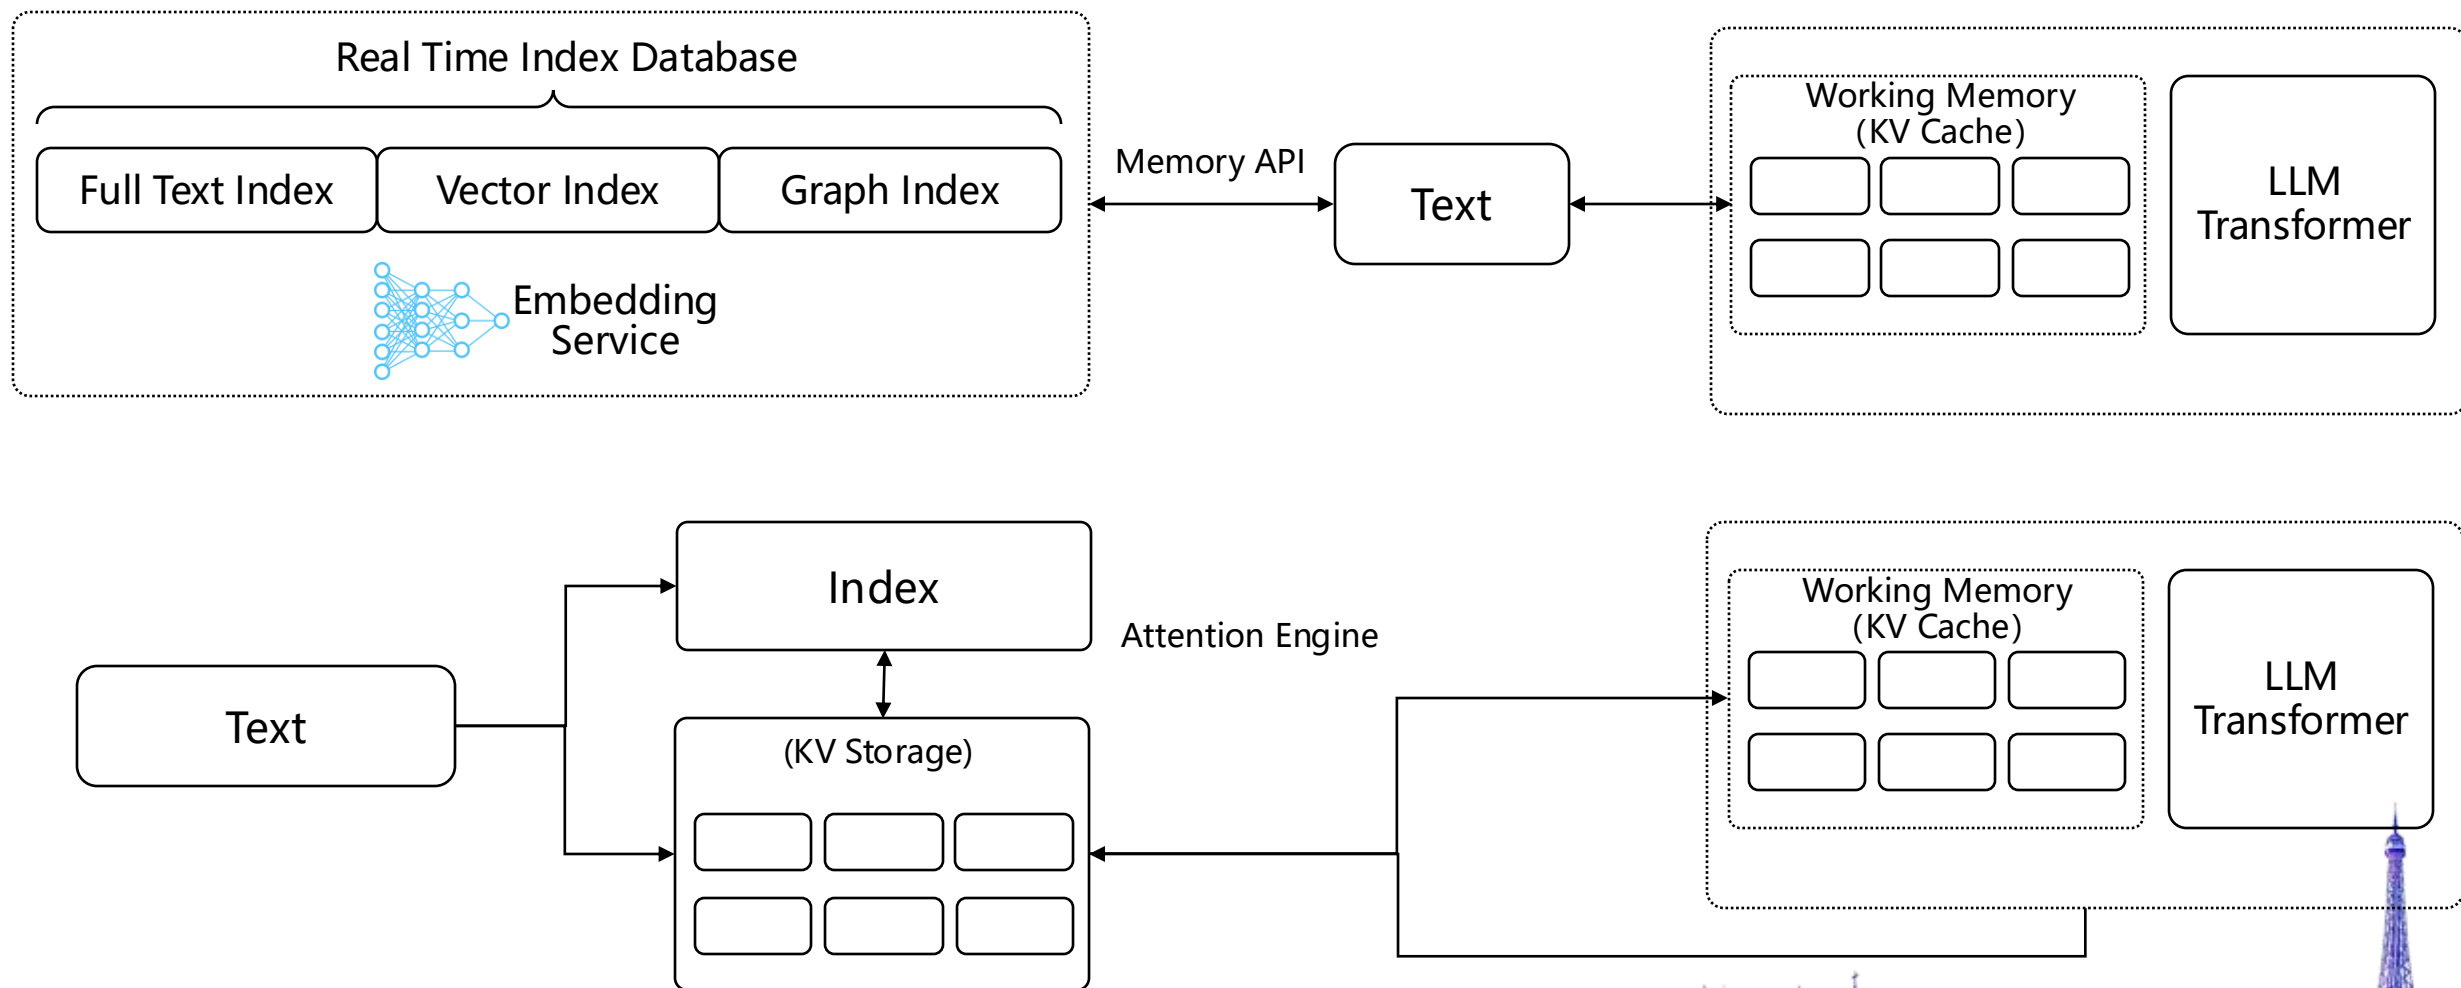


Agent Memory = RAG ?

- Retrieval Differences—Temporal/Weights
- Memory Decay
- Agent Facilities



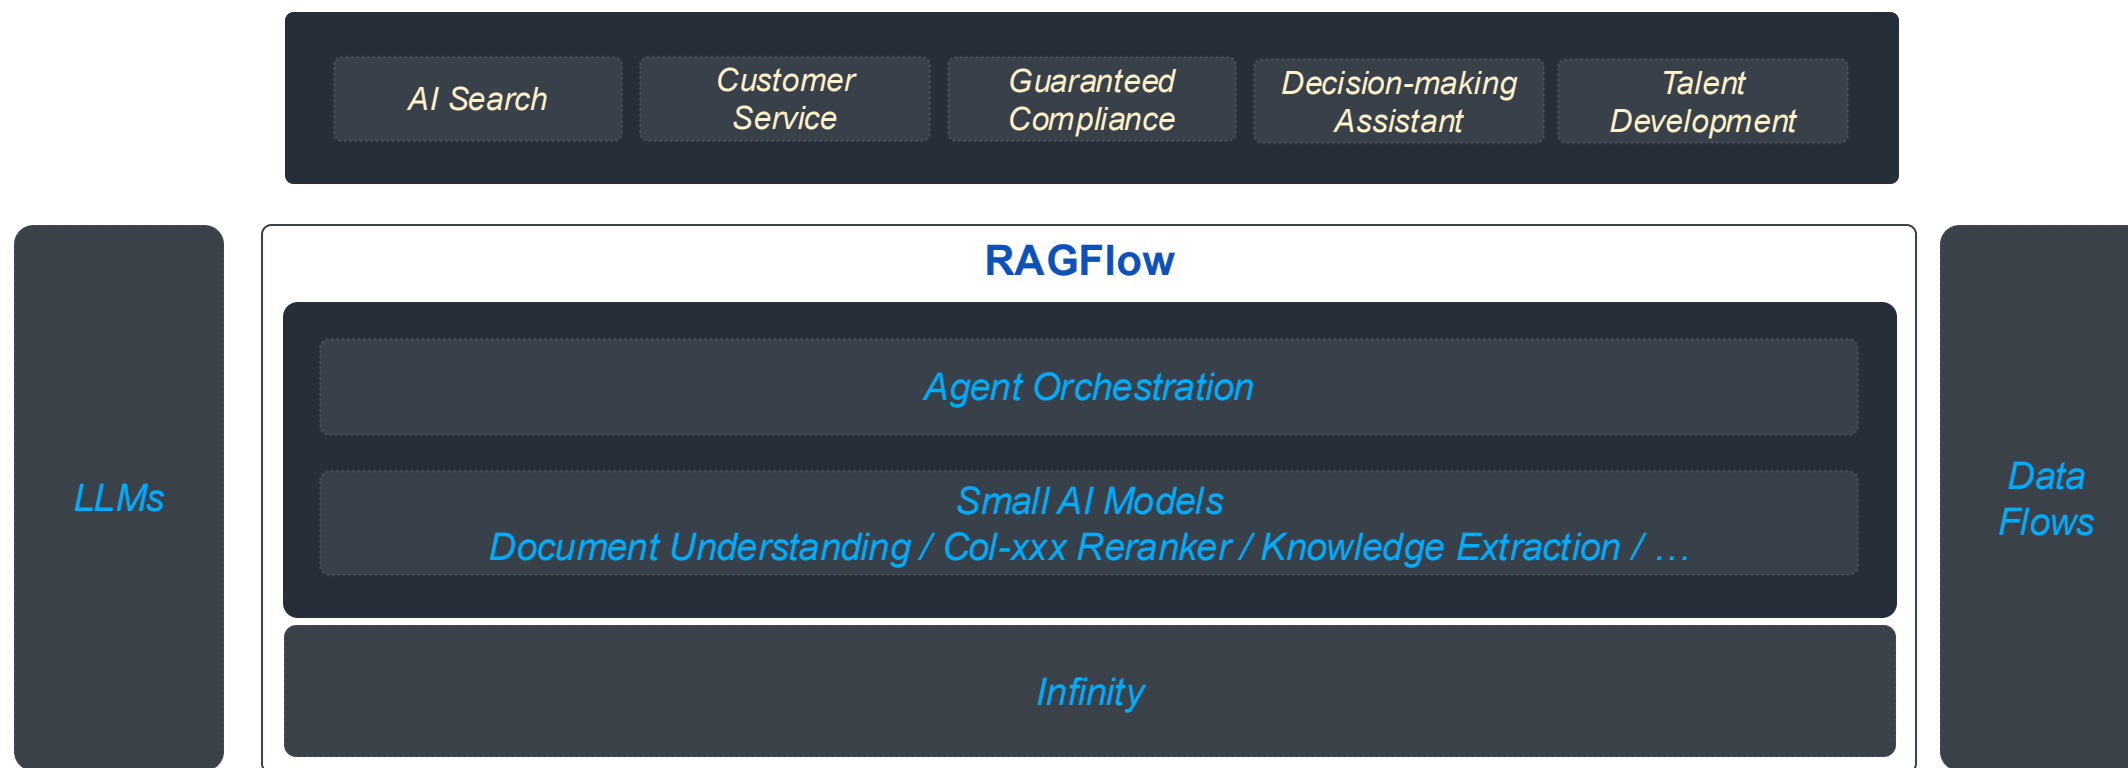
RAG/Memory Evolution



Summary



RAGFlow Ecosystem



RAG is not LLMOps, RAG is the database of LLM era



THANK YOU

<https://github.com/infiniflow/ragflow>
<https://github.com/infiniflow/infinity>

